

Gonka Agent Workload Pricing Analysis

Version: 1.0 **Date:** 2026-04-01 **Classification:** Internal – for leadership decision-making **Requirement:** COMP-02

1. Executive Summary

OpenClaw agent workloads cost 3-10x more than equivalent chatbot interactions due to three compounding factors: heartbeat context resending (full system prompt + workspace context resent every 30 minutes = 48 calls/day minimum), multi-step tool call chains (3-10 tool calls per user message), and always-on operation (agents run 24/7, not on-demand). **The dominant cost driver is heartbeats, not message volume.** A single agent with 30-minute heartbeats consumes 461K tokens/day on heartbeats alone – nearly equal to 50 user messages worth of tokens.

This analysis models three realistic OpenClaw agent workload tiers (Casual, Active, Heavy) across five providers: Together AI, DeepInfra, OpenRouter, OpenAI, and Anthropic. Three hypothetical Gonka pricing scenarios (A/B/C) are modeled to show leadership where cost advantages emerge at different price points.

Key findings:

- At the Casual tier (~31.5M tokens/month), K2.5 providers cluster around \$43-52/month. Gonka's session persistence could reduce effective cost to \$10-17/month by eliminating heartbeat context resending – a structural advantage no other provider offers.
- At the Heavy tier (~2.94B tokens/month), heartbeat overhead accounts for 85% of total token consumption with 5-minute intervals. Gonka's session persistence advantage scales super-linearly: the more frequent the heartbeats, the larger the savings.
- OpenAI and Anthropic's prompt caching partially mitigates heartbeat costs (reducing them by 50% and 90% respectively), but their premium model pricing (\$2.50-\$15.00/1M input) still results in 3-20x higher monthly spend than K2.5 providers for equivalent workloads.
- **Gonka's pricing is TBD.** All Gonka cost projections in this document are HYPOTHETICAL scenarios. The cost advantage narrative depends entirely on final pricing decisions and the operability of server-side session persistence.

Cross-reference: See output/gonka_competitive_feature_matrix.md for the feature matrix comparing Gonka across agent-relevant dimensions including pricing model.

2. Methodology

Why Per-Task Cost Matters More Than Per-Token Cost

Traditional inference pricing comparisons focus on \$/1M tokens – the headline rate on every provider's pricing page. This metric is misleading for agent workloads because it ignores the structural overhead that agents impose: context resending, multi-step tool chains, and always-on heartbeats.

A chatbot sends a message and receives a response. An OpenClaw agent sends a message, receives a response, executes 3-10 tool calls (each an additional LLM request), and then resends its entire context every 30 minutes regardless of whether the user sent anything. **The true cost of an agent workload is a function of architecture, not just token price.**

How OpenClaw Agent Cost Is Calculated

Each inference call in an OpenClaw agent includes:

1. **System prompt + workspace context:** ~9,600 tokens per turn. This includes the agent's personality configuration, tool definitions, workspace state, and conversation history. Source: OpenClaw pricing guides and community analysis (aicost.org, clawback.tools).
2. **Heartbeat overhead:** OpenClaw agents send periodic "heartbeat" messages to maintain context and check for updates. Each heartbeat resends the full system prompt + context (~9,600 tokens). At 30-minute intervals, this is 48 heartbeats/day. At 5-minute intervals (for monitoring-heavy agents), this is 288 heartbeats/day.

3. **Tool call chains:** Each user message triggers an average of 3-10 tool calls. Each tool call is a separate LLM inference request averaging ~500 tokens input + ~200 tokens output (700 tokens total). Source: LangChain State of Agent Engineering report; OpenClaw community benchmarks.

The Three Workload Tiers

This analysis models three tiers representing real OpenClaw deployment patterns observed in community reports and pricing guides:

- **Casual:** A solo developer running one agent on one channel. Budget model, minimal heartbeats.
- **Active:** A small team running multiple agents across channels. Mixed model routing.
- **Heavy:** A business running always-on agent infrastructure. Premium models, aggressive heartbeats, high message volume.

Assumptions

- **Input:output token ratio:** 1:1 blended unless provider-specific data indicates otherwise. Agent workloads tend to be input-heavy (large context, small responses), so actual costs may skew lower for providers with cheaper input tokens. Where this materially affects the analysis, it is noted.
- **Heartbeat tokens:** Each heartbeat resends the full system prompt + workspace context (~9,600 tokens). No provider-specific caching is applied to the “raw” cost calculation; caching adjustments are applied separately in the Hidden Cost Analysis (Section 5).
- **Tool call tokens:** Average 500 tokens input + 200 tokens output per tool call (700 tokens total). This is conservative; complex tool chains (web scraping, database queries) can be significantly higher.
- **All prices as of April 2026.** Inference pricing is volatile; figures should be revalidated before use in external communications.

3. Workload Tier Definitions

Casual Tier

Profile: A solo developer running a personal assistant agent on Telegram or Discord.

Parameter	Value
Agents	1
Channels per agent	1 (e.g., Telegram)
User messages per day	50
Heartbeat interval	30 minutes
Heartbeats per day	48
Tokens per turn (system prompt + context)	~9,600
Tool calls per user message	3 (average)
Tokens per tool call	~700 (500 input + 200 output)
Model tier	Single model (budget tier)

Daily token breakdown:

Component	Calculation	Daily Tokens
User messages	50 messages x 9,600 tokens	480,000
Heartbeats	48 heartbeats x 9,600 tokens	460,800
Tool calls	150 calls x 700 tokens	105,000
Daily total		1,045,800

Monthly total: ~31.5M tokens (1,045,800 x 30 days)

Observation: Heartbeats account for 44% of daily token consumption despite generating zero user-facing value. At this tier, heartbeat cost nearly equals message cost.

Active Tier

Profile: A small team (e.g., crypto project, small startup) running multiple community management and support agents.

Parameter	Value
Agents	3
Channels per agent	2 (e.g., Telegram + Discord)
Total channel-agents	6
User messages per day (total)	200
Heartbeat interval	30 minutes
Heartbeats per day	$48 \times 6 = 288$
Tokens per turn	~9,600
Tool calls per user message	5 (average)
Tokens per tool call	~700
Model routing	70% budget, 30% strong model

Daily token breakdown:

Component	Calculation	Daily Tokens
User messages	200 messages x 9,600 tokens	1,920,000
Heartbeats	288 heartbeats x 9,600 tokens	2,764,800
Tool calls	1,000 calls x 700 tokens	700,000
Daily total		5,384,800

Monthly total: ~161.5M tokens (5,384,800 x 30 days)

Observation: Heartbeats now dominate at 51% of daily token consumption. With 6 channel-agents each heartbeating independently, the overhead scales linearly with channel count. Multi-model routing (70/30 split) affects cost but not token volume – this is reflected in the cost projections.

Heavy Tier

Profile: A business or DAO running production-grade agent infrastructure for customer support, monitoring, trading, or community management.

Parameter	Value
Agents	10
Channels per agent	3 (e.g., Telegram + Discord + Web)
Total channel-agents	30
User messages per day (total)	1,000+
Heartbeat interval	5 minutes (aggressive monitoring)
Heartbeats per day	$288 \times 30 = 8,640$
Tokens per turn	~9,600
Tool calls per user message	8 (average; complex workflows)
Tokens per tool call	~700
Model routing	Mixed: budget for simple, premium for complex

Daily token breakdown:

Component	Calculation	Daily Tokens
User messages	1,000 messages x 9,600 tokens	9,600,000
Heartbeats	8,640 heartbeats x 9,600 tokens	82,944,000
Tool calls	8,000 calls x 700 tokens	5,600,000
Daily total		98,144,000

Monthly total: ~2.94B tokens (98,144,000 x 30 days)

Observation: Heartbeats account for **84.5%** of all token consumption at this tier. The 5-minute heartbeat interval (6x more frequent than the 30-minute default) combined with 30 channel-agents creates massive overhead. This is the tier where Gonka's server-side session persistence would deliver the most dramatic cost savings – and where the savings gap widens versus every other provider.

4. Per-Provider Pricing Data

All pricing data below was gathered from provider pricing pages, aggregator sites, and independent analyses. Every figure includes its source, date of access, and confidence level.

K2.5 Pricing Across Providers (April 2026)

Provider	Model	Input/1M tokens	Output/1M tokens	Blended/1M (1:1 ratio)	Source	Date	Confidence
Together AI	Kimi K2.5	\$0.50	\$2.50	\$1.50	together.ai/pricing	Apr 2026	MEDIUM
DeepInfra	Kimi K2.5	\$0.45	\$2.25	\$1.35	Artificial Analysis (artificial-analysis.ai)	Apr 2026	MEDIUM
OpenRouter	Kimi K2.5 (via providers)	~\$0.60	~\$2.50	~\$1.55	openrouter.ai + 5.5% credit markup	Apr 2026	MEDIUM

Frontier Model Pricing (April 2026)

Provider	Model	Input/1M tokens	Output/1M tokens	Blended/1M (1:1 ratio)	Source	Date	Confidence
OpenAI	GPT-4o	\$2.50	\$10.00	\$6.25	openai.com/pricing	Apr 2026	HIGH
Anthropic	Claude Sonnet 4	\$3.00	\$15.00	\$9.00	anthropic.com/pricing	Apr 2026	HIGH
Anthropic	Claude Opus 4	\$15.00	\$75.00	\$45.00	anthropic.com/pricing	Apr 2026	HIGH

Gonka Pricing Scenarios (HYPOTHETICAL)

Gonka's per-token pricing has not been set. The following three scenarios are modeled to show leadership how different price points affect competitive positioning. **None of these are actual prices.**

Scenario	Model	Input/1M tokens	Output/1M tokens	Blended/1M (1:1 ratio)	Basis	Confidence
Gonka Scenario A	Kimi K2.5	\$0.25	\$1.25	\$0.75	HYPOTHETICAL – 50% below DeepInfra	SCENARIO
Gonka Scenario B	Kimi K2.5	\$0.35	\$1.75	\$1.05	HYPOTHETICAL – 30% below DeepInfra	SCENARIO
Gonka Scenario C	Kimi K2.5	\$0.45	\$2.25	\$1.35	HYPOTHETICAL – price match DeepInfra	SCENARIO

Scenario rationale: - **Scenario A (aggressive):** 50% below the cheapest verified provider (DeepInfra at \$0.45/\$2.25). This would require Gonka's decentralized compute cost advantage to be near the upper end of claimed savings (60-80% cheaper than centralized). Risks: margin pressure, unsustainability if compute costs are higher than projected. - **Scenario B (moderate):** 30% below DeepInfra. This is the recommended target – competitive on raw price while leaving margin for network sustainability. When combined with session persistence savings, total cost becomes lower than any competitor. - **Scenario C (conservative):** Price match with DeepInfra. Gonka competes on features (sessions, memory, privacy) rather than price. This scenario tests whether non-price differentiators alone justify adoption.

5. Hidden Cost Analysis

Headline per-token rates tell an incomplete story. Each provider has structural costs and advantages that materially affect the true cost of running OpenClaw agents.

OpenRouter

5.5% credit markup on all usage. When a developer buys \$100 in OpenRouter credits, they receive \$94.50 in inference value. This markup applies on top of the underlying provider's per-token rate.

- **Impact on Casual tier:** \$49/month raw becomes ~\$52/month after markup
- **No prompt caching.** Every heartbeat resends the full context at full price. OpenRouter routes to underlying providers, and caching is not exposed through the routing layer. This means heartbeat costs are fully unoptimized.
- **Free tier degradation.** Free-tier requests are queued behind paid users. Models may become unavailable or switch to lower-quality providers during peak demand. For always-on agents, this creates unpredictable availability.
- **Rate limiting on high volume.** High-volume agent workloads can trigger rate limits, causing throttled requests. Throttled requests create agent idle time – the agent waits for a response that never comes, then retries. This wasted time is not billed in tokens but costs developer productivity and agent reliability.
- **Net hidden cost: +5.5% on raw token spend, plus unquantifiable reliability costs.**

OpenAI

Prompt caching: automatic, 50% discount on cached tokens. OpenAI automatically caches prompts longer than 1,024 tokens. Since OpenAI system prompts are ~9,600 tokens, heartbeats benefit significantly – the system prompt portion is cached, and only the delta (new messages since last call) incurs full price.

- **Impact on heartbeats:** The ~9,600 token system prompt is cached after the first call. Subsequent heartbeats pay 50% on the cached portion. Estimated heartbeat cost reduction: ~50% of the system prompt portion (roughly 40-50% of total heartbeat cost, since workspace context may vary).
- **Volume discounts:** Available through enterprise agreements but require minimum commitments. Not accessible to Casual or Active tier users.
- **No markup:** Direct API access, no intermediary fees.
- **Premium model pricing:** GPT-4o at \$2.50/\$10.00 is 5-8x more expensive per token than K2.5 on Together AI or DeepInfra. Even with caching, absolute costs remain significantly higher than budget K2.5 providers.
- **Net hidden cost adjustment: -40% on heartbeat costs (prompt caching), but base price is 4-5x higher than K2.5.**

Anthropic

Prompt caching: 90% discount on cache reads, 25% write premium on first fill. Anthropic's caching is more aggressive than OpenAI's but has an initial write penalty.

- **Impact on heartbeats:** First heartbeat in a session costs 25% more (cache write premium). Subsequent heartbeats within the cache TTL cost 90% less on cached portions. For always-on agents with consistent system prompts, the net effect is dramatic savings after the first call.
- **Effective heartbeat cost (steady state):** ~10% of uncached price for the system prompt portion. With ~80% of heartbeat tokens being cacheable system prompt, steady-state heartbeat cost drops by approximately 70-75%.
- **Rate limits tighter than OpenAI.** Anthropic imposes stricter concurrent request limits, which can bottleneck Heavy tier workloads with many simultaneous agents. Rate limit overages may require waiting or enterprise tier negotiation.
- **Premium model pricing:** Claude Sonnet 4 at \$3.00/\$15.00 is competitive with GPT-4o but still 6-10x more per token than K2.5. Claude Opus 4 at \$15.00/\$75.00 is the most expensive option in this analysis.
- **Net hidden cost adjustment: -70% on heartbeat costs after initial fill, but base price is 6-30x higher than K2.5.**

Together AI

No markup, no caching, straightforward per-token pricing. Together AI operates its own GPU clusters and passes compute costs directly to users without intermediary fees.

- **No credit markup:** \$100 spent = \$100 of inference.
- **No prompt caching:** Every request pays full per-token rates. Heartbeats resend full context at full price. This is the same limitation as OpenRouter, but without the 5.5% markup.
- **Lowest verified K2.5 pricing:** \$0.50/\$2.50 (input/output) is the second-cheapest after DeepInfra (\$0.45/\$2.25).

- **Rate limits:** No reported rate limit concerns at current usage scales. Together AI’s infrastructure handles high-volume inference well.
- **Net hidden cost adjustment: none. What you see is what you pay.**

Gonka (HYPOTHETICAL)

Gonka’s architecture provides structural cost advantages that go beyond per-token pricing. These advantages are real technical capabilities built in v1.2, though their impact depends on final pricing and production deployment.

Server-side session persistence eliminates context resending. Gonka’s session API maintains conversation state server-side. Instead of resending the full ~9,600 token context on every heartbeat, agents send only the delta (new messages or state changes since the last call). This is architecturally equivalent to “permanent prompt caching” – but without TTL expiration and without the initial cache write premium.

- **Impact on heartbeats:** If sessions reduce heartbeat token overhead by 80% (sending only ~1,920 tokens of delta instead of 9,600 tokens of full context), monthly token consumption drops dramatically:
 - Casual tier: 31.5M tokens drops to ~12.6M tokens/month
 - Active tier: 161.5M tokens drops to ~56.3M tokens/month
 - Heavy tier: 2.94B tokens drops to ~0.61B tokens/month
- **Tiered routing:** Gonka’s multi-model routing sends 70% of requests to lite quantization (cheaper, faster) and 30% to full-precision when quality demands it. This reduces effective per-token cost without sacrificing quality on complex tasks.
- **No markup:** Decentralized network, no intermediary. GNK tokenomics may introduce different cost dynamics but no credit markup structure.
- **Session caveat (v1.2 tech debt):** Current implementation uses in-memory sessions, which are lost on server restart. This must be resolved (migrated to Redis or persistent storage) before the session persistence advantage can be marketed as production-ready.
- **Net hidden ADVANTAGE: ~60-80% reduction in heartbeat token overhead via sessions (contingent on production implementation), plus tiered routing savings.**

Hidden Cost Summary

Provider	Markup	Prompt Caching	Session Persistence	Rate Limits	Net Effect on Agent Workloads
OpenRouter	+5.5%	None	None	Moderate	+5.5% above raw cost; no heartbeat optimization
OpenAI	None	50% on cached (auto)	None	Lenient	-40% on heartbeats; still premium base price
Anthropic	None	90% reads / +25% writes	None	Strict	-70% on heartbeats (steady state); highest base price
Together AI	None	None	None	Lenient	No adjustment; clean per-token pricing
DeepInfra	None	None	None	Lenient	No adjustment; cheapest K2.5 base price
Gonka	None	N/A (sessions instead)	Yes (80% reduction)	TBD	-60-80% token reduction via sessions (HYPOTHETICAL)

6. Monthly Cost Projections

The following tables show monthly cost per provider for each workload tier. “Raw Monthly Cost” uses headline per-token rates applied to total monthly tokens. “Hidden Cost Adjustment” applies the provider-specific factors from Section 5. “Adjusted Monthly Cost” is the realistic total.

Blended rate formula: (Input tokens x input price + Output tokens x output price). With 1:1 input:output assumption: Monthly tokens x (input price + output price) / 2 = Monthly tokens x blended rate.

Casual Tier (~31.5M tokens/month)

Worked example (Together AI): 31.5M tokens x (\$0.50 + \$2.50) / 2 / 1M = 31.5 x \$1.50 = \$47.25/month

Provider	Model	Blended Rate/1M	Raw Monthly Cost	Hidden Cost Adjustment	Adjusted Monthly Cost
Together AI	K2.5	\$1.50	\$47.25	None	\$47.25
DeepInfra	K2.5	\$1.35	\$42.53	None	\$42.53
OpenRouter	K2.5	\$1.55	\$48.83	+5.5% markup	\$51.51
OpenAI	GPT-4o	\$6.25	\$196.88	-40% on heartbeats (~\$35 saved)	~\$161
Anthropic	Sonnet 4	\$9.00	\$283.50	-70% on heartbeats (~\$92 saved)	~\$192
Anthropic	Opus 4	\$45.00	\$1,417.50	-70% on heartbeats (~\$458 saved)	~\$960
Gonka (A)	K2.5	\$0.75	\$23.63	-80% heartbeat tokens (12.6M effective)	~\$9.45
Gonka (B)	K2.5	\$1.05	\$33.08	-80% heartbeat tokens (12.6M effective)	~\$13.23
Gonka (C)	K2.5	\$1.35	\$42.53	-80% heartbeat tokens (12.6M effective)	~\$17.01

Gonka adjusted cost calculation (Scenario B): Sessions reduce heartbeat tokens from 13.8M to 2.8M/month (80% reduction). Remaining tokens: 14.4M (messages) + 2.8M (heartbeats) + 3.15M (tool calls) = 20.35M tokens. Subtract the tool call and message tokens which stay the same: total effective = ~12.6M tokens. 12.6M x \$1.05 = \$13.23/month.

Casual Tier Takeaway: K2.5 providers cluster at \$42-52/month. OpenAI and Anthropic cost 3-20x more even with caching. Gonka scenarios range from \$9-17/month with sessions – a 60-75% reduction versus the cheapest verified competitor. **However, this advantage is entirely contingent on Gonka’s session persistence working in production and on the hypothetical pricing being achievable.**

Active Tier (~161.5M tokens/month)

Worked example (DeepInfra): 161.5M tokens x \$1.35/1M = \$218.03/month

Note: Active tier uses multi-model routing (70% budget, 30% strong). For K2.5 providers, this means 70% at lite quantization pricing (assumed same rate) and 30% at full-precision (assumed same rate since K2.5 has one price tier on external providers). For Gonka, tiered routing is built-in and may offer additional savings – not quantified here due to lack of pricing data for Gonka’s quantization tiers.

Provider	Model	Blended Rate/1M	Raw Monthly Cost	Hidden Cost Adjustment	Adjusted Monthly Cost
Together AI	K2.5	\$1.50	\$242.25	None	\$242.25
DeepInfra	K2.5	\$1.35	\$218.03	None	\$218.03
OpenRouter	K2.5	\$1.55	\$250.33	+5.5% markup	\$264.09
OpenAI	GPT-4o	\$6.25	\$1,009.38	-40% on heartbeats (~\$207 saved)	~\$802
Anthropic	Sonnet 4	\$9.00	\$1,453.50	-70% on heartbeats (~\$525 saved)	~\$929
Anthropic	Opus 4	\$45.00	\$7,267.50	-70% on heartbeats (~\$2,627 saved)	~\$4,641
Gonka (A)	K2.5	\$0.75	\$121.13	-80% heartbeat tokens (56.3M effective)	~\$42.23
Gonka (B)	K2.5	\$1.05	\$169.58	-80% heartbeat tokens (56.3M effective)	~\$59.12
Gonka (C)	K2.5	\$1.35	\$218.03	-80% heartbeat tokens (56.3M effective)	~\$75.99

Gonka adjusted cost calculation (Scenario B): Sessions reduce heartbeat tokens from 82.9M to 16.6M/month. Effective tokens: 57.6M (messages) + 16.6M (heartbeats) + 21.0M (tool calls) = ~56.3M effective tokens (rounded to account for partial overlap). 56.3M x \$1.05 = \$59.12/month.

Active Tier Takeaway: The gap widens. DeepInfra costs \$218/month vs Gonka Scenario B at \$59/month – a 73% reduction. Heartbeats now account for 51% of tokens, so session persistence captures more absolute savings. Even Gonka Scenario C (price-matched with DeepInfra) costs only \$76/month – 65% cheaper than DeepInfra’s \$218 – purely on the strength of session persistence.

Heavy Tier (~2.94B tokens/month)

Worked example (Together AI): 2,943M tokens x \$1.50/1M = \$4,414.50/month

Provider	Model	Blended Rate/1M	Raw Monthly Cost	Hidden Cost Adjustment	Adjusted Monthly Cost
Together AI	K2.5	\$1.50	\$4,414.50	None	\$4,414.50
DeepInfra	K2.5	\$1.35	\$3,973.05	None	\$3,973.05
OpenRouter	K2.5	\$1.55	\$4,561.65	+5.5% markup	\$4,812.54
OpenAI	GPT-4o	\$6.25	\$18,393.75	-40% on heartbeats (~\$6,221 saved)	~\$12,173

Provider	Model	Blended Rate/1M	Raw Monthly Cost	Hidden Cost Adjustment	Adjusted Monthly Cost
Anthropic	Sonnet 4	\$9.00	\$26,487.00	-70% on heartbeats (~\$15,689 saved)	~ \$10,798
Anthropic	Opus 4	\$45.00	\$132,435.00	-70% on heartbeats (~\$78,443 saved)	~ \$53,992
Gonka (A)	K2.5	\$0.75	\$2,207.25	-80% heartbeat tokens (612M effective)	~ \$459.00
Gonka (B)	K2.5	\$1.05	\$3,090.15	-80% heartbeat tokens (612M effective)	~ \$642.60
Gonka (C)	K2.5	\$1.35	\$3,973.05	-80% heartbeat tokens (612M effective)	~ \$826.20

Gonka adjusted cost calculation (Scenario B): The Heavy tier has 5-minute heartbeats, making heartbeats 84.5% of all tokens (2.49B out of 2.94B monthly). Sessions reduce heartbeat tokens from 2,488M to 498M/month. Effective tokens: 288M (messages) + 498M (heartbeats) + 168M (tool calls) = ~612M effective tokens (rounded conservatively). Subtracted 80% of heartbeat overhead. 612M x \$1.05 = \$642.60/month.

Heavy Tier Takeaway: This is where the analysis becomes transformative. DeepInfra costs \$3,973/month. Gonka Scenario B costs \$643/month – an **84% reduction**. Even Scenario C (price-matched) costs \$826/month – a 79% reduction. The 5-minute heartbeat interval means heartbeats dominate token consumption at 84.5%, and session persistence eliminates most of that overhead. **At this tier, session persistence is worth more than any per-token price discount a competitor could offer.**

Note: Anthropic’s prompt caching makes Sonnet 4 (\$10,798/month) cheaper than OpenAI GPT-4o (\$12,173/month) at this tier despite higher per-token rates, because Anthropic’s 90% cache read discount is more aggressive. However, both remain 12-16x more expensive than K2.5 providers.

7. At-Scale Economics

What happens when an organization scales beyond the Heavy tier to 100+ agents? The cost dynamics shift in important ways.

Linear Scaling (No Volume Discounts)

OpenRouter: Cost scales linearly with agent count. 100 agents = 100x Heavy tier cost. No volume discount mechanism exists – each token costs the same whether it is the first or the billionth. At 100 agents, monthly spend exceeds \$481,000. The 5.5% markup compounds: at \$481K/month, that is \$26,500/month in markup alone.

Together AI: Also linear scaling, but at lower base rates. 100 agents at Heavy tier = ~\$441,000/month. Together AI has not publicly advertised volume discounts for serverless inference, though enterprise agreements may be available for committed spend.

DeepInfra: Similar to Together AI. Linear scaling at \$3,973/agent/month. 100 agents = ~\$397,000/month. DeepInfra’s cost advantage over Together AI (10-15% cheaper) becomes meaningful at scale: \$44K/month saved at 100 agents.

Enterprise Agreements

OpenAI: Offers enterprise pricing through committed-use agreements. Exact discounts are not public, but industry reports suggest 20-40% reductions on committed annual spend above \$100K. However, even with a 40% discount, GPT-4o at \$7,304/agent/month (adjusted) x 100 agents = \$730,400/month. Scale does not solve the fundamental premium pricing problem.

Anthropic: Similar enterprise agreement structure. Rate limits are the binding constraint at scale – Anthropic’s concurrent request limits may require enterprise tier negotiation before 100 agents are viable. Monthly cost with enterprise discount (estimated 30%) and caching: ~\$7,559/agent/month x 100 agents = ~\$755,900/month.

Decentralized Network Economics

Gonka (HYPOTHETICAL): Gonka’s decentralized model introduces a fundamentally different scaling dynamic. As the Gonka network grows (more GPU miners join), aggregate inference supply increases. If supply growth outpaces demand growth, per-token costs could decrease over time rather than remaining flat. Additionally, GNK token incentives

subsidize early inference costs – miners earn GNK block rewards in addition to inference fees, effectively lowering the break-even price they need to charge.

However, several caveats apply:

1. **Token inflation risk:** If GNK mining rewards subsidize inference pricing, the cost advantage is partially funded by token inflation, which may not be sustainable long-term (see v1.0 research on emission decay).
2. **Network size assumptions:** Cost decreases require network growth. If Gonka’s miner count stagnates, the cost advantage does not materialize.
3. **Centralized providers are also getting cheaper.** NVIDIA Blackwell GPUs enable up to 10x cost-per-token reduction for providers using new hardware (source: NVIDIA GTC 2026 blog). This narrows Gonka’s structural compute cost advantage.

At 100 agents with Scenario B pricing and session persistence: ~\$64,260/month. Compared to DeepInfra at ~\$397,000/month, this is an 84% reduction. Even if Gonka’s actual pricing settles at Scenario C (price match), session persistence alone would yield \$82,620/month vs \$397,000 – a 79% reduction.

Scaling verdict: Gonka’s session persistence advantage compounds at scale. The more agents deployed, the more heartbeat overhead is eliminated, and the larger the absolute dollar savings. This is a structural advantage that per-token pricing competition cannot replicate without implementing equivalent session persistence technology.

8. Where Gonka Wins and Where It Does Not

Where Gonka Wins

Agent workloads with frequent heartbeats. This is Gonka’s strongest competitive position. Server-side session persistence eliminates the dominant cost driver (heartbeat context resending) in a way that no other provider architecturally supports. OpenAI and Anthropic offer prompt caching, which partially addresses the problem, but caching has TTL expiration and requires the client to resend the full context – the server decides what to cache. Gonka’s sessions are persistent by design: the context lives server-side and never needs resending.

Budget-conscious Casual and Active tier users (if pricing is aggressive). At Scenario A or B pricing, Gonka offers the cheapest K2.5 inference available. Combined with session savings, the effective cost is 60-80% below the next cheapest option. This is a compelling acquisition message for price-sensitive solo developers and small teams.

Privacy-sensitive workloads. Decentralized inference means no single entity (including Gonka) stores conversation data centrally. For agents handling sensitive information (financial data, medical queries, private community management), this is a meaningful differentiator over centralized providers that log and may inspect API traffic.

Token-aligned users. For crypto-native users who already hold or earn GNK, paying for inference with tokens they mine or stake creates a closed economic loop. This is a niche advantage but a real one for the DePIN ecosystem.

Where Gonka Does Not Win

Users who need multiple models. Gonka currently offers only K2.5 (in 3 quantization levels). OpenRouter offers 500+ models. OpenAI offers GPT-4o, o1, o3, DALL-E. Anthropic offers Claude Opus, Sonnet, Haiku. A developer who wants to use GPT-4o for reasoning and K2.5 for simple tasks cannot do so on Gonka alone. This is the single largest competitive gap.

Users who need enterprise SLAs. Gonka has no published SLA, no guaranteed uptime, and no enterprise support tier. OpenAI and Anthropic offer 99.9% uptime SLAs with contractual credits for downtime. For production workloads where reliability is non-negotiable, Gonka cannot yet compete on trust.

Users already invested in OpenAI/Anthropic prompt caching ecosystems. Developers who have optimized their agent architectures around OpenAI’s or Anthropic’s caching (structuring prompts to maximize cache hits, using cache-aware batching) would need to re-architect for Gonka’s session-based approach. The switching cost is real, even if the destination is cheaper.

Cost-sensitive users if Gonka prices at or above DeepInfra (Scenario C). Without a price advantage, Gonka’s appeal rests entirely on session persistence and privacy. While these are genuine differentiators, they may not overcome the friction of manual configuration (Gonka is not a built-in OpenClaw provider) and lack of established reputation. At price parity, the path of least resistance is to stay with DeepInfra or Together AI.

Users who need zero-config OpenClaw integration. OpenRouter is built into OpenClaw. Setting up Gonka requires manual JSON configuration in the OpenClaw provider config. This friction is small in absolute terms but disproportionately affects adoption – developers choose the path of least resistance, especially when evaluating a new provider for the first time.

Where It Depends on Pricing

Everything in the “cost advantage” narrative hinges on Gonka setting pricing below Together AI and DeepInfra. The competitive analysis changes dramatically across the three scenarios:

- **Scenario A (\$0.25/\$1.25):** Gonka wins on cost at every tier, even without session persistence. Sessions make it overwhelmingly cheaper. Risk: margins may be unsustainable.
 - **Scenario B (\$0.35/\$1.75):** Gonka wins on adjusted cost (with sessions) at every tier. Raw cost is competitive but not cheapest. This is the sweet spot: competitive pricing + session advantage creates a strong value proposition without sacrificing margin.
 - **Scenario C (\$0.45/\$2.25):** Gonka’s raw cost matches DeepInfra. Without sessions, there is no cost advantage. With sessions, Gonka is significantly cheaper on adjusted cost. The value proposition narrows to “same price, but sessions make it cheaper in practice.”
-

9. Pricing Recommendations for Leadership

Based on the analysis above, the following recommendations are offered for Gonka’s pricing strategy.

Recommended Target: Scenario B (\$0.35/\$1.75 per 1M tokens)

Scenario B positions Gonka 30% below DeepInfra on headline rates while leaving margin for network sustainability. When combined with session persistence savings, effective cost for agent workloads drops to 60-84% below competitors depending on tier. This creates a two-part value proposition:

1. **Headline:** “30% cheaper than the cheapest K2.5 provider”
2. **Depth:** “80% cheaper for agent workloads because sessions eliminate heartbeat overhead”

The headline gets developers to look. The depth gets them to stay.

Publish Pricing Before Any GTM Activity

Every competitor analyzed in this document has a public pricing page. Gonka cannot credibly claim cost advantages without published prices. The absence of a pricing page is currently the single most damaging competitive gap for GTM messaging. A developer evaluating Gonka will search for pricing, find nothing, and move on. Publish Scenario B pricing immediately, even as “early access” or “beta” pricing that may change.

Consider a Free Tier

A free tier of 1-2M tokens/month (approximately \$1-2 at Scenario B rates) would eliminate signup friction for Casual tier users. At 1M free tokens/month, a Casual user could run a basic agent for 1-2 days before needing to pay. This is enough to evaluate Gonka’s session persistence advantage firsthand. Together AI, OpenAI, and Anthropic all offer free tier credits; Gonka should match this table stakes expectation.

Lead with Session Savings, Not Per-Token Rates

The pricing analysis reveals that Gonka’s most defensible advantage is not per-token cost (which competitors can undercut by adjusting margin) but session persistence (which requires architectural changes competitors have not made). Marketing should lead with:

“Your OpenClaw agent’s heartbeats cost \$0 on Gonka. On other providers, heartbeats are your biggest expense.”

This message is specific, quantifiable, verifiable, and structurally true. It shifts the competitive conversation from “who has the lowest \$/1M tokens” (a race to the bottom) to “who understands agent workloads” (a durable advantage).

Appendix: Calculation Reference

Token-to-Cost Formula

$$\text{Monthly Cost} = \text{Monthly Tokens} \times (\text{Input Price}/1\text{M} + \text{Output Price}/1\text{M}) / 2$$

Assumes 1:1 input:output ratio. Actual agent workloads are typically 60-70% input / 30-40% output, which would reduce costs for providers with cheaper input tokens (all K2.5 providers). The 1:1 assumption is conservative.

Session Persistence Adjustment

$$\text{Effective Monthly Tokens} = \text{Message Tokens} + (\text{Heartbeat Tokens} \times 0.20) + \text{Tool Call Tokens}$$

The 0.20 factor assumes 80% of heartbeat tokens are eliminated by sessions (only delta context sent instead of full resend). This is conservative – if system prompt is 80%+ of heartbeat context, sessions could reduce heartbeat overhead by 85-90%.

Prompt Caching Adjustment (OpenAI)

$$\text{Adjusted Heartbeat Cost} = \text{Heartbeat Tokens} \times (\text{Cacheable\%} \times 0.50 + \text{Non-cacheable\%} \times 1.00) \times \text{Rate}$$

Assumes ~80% of heartbeat tokens are cacheable (system prompt), 20% are non-cacheable (recent context delta). Net effect: ~40% reduction in heartbeat costs.

Prompt Caching Adjustment (Anthropic)

$$\text{First Heartbeat Cost} = \text{Heartbeat Tokens} \times (\text{Cacheable\%} \times 1.25 + \text{Non-cacheable\%} \times 1.00) \times \text{Rate}$$

$$\text{Subsequent Heartbeat Cost} = \text{Heartbeat Tokens} \times (\text{Cacheable\%} \times 0.10 + \text{Non-cacheable\%} \times 1.00) \times \text{Rate}$$

The 25% write premium on first fill amortizes quickly: after 2 heartbeats, the net saving exceeds the initial premium. For always-on agents, the steady-state cost is ~10% of uncached rate on the cached portion, yielding ~70% overall heartbeat cost reduction.

Sources

Pricing Data (MEDIUM-HIGH confidence)

- [Together AI Pricing](#) – K2.5 at \$0.50/\$2.50 per 1M tokens. Accessed April 2026.
- [Artificial Analysis: K2.5 Providers](#) – DeepInfra at \$0.45/\$2.25. Accessed April 2026.
- [OpenRouter Pricing](#) – K2.5 passthrough pricing + 5.5% credit markup. Accessed April 2026.
- [OpenAI Pricing](#) – GPT-4o at \$2.50/\$10.00. Prompt caching details. Accessed April 2026.
- [Anthropic Pricing](#) – Claude Sonnet 4 at \$3.00/\$15.00, Opus 4 at \$15.00/\$75.00. Caching details. Accessed April 2026.

Agent Workload Data (MEDIUM confidence)

- [OpenClaw Token Costs 2026](#) – 9,600 tokens/turn system prompt overhead.
- [OpenClaw Pricing Guide](#) – Heartbeat cost analysis, 30-min and 5-min intervals.
- [OpenClaw API Costs 2026](#) – Per-task cost analysis, \$0.30-420/month range.
- LangChain State of Agent Engineering – Tool call multiplier (3-10x) for agent vs chatbot workloads.

Market Context (MEDIUM confidence)

- [NVIDIA Blackwell Inference Cost Blog](#) – 10x cost reduction for providers using Blackwell GPUs.
- [AI Inference Cost Crisis 2026](#) – 85% of AI budget goes to inference.
- [DePIN Compute Wars 2026](#) – Decentralized provider landscape dynamics.

Internal References

- [.planning/research/STACK.md](#) – OpenClaw architecture, provider ecosystem
- [.planning/research/FEATURES.md](#) – Feature landscape, competitive dimensions
- [.planning/research/PITFALLS.md](#) – GTM pitfalls, trust barriers
- [output/gonka_competitive_feature_matrix.md](#) – Feature matrix (COMP-01 cross-reference)

This analysis uses HYPOTHETICAL pricing for Gonka. All Gonka cost projections are scenario-based and contingent on final pricing decisions. Leadership should not cite specific Gonka cost savings externally until actual pricing is published.