

Gonka Tokenomics: Macro-Economic Research Synthesis

Research Date: February 2026 (last verified against live sources: July 18, 2026) **Version:** 2.5 **Scope:** Novel macro-economic studies and tokenomics research applied to Gonka's decentralized AI compute network **Update Note (v2.5, July 18, 2026):** Fifth verification pass. Key corrections: GENIUS Act stablecoin rule deadline (Jul 18, 2026) was **missed** – eight proposals published, zero finalized, regime now defaulting toward the Jan 18, 2027 statutory backstop; Chutes figures corrected to the tao.media source (~120B tokens/day steady-state with 160B peaks, 34T+ cumulative, 696K+ users, \$5.5M annualized revenue) and its OpenRouter listing corrected to 9 models (DeepSeek not confirmed in current listing); Bitwise's TAO Strategy ETF filing added alongside Grayscale's GTAO conversion (two TAO ETF decisions in the August 2026 window); Bittensor Subnet 108 "Frontier Compute" (agent workflow orchestration, May 2026) added to the watch list; Aethir's own Q1 2026 chain-migration roadmap noted as a structural risk symmetric to Akash's; Nosana June 2026 recap items added; BTC dominance widened to ~57-60% depending on source; CryptoRank added to the GNK aggregator-divergence caveat. GNK price anchor (~\$0.13) re-verified as of Jul 18-19, 2026. **Update Note (v2.4, July 18, 2026):** Fourth verification pass. Key corrections: live Gonka network size revised down again – joingonka.ai's live counter shows ~1,178 GPUs active (consistent with tracker.gonka.vip ~1,214), far below both the April 2026 ~4,648 figure and CoinMarketCap's stale "~5,000" description; SaladCloud is now *live* as Render's third subnet (~60,000 GPUs – the v2.3 withdrawal of that figure was wrong; it belongs to the Salad subnet, distinct from RNP-021's ~1,200 H200-equivalent cap) and Render's Dispersed subnet added an OpenClaw agent-workload recipe; Grayscale's GTAO spot-ETF conversion has an SEC decision expected by August 2026; Kraken listed Bittensor subnet alpha tokens including Chutes (Jun 29, 2026); io.net's 15.96M IO July 11 unlock noted against its 12M year-one burn target; CLARITY Act reframed from "vote targeted week of July 20" to "no vote scheduled, odds near coin-flip"; GENIUS Act stablecoin rule deadline (Jul 18, 2026) added; TrendForce's Jul 8-9 upward revision of Q3/Q4 DRAM forecasts noted; Vera Rubin ramp now slightly delayed (share of 2026 shipments cut ~29% to ~22%; Rubin Ultra reportedly cancelled/scaled back); H100 July cohort median firmed to ~\$3.15/hr; AI-token sector cap refreshed to ~\$22B; OpenRouter demand anchored to the official 25T tokens/week (May 2026) figure. **Update Note (v2.3, July 18, 2026):** Third verification pass. Major corrections: live Gonka network size is far below the ~14,000 H100-equivalent February figure – ecosystem trackers report ~4,648-5,000 H100-equivalents as of mid-2026, which roughly triples per-GPU mining share versus the 14,000-denominator math; Chutes (Bittensor Subnet 64) added as the largest decentralized inference provider and a HIGH-relevance competitor; memory supercycle reframed from "intensifying" to "still rising but decelerating" (TrendForce Jul 2026: Q3 DRAM +13-18% QoQ vs ~60% in Q2); canonical TAO on Solana corrected to May 5, 2026; Vera Rubin first shipments pulled forward to July 2026; CLARITY Act status updated (floor vote targeted week of July 20; three Democratic senators formally opposed); SEC Innovation Exemption caveated (still at OIRA, eligibility limits Gonka may not meet); "AI tokens only profitable sector in Q1 2026" softened to "smallest sector decline"; OpenRouter 100T-token empirical data added; Render RNP-021 capacity corrected (~1,200 H200-equivalents, not ~60,000 GPUs) plus Coinbase listing; Nosana, Gensyn, Akash-RFP, and B200-pricing details refreshed; Tianrong (TIPS) DEPINfer noted as a new entrant. **Update Note (v2.2, July 18, 2026):** Second verification pass against live sources. Corrections: TAO market cap fixed (~\$2B, not ~\$3.97B); Vera Rubin is in full production (June 1, 2026), Q3 launch confirmed; H100 *rental* deflation "stalled" but did not reverse (July 2026 on-demand average down ~4% YoY) – the cost reversal is in hardware acquisition and some hyperscaler list prices, driven by the memory supercycle; io.net's IDE burn is currently emission-funded per independent explorer analysis, not yet the marketed revenue buyback. Additions: Bittensor Robin τ upgrade + TAO ETF flows + canonical TAO on Solana; Gensyn \$AI TGE; BitTorrent BTTInferGrid; Prime Intellect >\$100M annualized revenue; SEC's pending July 2026 token-sale rulemaking; Q2 2026 memory-supercycle trajectory; 2026 capex estimates above Deloitte's \$600B. **Update Note (v2.1, July 2026):** Corrected against July 2026 market data. Key changes since v2.0: GNK now trades (~\$0.13, ~95% below its Jan 2026 ATH of \$2.61); crypto is in a pronounced bear market; Akash and io.net have both shipped buyback/burn mechanisms, ending Gonka's "first mover" window on that axis; B200 shipped in early 2025 (not mid-2026); and H100 rental deflation has stalled due to the memory supercycle. All price-anchored scenarios (\$1.00 GNK, \$0.45/\$0.30 floor triggers, POL ranges) require recalibration. **Update Note (v2.0):** Updated with deep research from 5 parallel investigation agents covering POL strategy, real yield mechanisms, ve-

tokenomics governance, fee transition stress testing, GPU economics, developer growth strategies, and floor price defense

Executive Summary

This document synthesizes cutting-edge research from 10 initial parallel investigations and 5 subsequent deep-dive research agents into macro-economics and tokenomics, applying findings to Gonka's unique position as a decentralized AI compute network. The research draws from academic papers, industry analyses, and empirical data from 2024-2026 to provide comprehensive insights for Gonka's economic design.

Key Findings:

- Gonka's exponential decay emission ($\exp(-0.000475 * \text{epochs})$) represents optimal design based on latest tokenomics research
- Sprint Consensus's AI-productive PoW addresses critical sustainability concerns facing pure hash-based PoW
- The 20% collateral base weight + 80% collateral-weighted system aligns with emerging best practices
- Dynamic pricing mechanisms mirror successful implementations like EIP-1559
- The decentralized AI compute market projects to reach \$100B+ by 2032

New Findings (v2.0 – February 2026 Deep Research, revised July 2026):

- **Protocol-Owned Liquidity:** 20-25M GNK from Community Pool deployed as concentrated liquidity on Uniswap v3 achieves deep liquidity at \$0.50 per \$1 TVL (vs \$10 per \$1 for mercenary LM). Dollar sizing was modeled at \$1.00 GNK; at the July 2026 price (~\$0.13) the same GNK allocation supports roughly one-eighth the USD depth
- **Real Yield:** Enhanced revenue split (20/70/5/5) with continuous TWAP buyback-and-burn creates dual deflationary pressure alongside EIP-1559 base fee burns
- **veGNK Governance:** 1 month - 2 year lock range with linear time-weighting fully mitigates flash loan governance attacks; projected 35-50% lock rate at steady state
- **Fee Transition:** Conservative crossover at Year 4 (\$1 GNK) to Year 8-9 (\$5 GNK); moderate scenario achieves fee dominance by Year 1-3
- **GPU Pricing (revised):** H100 pricing collapsed 64-81% from Q4 2024 to Q1 2026, then deflation stalled – July 2026 H100 median \$2.29-3.12/hr, with on-demand averages roughly flat to slightly down YoY (Thunder Compute: \$3.89 Jul 2025 to \$3.72 Jul 2026). The memory supercycle is instead reversing *hardware acquisition* costs (+30-50% on GPU servers) and some hyperscaler list prices. B200 shipped in early 2025 and is live rental inventory, not a future catalyst. Oracle-based USD pricing still recommended to eliminate dual volatility
- **Competitive Moat (revised):** Buyback/burn is no longer a differentiator – Akash activated Burn-Mint Equilibrium in March 2026 and io.net's Incentive Dynamic Engine went live June 2026. Gonka's remaining moat is direct yield distribution to veGNK lockers, which remains rare
- **Floor Defense (revised):** Programmatic TWAP buybacks triggered at 75% of 30-day TWAP, with \$0.45 absolute floor (25% below the \$0.60 Bitfury Schelling point). As of July 2026, GNK trades ~\$0.13 – both the \$0.45 and \$0.30 triggers are far breached and the mechanism was never deployed; trigger levels need rebasing before implementation

Table of Contents

1. [Bitcoin Halving Economics & Emission Design](#)
2. [PoW vs PoS Economic Models](#)
3. [Token Emission Curve Innovations](#)
4. [Decentralized Compute Network Economics](#)
5. [AI Compute Market Dynamics](#)
6. [DeFi Mechanism Design Innovations](#)
7. [Network Effects & Adoption Economics](#)
8. [Collateral & Slashing Economics](#)
9. [Dynamic Pricing Mechanisms](#)
10. [Decentralized Governance Economics](#)
11. [Gonka-Specific Applications](#)
12. [Protocol-Owned Liquidity Strategy](#)
13. [Real Yield & Revenue Distribution](#)
14. [Fee-to-Emission Transition Stress Test](#)
15. [Competitive Positioning & Developer Growth](#)
16. [Research Sources & Citations](#)

1. Bitcoin Halving Economics & Emission Design

1.1 Stock-to-Flow Model: Current Status (2024-2026)

The Stock-to-Flow (S2F) model, once a dominant Bitcoin valuation framework, has faced significant empirical challenges:

Model Performance:

- Precision broke down after 2021, with Bitcoin trading significantly below S2F predictions
- A 2024 academic paper found S2F predictions help explain Bitcoin returns *in-sample* but have **limited to no ability to predict out-of-sample returns**
- Academic research identified 80.57% Pearson correlation between S2F estimates and the logarithm of time since Bitcoin's genesis block – when time fixed-effects are introduced, “statistically significant” regression results become insignificant

Current Assessment:

“Stock-to-flow works best as a conceptual baseline framework rather than a predictive model – it correctly identifies that Bitcoin's scarcity increases over time and supports long-term value appreciation, but treating S2F predictions as price targets leads to disappointment.”

Implications for Gonka:

- Scarcity models provide directional guidance but not precise valuations
- Gonka's fixed 1B supply with exponential decay creates comparable scarcity dynamics
- Network utility (AI compute) provides fundamental value beyond pure scarcity

1.2 Halving Events & Network Security

2024 Halving Impact (Bitcoin):

- Block reward dropped from 6.25 BTC to 3.125 BTC (April 19-20, 2024)
- Hash price fell ~60% from April 2024 levels while hashrate increased 40%
- Top 10 listed miners: ~\$45,000 per Bitcoin all-in operating costs post-halving

Security Implications:

- Academic research (Sedlmeir et al., SSRN) identifies halving as an “imminent security risk” by destabilizing mining profitability
- Estimated cost of one-hour 51% attack: **\$5-20 billion** as of February 2024
- Hashrate declines reduce attack costs proportionally

Bittensor First Halving (executed ~December 12-15, 2025): (Updated v2.1)

- Supply-triggered at 10.5M TAO mined; block reward cut 1.0 TAO to 0.5 TAO, daily issuance 7,200 to 3,600 TAO
- Subnets without strong demand-side revenue saw significant miner churn
- Post-halving, TAO fell from a Q4 2025 rejection at \$535 to ~\$189-199 (mcap ~\$1.9-2.2B; ~71-75% below the \$757 ATH) by July 18, 2026
- Validates Gonka's dual-income model (emission + inference fees) as critical hedge against halving shocks

Mining Economics Post-Halving:

- Power costs: 75-85% of miners' total operating expenses
- Industry consolidation accelerating toward energy-efficient infrastructure
- Bitcoin mining network: 104% hashrate increase in 2024 (following 90% growth in 2023)

1.3 Exponential Decay vs Step-Function Halvings

Design Approach	Characteristics	Trade-offs
Step-Function (Bitcoin)	Discrete halvings every ~4 years	Predictable, transparent, creates visible supply shocks
Exponential Decay (Gonka)	Continuous reduction proportional to current value	Smoother supply profile, reduces periodic disruptions

Gonka's Emission Formula:

```
current_epoch_reward = 323,000 * exp(-0.000475 * epochs)
```

- Achieves ~50% reduction every ~1,460 epochs (~4 years)
- Provides Bitcoin-like scarcity with smoother distribution curve
- Eliminates “halving shock” market disruptions

Key Emission Milestones (v2.0 – from stress test research):

Year	Epoch Reward (GNK)	% of Initial	Annual Inflation Rate
0	323,000	100%	33.7%
4	152,440	47.2%	11.6%
8	71,929	22.3%	4.9%
12	33,936	10.5%	2.5%
20	7,582	2.3%	1.0%

2. PoW vs PoS Economic Models

2.1 Capital Efficiency Analysis

Proof-of-Work Capital Requirements:

- Substantial upfront CapEx for ASIC procurement and infrastructure
- Significant ongoing OpEx on electricity and maintenance
- Rising difficulty increases capital requirements continuously

Proof-of-Stake Advantages:

- Ethereum validators: ~3.15% APY average risk-adjusted reward (Q2 2025)
- Staking rewards in 2025: up to 7% APR
- Eliminates hardware arms race; validators operate on standard laptops (8 GB RAM)

Capital Immobilization Trade-off:

- PoS requires bonding capital (immobilized), incurring opportunity costs
- However, immobilization discourages frivolous participation and underpins network durability

2.2 Security Budget Comparison

PoW Attack Economics:

- Duke University (2025): Week-long 51% attack on Bitcoin could cost ~\$6 billion
- Four mining pools control ~75% of Bitcoin’s hashrate (centralization risk)

PoS Attack Thresholds:

- 33%+ of bonded stake required to halt finality
- 66%+ required to rewrite finalized checkpoints
- Slashing mechanisms counter “Nothing-at-Stake” problem

Slashing Statistics (Ethereum, Feb 2024):

- Only 414 validators slashed out of ~1,174,000 deposited (<0.04%)
- Correlated slashing penalties scale with simultaneous events (Day 18 “correlation penalty”)

2.3 Decentralization Metrics (2025)

Bitcoin Mining Concentration:

- Four pools (Foundry USA, AntPool, ViaBTC, F2POOL): ~75% of total hashrate
- High infrastructure barriers concentrate mining in areas with cheap electricity

Ethereum Staking Concentration:

- Lido: 29-31% of staked ETH
- 1.04-1.06 million active validators with 30-34 million ETH staked (~28-30% of supply)
- Dynamic analysis shows trend toward less concentration over time

2.4 Environmental Economics

Metric	PoW	PoS
Bitcoin annual consumption	~120 TWh/year	N/A
Ethereum post-Merge reduction	N/A	99.95% vs PoW

Metric	PoW	PoS
Country equivalent (Bitcoin)	Argentina	N/A

2.5 Gonka's Hybrid Approach: Sprint Consensus

Sprint Consensus Innovation: Gonka's Transformer-based PoW addresses critical PoW limitations:

1. **Productive Computation:** Unlike hash-based PoW, Sprint's Transformer operations contribute to AI workloads
2. **Verification Efficiency:** 1-10% randomized task verification (vs 100% in other networks)
3. **Resource Allocation:** ~100% of computational resources allocated to AI workloads
4. **Synchronized Racing:** ~5 minutes per day of network-wide Sprint computation

Comparison to Pure PoW:

- Eliminates "wasted" computational work
- Maintains PoW's security properties through economic incentives
- Aligns mining rewards with network utility generation

3. Token Emission Curve Innovations

3.1 Emission Model Comparison

Exponential Decay:

- Aggressive early growth that tapers over time
- Sustainable long-term trajectory
- Example: Bittensor (TAO) halving mid-December 2025

Linear Models:

- Consistent token release over defined period
- Example: Curve Finance (274M to 137M tokens/year, 2020-2024)

Step-Function (Halving):

- Periodic discrete reductions
- Bitcoin April 2024: 6.25 to 3.125 BTC per block
- Next halving: April 2028

Hybrid Approaches:

- Multi-phase designs increasingly common
- Linear distribution for aggressive growth, then exponential decay for sustainability, then flat line at cap

3.2 Inflation Targeting in Crypto

Algorand (2025):

- Transitioned from governance rewards to staking rewards
- Vesting schedule for early backers concluded in 2024
- Current annual inflation: ~6.37%

Cosmos (ATOM) Debate:

- 7-20% annual inflation range
- Community debates cutting to Ethereum-level 2-4%
- Failed 2022 governance vote (63% rejected reduction)

3.3 The Tail Emissions Debate

Monero's Model (Since mid-2022):

- Continuous "tail emission" of 0.6 XMR per block (2-minute target)
- Modest inflation – over a century to double supply
- Ensures miners not 100% reliant on transaction fees

Economic Rationale:

- "Minimum subsidy" keeping fees low
- Provides lower bound of network security
- Enables dynamic block sizes
- Tail emission inflation likely equals rate of lost coins

Contrast with Fixed Supply (Bitcoin):

- Bitcoin's transition to fee-dominated economics (post-2040) remains incompletely studied

- Fundamental question: Will transaction demand compensate for diminishing block rewards?

Gonka Relevance (v2.0): Gonka's fee transition stress test (Section 14) models this exact question – at conservative 10% developer growth, fee dominance occurs by Year 4-10 depending on GNK price. Gonka's contingency plan includes governance-activatable tail emissions if fee growth falls below threshold.

3.4 Vesting Schedule Best Practices (2024-2025)

Industry Standards:

- 58% of term sheets include founder vesting (2024)
- Most common: 4-year vesting with 1-year cliffs
- Trend: Longer vesting, shorter cliffs

Price Stability Correlation:

- Projects with >70% tokens vested experience lower volatility
- Extended vesting reduces market manipulation risks

Gonka Application:

- 200M GNK to founders with vesting schedule
- Structured distribution prevents early dumping
- 120M community allocation supports ecosystem growth

4. Decentralized Compute Network Economics

4.1 Market Overview

Network	Latest Revenue (as of Jul 2026)	Key Mechanism	Cost Advantage
Aethir	\$127.8M FY2025 revenue; end-2025 ARR \$147-166M (no 2026 figures published yet)	Enterprise compute contracts (150+ clients)	Enterprise pricing, real cash flow
Render	1M+ cumulative RENDER burned by Dec 2025	Burn-and-Mint Equilibrium	Dynamic based on demand
Akash	\$253,250 Q1 2026 lease revenue (Messari); ~\$5M self-reported gross compute spend same quarter	Reverse Auction + BME (Mar 2026)	90% lower than AWS
io.net	~\$12.5M self-reported annualized revenue mid-2026 (unaudited)	Aggregation + IDE (Jun 2026)	70-90% cheaper than AWS
Bittensor	~\$43M Q1 2026 revenue from AI customers	Subnet tokenization, flow-based emissions	Varies by subnet
Chutes (Bittensor SN64)	\$5.5M annualized revenue (75% organic, 25% sponsored; tao.media)	Alpha-token emissions (~\$54K/day TAO pre-halving basis, ~\$27K post) + paid inference	Decentralized inference at commodity token prices

Note: Akash lease revenue (Messari, on-chain leases) and Akash's self-reported gross compute spend are different metrics; do not conflate them. Chutes is broken out separately because it is the single most direct decentralized-inference competitor (see §4.5).

Updated Competitive Data (v2.1 – July 2026):

Network	Real Yield to Token Holders	Buyback Mechanism	Revenue Distribution
Akash	Take Pool to AKT stakers	Live: Burn-Mint Equilibrium (Mainnet 17, Mar 23, 2026) – on-chain spend triggers market buy-and-burn of AKT	Community Pool (governance)
Render	None (burn-mint, no distribution)	Burn-and-Mint	No direct token holder yield
Bittensor	None (emission-based rewards; ~\$43M Q1 2026 customer revenue)	None	50% validators / 50% miners

Network	Real Yield to Token Holders	Buyback Mechanism	Revenue Distribution
io.net	None direct	Live in name: IDE (Jun 11, 2026) – marketed as $\geq 50\%$ of surplus revenue buying back and burning IO ($\geq 12\text{M}$ IO year-one target); independent explorer analysis (Jul 2026) shows the burn is currently emission-funded, not revenue-funded	Emissions tied to real network earnings (demand-driven emissions not yet switched on)
Gonka (proposed)	5% to veGNK + Fund surplus	5% continuous TWAP burn (proposed)	20/70/5/5 split

Key Insight (revised July 2026): Buyback-and-burn is now table stakes on paper – Akash (BME) and io.net (IDE) both shipped in H1 2026 – but io.net’s IDE burn is so far emission-funded rather than the marketed revenue buyback. What remains rare among decentralized AI compute networks is **direct fee-revenue distribution to token lockers**, funded by real revenue. Gonka’s moat claim should rest on the veGNK yield component (and on revenue actually funding it), not the burn component.

4.2 Render Network: Burn-and-Mint Equilibrium (BME)

Mechanism:

- All jobs priced in USD
- Creators burn RENDER tokens equivalent to job value
- Node operators receive newly minted RENDER from capped, declining schedule

Performance:

- Burns Jan-Sep 2025: 530,171 RENDER (+278.9% vs same 2024 period); cumulative BME burns passed 1M RENDER by December 2025
- Q3 2025: 37% burn efficiency increase QoQ
- RENDER at \$1.51 (Jul 16, 2026), ~89% below its \$13.53 ATH; market cap ~\$784M

Growth Phase Reality:

- Monthly emissions still outpace burns
- Requires accelerating usage to achieve net deflation

AI Compute Expansion (v2.4 – July 2026): Render is no longer a rendering-only network. The RCN framework and Dispersed subnet support AI/ML inference and general GPU workloads – and per the Render Foundation’s June 2026 report, Dispersed added an **OpenClaw recipe**, i.e., Render now explicitly supports OpenClaw agent workloads. RNP-021 onboards enterprise GPUs (NVIDIA H100/H200/A100, AMD MI300, Intel Max) for AI training and inference, capped at ~1,200 H200-equivalents per the RNP-021 text. Separately, **SaladCloud is now live as Render’s third subnet** (integration milestones 1 and 2 shipped per the June 2026 Foundation report and Messari’s July 7, 2026 “Understanding Dispersed”): customers fund SaladCloud with RENDER and GPU providers redeem earnings in RENDER to Solana wallets, adding approximately **60,000 consumer GPUs** aimed at agentic-AI compute shortages – one of the largest single expansions in Render’s history. (v2.3 wrongly withdrew the ~60,000 figure as conflicting with RNP-021; both numbers are true of *different* subnets – ~60,000 for Salad, ~1,200 H200-equivalents for RNP-021’s enterprise tier.) Mid-July 2026 developments deepen the AI pivot: Coinbase added RENDER support on July 10, 2026, and RENDER became an accepted payment inside OTOY Studio’s AI creative suite (July 14, 2026) for generation across 30+ AI models including Seedance 2.0 and Kling; the Salad subnet and MCP integration debuted at RenderCon 2026. Render now overlaps directly with Gonka’s target market – its threat level should no longer be scored LOW.

4.3 Akash Network: Reverse Auction Model

Core Mechanism:

- Users specify maximum price for compute resources
- Providers bid downward; lowest bidder wins
- Achieves 90% cost reduction vs AWS

Metrics (updated to Q1 2026, per Messari):

- Lease revenue: \$253,250 (-45% QoQ) – the earlier upward extrapolation did not hold
- New leases: 43,540 (+27.1% QoQ, third consecutive quarterly recovery), but active leases and provider count compressed

- Akash separately self-reports ~\$5M gross compute spend for the same quarter (an all-time high; a broader metric than Messari lease revenue – cite them separately)
- GPU utilization ~33.7%

Lease Revenue History:

Quarter	Lease Revenue (Messari)	QoQ Growth
Q1 2024	\$320K	-
Q2 2024	\$485K	+51.6%
Q3 2024	\$851K	+75.5%
Q1 2026	\$253K	-45%

Token Economics:

- 20% of lease fees flow to Take Pool
- Distributed to AKT holders based on “Stake Weight”
- Comparable to Uber (23%) and Apple (30%) take rates
- **BME (Mar 23, 2026):** Burn-Mint Equilibrium activated via Mainnet 17 – all on-chain compute spend triggers a market buy-and-burn of AKT (minting ACT stable credit for settlement), cutting effective inflation to ~7.1%; Akash’s first deflationary mechanism

Chain Migration Risk (v2.3 – July 2026): Akash announced plans to deprecate its sovereign Cosmos L1 and migrate to a shared-security model. As of July 2026 the migration remains in the evaluation phase: the open RFP process (reported to formally launch mid-October 2026) has 15+ blockchain foundations pitching, with no governance vote, no confirmed destination chain, and no fixed migration date; CEO Greg Osuri publicly calls Solana a “strong contender.” This is a major structural change and execution risk for Gonka’s closest decentralized competitor.

4.4 io.net: Aggregation Economics

Network Scale (caveat – registered vs verified):

- Registered GPUs: 327,000+ claimed (March 2025); io.net’s own docs now say “more than 100,000 devices”
- Daily-verified GPUs: only ~6,720 (~2% of the headline figure), a hangover from the 2024 Sybil attack that inflated registrations; io.net’s 2025 year-in-review cited 2,752 verified GPUs across 138+ countries
- Use verified figures, not the 300K+ headline, in any competitive benchmark

Revenue Metrics:

- Annualized TNE: \$18.4 million (November 2024)
- Self-reported annualized revenue ~\$12.5M by mid-2026 (no independent audit); cited an \$8M enterprise contract as baseline demand
- Critics note closed-source core and unaudited self-reported revenue

Token Innovation – IDE (launched June 11, 2026):

- Incentive Dynamic Engine live (announced alongside an \$8M enterprise deal, ~\$650K/month, and 4B daily AI tokens): emissions/burns tied to real network earnings via a “sustainability ratio”
- Suppliers paid to a stable dollar target; only as much IO released as needed at live price
- Marketed design: $\geq 50\%$ of surplus revenue buys back and burns IO; targets $\geq 12\text{M}$ IO removed in year one and retirement of $\geq 115\text{M}$ of the 231M non-emitted reward pool (a reward-pool reduction, not a 50% circulating-supply cut)
- **Implementation caveat (Jul 2026):** independent analysis of io.net’s own explorer shows the burn is currently funded from emissions, not revenue, with demand-driven emissions still switched off – the revenue-buyback loop is not yet operating as marketed. This strengthens Gonka’s “real revenue to lockers” differentiation argument
- **Offsetting supply event (v2.4):** io.net had a 15.96M IO token unlock on July 11, 2026 – a *single monthly unlock* larger than the IDE’s entire $\geq 12\text{M}$ first-year burn target, landing while the \$8M contract contributes only ~\$650K/month in on-chain earnings. Net supply pressure remains strongly positive; cite this alongside the emission-funded-burn caveat

4.5 Bittensor: Flow-Based Subnet Economics

Dynamic TAO (February 2025):

- Each subnet has own Alpha token tradeable against TAO
- Subnets became directly investible
- Expanded to 129 active subnets

Flow-Based Emissions (November 2025):

- Emissions determined by net TAO inflows (staking minus unstaking)
- Uses EMA with ~86.8-day window and 30-day half-life
- Starves “zombie subnets” with no activity

First Halving (executed ~December 12-15, 2025):

- Reduction: 1 TAO/block to 0.5 TAO/block; daily issuance 7,200 to 3,600 TAO
- Post-halving: subnets without strong usage-based revenue experienced significant miner churn (see §1.2 for price impact)
- Demand side improving: ~\$43M Q1 2026 revenue from AI customers – Bittensor is no longer purely “emission-only, no demand revenue”

Post-Halving Developments (v2.2 – May-July 2026):

- **Robin τ upgrade (May 2026):** Doubled subnet capacity from 128 to 256 slots; each new registration locks ~700 TAO (potentially ~\$32M of TAO if slots fill). 120+ subnets active – a structural TAO demand sink relevant to any emission/demand analysis
- **Institutional flows:** Spot TAO ETF filings and reported ~\$620M institutional bets became a major TAO narrative from May 2026, even as price sat ~71-75% below ATH. **Decisions imminent (v2.5):** there are now **two** TAO ETF decisions in the August 2026 window, not one: the Grayscale Bittensor Trust (GTAO) spot-ETF conversion (S-1 amendments filed through May 2026; Coinbase Custody as anticipated custodian; still pending as of July 18, 2026) and Bitwise’s N-1A filing for a TAO Strategy ETF (part of an 11-fund crypto-strategy batch filed Dec 30, tracked for the same window) – binary approval/denial catalysts for Gonka’s largest decentralized competitor’s token within weeks of this document’s date
- **Canonical TAO on Solana (May 5, 2026):** Launched via Wormhole Labs’ Sunrise gateway (announced at Solana Accelerate USA), enabling native trading on Jupiter/Meteora; it coincided with Grayscale opening its Bittensor Trust for private placement. Notable because Akash is simultaneously evaluating Solana as its migration target (§4.3) – Gonka’s two largest decentralized competitors are converging on Solana liquidity and distribution
- **Chutes (Subnet 64, Rayon Labs) – HIGH-relevance competitor (v2.3, figures corrected v2.5):** By mid-2026 Chutes is the largest decentralized inference provider: ~120B tokens processed/day steady-state post-monetization (peaks of 160B), 34T+ cumulative tokens, 696K+ users excluding OpenRouter traffic, and \$5.5M annualized revenue (75% organic) per tao.media (Feb 2026). It is independently visible as a provider on OpenRouter serving 9 models as of July 18, 2026 (Z.ai GLM, Qwen, MoonshotAI Kimi, Google Gemma, MiniMax; DeepSeek not confirmed in the current listing). Any developer on OpenRouter can already route to decentralized inference via Chutes with zero configuration – this directly undercuts any “only decentralized provider a developer can actually use” positioning for Gonka and must appear in every competitive landscape view. **Exchange liquidity (v2.4):** on June 29, 2026 Kraken listed seven Bittensor subnet alpha tokens including Chutes (SN64) and Targon Compute (SN4) – the first major-exchange access to subnet alphas, with leading alphas at tens-to-\$100M+ market caps – adding independent token liquidity and visibility on top of Chutes’ usage lead
- **Subnet 108 “Frontier Compute” (v2.5 – watch list):** Launched early May 2026 into the post-Robin- τ slot expansion – decentralized workflow orchestration that breaks complex prompts into coordinated multi-step processes (data retrieval, training, synthesis). This is agent-orchestration territory: a first crack in the “no decentralized provider has agent features” framing. Early stage; watch-list only

4.6 New Entrants & Expansions (v2.1 – July 2026)

Networks absent from the v2.0 landscape that now matter:

- **Aethir:** The highest-earning token-incentivized protocol in decentralized compute – FY2025 revenue \$127.8M (Q1 \$28.5M, Q2 \$32.7M, Q3 \$39.9M), end-2025 ARR \$147-166M depending on source, from 150+ active enterprise clients and 435K+ GPU containers in 93 countries across AI training, Web3 infra, and cloud gaming. Real contracted cash flow, not token farming. (No Q1/Q2 2026 figures published as of July 2026 – date the ARR to end-2025. Caveat: Aethir has itself also cited a lower \$126.0M ARR figure, and analysts flag the inconsistency across its self-reported numbers – treat the \$147-166M range as an upper band.) Any DePIN revenue comparison that includes Akash (\$253K quarterly lease revenue) but omits Aethir (\$30-40M quarterly) is materially incomplete. **Structural risk (v2.5):** Aethir’s own 12-month strategic roadmap scheduled a chain migration in Q1 2026, plus Aethir v2 and an EigenLayer ATH Vault upgrade – the same class of execution risk this document flags for Akash (§4.3) applies to Aethir too; migration completion was not verified as of July 2026.
- **Prime Intellect:** Raised a \$130M Series A at a \$1B valuation on July 8, 2026 (Radical Ventures lead; NVIDIA Ventures, Intel Capital, Dell Technologies Capital, Iconiq; total funding >\$150M), positioning as the “Open Superintelligence Stack” for enterprise agent training; released INTELLECT-3, a 106B MoE trained on 512 H200s. Commercially the most important datapoint: it reports **>\$100M annualized revenue and ~6,000**

customers (AI startups, neolabs, enterprises) – out-earning every token-incentivized compute network except Aethir, and rivaling Aethir's run-rate *without a token*.

- **Gensyn:** Launched mainnet April 22, 2026 with >5,000 H100-equivalent hashrate on day one, plus flagship app Delphi (a prediction-market app that burns 70% of protocol revenue). The \$AI token launched via TGE April 29, 2026 with listings on Binance Alpha, Coinbase, KuCoin, and Bitget – one of 2026's most volatile debuts. As of mid-July 2026, \$AI trades ~\$0.023 (~\$32M market cap, rank ~#579), roughly –78% from its \$0.1035 ATH; it has since gained an Upbit listing (KRW/BTC/USDT pairs) and a canonical Uniswap V3 deployment on the Gensyn L2 supporting an automated buyback mechanism. Mainnet and Delphi remain live.
- **BitTorrent BTTInferGrid:** Launched June 17, 2026 – a decentralized AI-inference DePIN network, in bootstrapping phase through 2026 with enterprise/developer APIs planned. A direct decentralized-inference entrant in Gonka's segment.
- **Nosana:** Continues operating its Solana GPU inference marketplace – crossed 3 million completed jobs and 3 million job-hours by January 2026, running ~2,000 active nodes (no shutdown or pivot since the 2023 CI/CD-to-AI pivot). On February 5, 2026 it announced integration of its GPU supply into OpenGPU's network, relaunched its site with \$50 instant free GPU credits, added Sallar and Alio partnerships and a Learning Hub in early 2026, and plans enterprise features (fiat ramps, billing) in H2 2026. July 2026 grants to Voight (Jul 8 – observability/identity/deployment infrastructure for production AI agents on Solana) and AnveVoice (Jul 16 – voice/agent interaction layer) show Nosana funding agent-infrastructure tooling – mild overlap with Gonka's agent-native positioning. Its June 2026 recap (published Jul 1) adds a live Decentralize AI Hackathon, NVIDIA Cosmos 3 Nano support, crypto payments, and a Singapore event with 200+ builders creating AI agents. NOS trades ~\$0.27 with supply 100% unlocked.
- **Tianrong (OTC: TIPS):** A minor but new US-listed entrant – Tianrong Internet Products entered decentralized AI inference with its DEPINfer GPU-aggregation ecosystem and launched DEPIN Studios on July 6, 2026. Early stage; watch-list only.

5. AI Compute Market Dynamics

5.1 Market Size Projections (2024-2030)

Source	2024 Value	2030 Projection	CAGR
Markets and Markets	\$135.81B	\$394.46B	19.4%
Grand View Research	\$35.42B	\$223.45B	30.4%
TechSci Research	\$132.52B	\$371.37B	18.74%
Mordor Intelligence	\$87.60B (2025)	\$197.64B	17.71%

GPU Cloud Market (v2.0):

- Global GPU cloud market: \$3.34B (2023), projected \$33.91B by 2032 (CAGR 29.4%)
- Decentralized GPU compute: fastest-growing segment within this market

Market Composition (2024):

- Hardware: 72.1% revenue share
- GPUs: 67.4% of AI infrastructure market
- North America: 47.7% global share
- Asia-Pacific: Fastest growth at 19.1% CAGR

5.2 GPU Supply/Demand Economics

H100 Price History (v2.1 – Updated with July 2026 data):

Period	H100 Price Range (\$/hr)	Context
Q1 2024	\$8.00-10.00	Supply-constrained, extreme demand
Q3 2024	\$4.00-6.00	Hyperscaler fleet expansion
Q1 2025	\$2.50-4.00	Decentralized marketplace competition
Q1 2026	\$1.50-2.99	Trough of the deflation cycle
Jul 2026	Median \$2.29-3.12 (range ~\$1.40 to \$7-8 hyperscaler on-demand)	Deflation stalled; mild firming at the margin

Aggregate H100 Decline: 64-81% from Q4 2024 to Q1 2026 – but the decline stopped there. Important scoping (v2.2, refined v2.4): the *rental market median did not reverse* – Thunder Compute's July 2026 tracking shows on-demand H100 averages down ~4% YoY (Jul 2025 \$3.89 to Jul 2026 \$3.72/hr) – but the most recent cohort data

(AIMultiple GPU index, July 2026) puts the H100 cohort median at ~\$3.15/hr, at or slightly above the top of the \$2.29-3.12 band: **mild firming, not pure stabilization**. Range across providers: \$1.40 (Thunder Compute) to \$8+ (AWS/GCP); IntuitionLabs tracks \$1.49-6.98 across 15+ providers. H200 cohort median ~\$4.11/hr (\$2.30 FluidStack to \$13.78 Azure). What IS rising sharply is GPU server/hardware acquisition cost (+30-50% from the memory supercycle) and some hyperscaler list prices. Budget decentralized/specialized providers hold \$1.99-3.50/hr.

Key Drivers of the 2024-2025 decline:

1. NVIDIA shipped 3.5M+ H100 units by end 2025
2. Decentralized marketplaces (Vast.ai, RunPod, Akash) brought idle capacity to market
3. Blackwell (B200) volume shipments from early 2025 depreciated Hopper
4. DeepSeek effect: frontier inference on fewer GPUs
5. Quantization advances (GPTQ, AWQ, GGUF) reduced H100 hours per million tokens by 40-60%

Blackwell Status (corrected): B200 shipped at the start of 2025, not mid/Q3 2026 as v2.0 projected. In 2026 Blackwell accounts for >70% of NVIDIA high-end GPU shipments, led by GB300/B300 (Blackwell Ultra); B200 is already the mid-generation part serving cost-sensitive buyers, renting at roughly \$3.20-18.53/hr on-demand as of July 2026 (median ~\$6.25; observed floor Runcrate \$3.20, Spheron \$3.70, Packet.ai \$3.75, Lambda \$4.99, RunPod \$5.89; GCP \$16.11-18.53 at the top; spot from ~\$2.70 – the earlier \$2.69 on-demand low was a stale RunPod figure). Gonka itself supports B200 and recommends it for optimal mining rewards. The 2026 roadmap story is Vera Rubin: in full production since June 1, 2026, with first shipments in July 2026 and mass shipments this summer to eight cloud partners (AWS, Azure, Google Cloud, Oracle, CoreWeave, Lambda, Nebius, Nscale). **But the ramp is slipping (v2.4):** KeyBanc (mid-July 2026) reports the Rubin ramp is slightly delayed by thermal heat-lid issues and SK Hynix HBM4 qualification – initial July volume is below plan, and Rubin's share of NVIDIA's 2026 GPU shipments was cut from ~29% to ~22% (~1.7-1.8M units). The Rubin Ultra four-die GPU was reportedly cancelled/scaled back due to packaging limits (TechTimes, Jul 1, 2026), and the next-gen Kyber rack system is pushed to 2028. Net effect: less 2026 Rubin supply than guided, supporting continued firmness in H100/H200/B200 rental prices.

B200 Specifications and Pricing:

Specification	H100 SXM	H200 SXM	B200 (shipping since early 2025)
HBM Capacity	80 GB	141 GB	192 GB HBM3e
Memory Bandwidth	3.35 TB/s	4.8 TB/s	8.0 TB/s
FP8 Performance	3,958 TFLOPS	3,958 TFLOPS	9,000 TFLOPS
Inference Perf/Watt	Baseline	~1.5x	~2.0x

GPU Price Trajectory (H100-equivalent \$/hr) – revised:

The v2.0 projection of continuous 30-50% annual deflation (\$0.75-1.50 in 2027, \$0.30-0.80 by 2028) is invalidated by the memory supercycle. Current view:

Year	Low	Mid	High	Note
2024	\$4.00	\$6.00	\$10.00	Actual
Jul 2026	\$1.40	\$2.29-3.12 (median)	\$7-8	Actual; deflation stalled
2027+	–	–	–	Depends on memory supply recovery; prior deflation projections withdrawn

Any host-breakeven math that assumed H100 at \$1.50/hr by 2028 (e.g., “breakeven falls to \$0.62”) must be rebased.

Memory-Driven Price Dynamics (v2.3 – still rising, but decelerating):

- Blended DRAM contract prices: ~+80% QoQ in Q1 2026 (some YoY figures >100%); memory prices rose ~246% in 2025; Q2 2026 saw conventional DRAM contract prices up ~60% QoQ
- **Deceleration (TrendForce, July 3, 2026):** Q3 2026 conventional DRAM contract prices forecast up only 13-18% QoQ (vs ~60% in Q2), HBM blended +8-13% (vs 53-58% prior quarter), NAND +10-15% – consumer demand has hit an affordability wall. The correct framing is “still rising but decelerating,” not “intensifying”
- **Deceleration already revised upward (v2.4):** on July 8-9, 2026 TrendForce *raised* its Q3 and Q4 DRAM contract-price forecasts – PC DRAM Q3 now +15-20% QoQ (up from 8-13%), server DRAM confirmed +13-18% with US CSP long-term agreements capping increases – and module maker ADATA reportedly sees Q3 DRAM up 20-30% and NAND up 35-40%. “Decelerating” remains directionally right for contract prices, but the 13-18% floor is now the conservative end, not the consensus; GPU-hardware-cost relief in 2027 looks less likely than the July 3 snapshot implied

- Server DRAM remains in shortage through Q3 2026 on AI-server demand; GPU-based server prices still +30-50%; NVIDIA reportedly cut RTX 50-series output 30-40% in H1 2026 on GDDR7 shortage
- HBM sold out through 2026; HBM3E price hikes ~20% for 2026; HBM takes ~23% of DRAM wafer capacity; HBM market \$35B (2025) toward \$100B (2028)
- This is the primary reason GPU rental deflation stalled; hardware acquisition costs remain elevated even as the rate of increase moderates

5.3 Inference vs Training Market Segmentation

Market Scale:

- AI Inference Market: \$106.15B (2025) to \$254.98B (2030), CAGR 19.2%
- By 2030: Inference market ~10x size of training market

Cost Structure:

- Training: One-time/occasional heavy cost
- Inference: Continuous, scales with user adoption
- Inference accounts for **80-90% of total compute dollars** over model lifecycle

2026 Confirmation (v2.2): Deloitte's late-2025 prediction puts inference at roughly two-thirds of all AI compute in 2026 (vs half in 2025), with inference-optimized chips exceeding \$50B and hyperscaler capex ~\$600B (+36% YoY, ~75% AI). Mid-2026 analyst estimates have moved past that: 2026 AI-infrastructure capex estimates now run \$690B (Futurum) to \$725-770B (NextMSC and others) – Deloitte's \$600B is the conservative floor, not the consensus, which strengthens the inference-demand thesis. Goldman Sachs projects token consumption multiplying ~24x to 120 quadrillion tokens/month between 2026-2030 on agent adoption; agentic workloads consume 5-30x more tokens per task, and average enterprise AI budgets grew from \$1.2M (2024) to \$7M (2026). Token prices fell ~280x over two years while total enterprise AI spend rose ~320% – volume-driven growth, the core Gonka demand thesis.

Live Empirical Data (v2.3 – OpenRouter/a16z “State of AI,” May 2026): The strongest inference-demand evidence is now empirical, not forecast. OpenRouter's official Series B release (Businesswire, May 26, 2026) puts weekly volume at **25T tokens** – up from ~5T only six months prior, i.e., faster than a 4x YoY framing – with ~100T tokens/month; a ~28T/week mid-2026 figure circulates but is unofficial, so cite 25T/week (May 2026) as the hard number (~1% of global inference). Agentic workloads generate more than half of all output tokens, surpassing human usage; 67% of enterprises consume >1B tokens/month (Deloitte 2026); and OpenRouter raised \$113M (CapitalG-led Series B, May 2026). Directly relevant to Gonka's Kimi/Qwen model strategy: Chinese open-source models capture ~61% of global OpenRouter token usage, and China processes ~140T tokens/day. These figures should supplement or replace the forecast-only Goldman/Deloitte citations in any external material.

Infrastructure Differences:

Aspect	Training	Inference
Hardware	Raw compute power, node count	Specialized, efficiency-optimized
Storage	High-capacity SSDs/NVMe	Varies by deployment
Location	Remote, power-rich markets	Metro-adjacent for latency
Scaling	Vertical	Horizontal, edge-capable

5.4 Moore's Law & AI Compute Demand

AI vs Moore's Law Mismatch:

- AI compute needs growing >2x faster than Moore's Law
- US alone: ~100 gigawatts new demand by 2030
- Industry transitioning to multidimensional scaling

GPU Price-Performance:

- Doubles approximately every 2.5 years
- 2025 GPU prices: ~26% of 2019 levels
- Annual deflation rate: –35% to –55% (\$/TFLOPS), –45% to –60% (inference \$/M tokens)

Capital Requirements:

- Meeting AI demand: ~\$500B annual data center investment
- \$7 trillion race to scale infrastructure

5.5 Enterprise vs Consumer AI Compute

Workload Distribution (confirmed for 2026):

- Inference workloads: ~67% of all compute (up from 33% in 2023; ~50% in 2025)
- Almost all AI computing in giant data centers or high-end enterprise servers

Market Size:

- Enterprise AI: \$13.8B (2024) to \$150-170B (2030)
- Consumer AI: \$92.24B (2024) to \$674.49B (2030)

Gonka Positioning:

- Targets enterprise/professional AI compute workloads
- Inference-focused infrastructure aligns with market direction
- Decentralized approach offers cost advantages

6. DeFi Mechanism Design Innovations

6.1 VE-Tokenomics (Vote-Escrowed) Models

Core Mechanism:

- Lock governance tokens for extended periods
- Receive non-transferable “ve” tokens proportional to amount and duration
- Time-weighted voting power governs protocol emissions

ve(3,3) Innovation:

- Combines Curve’s vote-escrow with Olympus DAO’s (3,3) cooperative game theory
- Replaces unsustainable liquidity mining with fee-aligned incentives
- Reduces “mercenary capital” through lock commitments

veGNK Design (v2.0 – from deep research):

Parameter	Proposed Value	Rationale
Min lock	1 month	Filters transient holders
Max lock	2 years	Aligns with AI infrastructure cycles
Voting power	Linear decay (Curve model)	Proven, no gaming
Max boost	2.5x on AI Training Fund yield	Industry standard
Early exit	Not permitted	Maximum commitment
Transferability	Non-transferable	Prevents governance concentration

Lock Rate Benchmarks (2026 data):

Protocol	Lock Rate	Avg Duration	Max Lock
Curve (veCRV)	45%	2.3 years	4 years
Velodrome (veVELO)	52%	1.8 years	4 years
Frax (veFXS)	42%	2.1 years	4 years
Balancer (veBAL)	35%	28 weeks	1 year
PancakeSwap (veCAKE)	25%	18 weeks	1 year
Average	40%	1.7 years	2.7 years

Gonka Projection: 35-50% lock rate at steady state (18+ months post-launch).

Effectiveness:

- Curve’s gauge voting: 2.5x boost for sufficient CRV lockers
- veNFTs (2024-2025): Tradeable locked positions – research shows they increase governance centralization; NOT recommended

Challenges:

- Governance concentration: Few large lockers can dominate
- Scalability: Epoch-based voting (10-day minimum) slows adaptation

6.2 Real Yield vs Emission-Based Yield

2024-2025 Shift:

- **77% of DeFi yields** came from real fee revenue (over \$6B in 2024)
- Fundamental transition from emission-dependent to revenue-sharing

Real Yield Protocol Comparison (v2.0):

Protocol	Revenue Source	Annual Revenue	Staker Distribution	Buyback
GMX	Trading fees	~\$150M	30% in ETH/AVAX	None
Aave	Lending spread	~\$250M	~10% + buyback	\$52M/yr
Hyperliquid	Trading fees	~\$1.2B	0%	97% burn
MakerDAO	Stability fees	~\$200M	After \$50M buffer	MKR burn
Synthetix	Exchange fees	~\$40M	100% to SNX stakers	None
Curve	Swap fees	~\$25M	50% to veCRV	None

Real Yield Advantages:

- Market-cycle independence
- Attracts long-term aligned capital
- Protocol stability through volatility

Emission Limitations:

- Reliance on new user acquisition
- Inflationary pressure on token value
- Creates farm-and-dump dynamics

6.3 Protocol-Owned Liquidity (POL)**Problem Solved:**

- DeFi 1.0 “mercenary capital” fled when incentives expired
- POL: Protocol itself owns and manages liquidity positions

Mercenary Liquidity Quantification (v2.0):

Protocol	Peak Incentivized TVL	TVL After Emissions End	Retention Rate
SushiSwap (2020-2021)	\$1.4B	\$180M (6 months later)	12.8%
Curve (Base chain 2024)	\$450M	\$95M (3 months later)	21.1%
Balancer Liquidity Mining	\$850M	\$200M (2021-2022)	23.5%
Average (10+ protocols)	-	-	15-25%

Cost Comparison:

- Traditional liquidity mining: **\$10 spent per \$1 retained** liquidity
- Protocol-Owned Liquidity: **\$0.50 deployed per \$1 TVL** (20x more efficient)

Economic Benefits:

- All LP fees accrue to treasury
- Stability enhancement, reduced slippage
- “Cheaper and more stable” than liquidity mining

6.4 AMM Mechanism Design Advances**Recent Research (2024-2025):**

- Undergirding Bonding Curves (UBC): Continuous value growth representation
- PABC framework: Design toolkit for efficient AMMs
- Sharded AMMs (SAMM): Scalability through horizontal partitioning

Pump.fun Success:

- 11 million+ token launches since January 2024
- \$800 million+ in fees generated
- Binance adopted similar bonding curve mechanism

Fee Structure Optimization (Uniswap v3):

- High-fee pools: 58% of liquidity, only 21% of volume
- v3 vs v2: 6x lower mean illiquidity (0.005 vs 0.030)
- Informed traders route to higher-fee pools with sufficient depth

6.5 Token Buyback Programs (2024-2025)**Scale of Adoption:**

- Top 12 DeFi protocols: ~\$800M on buybacks/dividends (2025)
- 400% increase from early 2024
- 28 projects: \$1.4B+ on buybacks in 2025

Notable Programs (v2.0 – expanded):

Protocol	Buyback Volume	Mechanism	Impact
Aave	\$52M/yr (\$1M/week)	Weekly TWAP buyback	Surplus buffer model
Hyperliquid	~\$1.2B/yr (97% of fees)	Continuous TWAP	380% price rally
BNB	~\$1B/quarter	Quarterly Auto-Burn formula	104M+ BNB burned total
MakerDAO	Variable (after \$50M buffer)	Smart Burn Engine	25,000+ MKR burned
Orca	\$10M treasury buyback	One-time + 25% supply burn	Immediate impact

Performance:

- Buyback projects outperformed non-buyback by 46.67% in 2024
- DWF Labs research: buybacks generate 15-30% higher long-term returns vs dividends for crypto tokens
- However, price response to announcements is mixed

Buyback Effectiveness Conditions:

1. Funded by real revenue (not inflationary or treasury-depleting)
2. Sustained over time (not one-off events)
3. Transparent and predictable (market can price in demand)
4. Meaningful relative to daily volume (>1-2%)
5. Combined with burn (permanent removal) rather than redistribution

7. Network Effects & Adoption Economics

7.1 Metcalfe's Law in Crypto

Validation (2024-2026):

- Network value proportional to square of user count
- Valid for evaluating cryptocurrencies in medium to long run
- Short-term applicability “highly debatable”

Strongest for Mature Networks:

- Ethereum: Smart contract infrastructure creates positive feedback
- More developers, then more users, then exponentially increased value

Limitations:

- Assumes equal user contribution to value
- Doesn't account for user quality/engagement
- Ignores external factors (regulation, technology issues)

7.2 Two-Sided Marketplace Economics

Tokenized Marketplace Architecture:

- Token incentives most effective targeting supply-side constraints
- “Actively supplied” networks have greater defensibility

Bootstrapping Strategy:

- Dynamically scale supply-side incentives
- Concentrate rewards where supply constraints exist
- First to greatest demand liquidity gains competitive moat

Market Context:

- Global marketplace GMV: \$7.2 trillion (2024)
- 35% of all digital commerce
- By 2030: Hybrid models (centralized UX + decentralized value capture) dominant

7.3 Token Incentive Design Evolution

2024-2025 Paradigm Shift:

- Moved from hype-driven to sustainable, incentive-aligned ecosystems
- Tokens designed to “motivate desired behavior”
- Rewards aligned to actual network health metrics

Distribution Approaches:

- Points-based systems (Blast, LayerZero, friend.tech)
- Restaking mechanisms (EigenLayer, EtherFi, Pendle)
- Over \$4B distributed via airdrops in 2024

Critical Finding:

“The most influential factor in a successful token economy is not profit maximization but fostering a user-centric community where engagement and empowerment are prioritized.”

7.4 Critical Mass & Tipping Point**10% Penetration Threshold:**

- Global crypto user penetration reached 11.02% in 2025
- Surpassed 10% threshold that triggers exponential growth acceleration
- Historical technology adoption (internet, smartphones) shows similar pattern

2025 Adoption Metrics (bull-cycle snapshots – treat with caution):

- 30% of American adults (70.4 million) own crypto
- One-third of SMBs use crypto (2x 2024 rate)
- 83% of institutional investors planned to increase allocation
- BlackRock IBIT ETF: \$50B+ AUM in under one year

Market Conditions Reality Check (v2.3 – July 18, 2026): Crypto is in a pronounced bear phase: total market cap ~\$2.19-2.27T, down ~43% year-over-year from the ~\$4.27T all-time high of October 6, 2025; BTC trades ~\$64,000 with BTC dominance ~57-60% depending on source (CMC vs Coinbase; BTC mcap ~\$1.29T against the ~\$2.2T total implies ~57-59%) and sustained ETF outflows. The “severe bear” contingency modeled elsewhere in this document (§15.4) is approximately the live base case. Counterpoint favorable to Gonka’s thesis: AI-linked tokens posted the *smallest sector decline* (~-14%) in Q1 2026 – with individual leaders TAO, FET, and RENDER posting outright gains while 38% of altcoins traded near all-time lows. (An earlier “only profitable sector” framing overstated this; “least-bad sector, with revenue-generating leaders posting gains” is the defensible claim.) Drivers were demand-side catalysts – Bittensor’s \$43M Q1 revenue, NVIDIA CEO Jensen Huang publicly endorsing decentralized AI training (comparing Bittensor to “a modern folding@home”), and the successful Covenant-72B decentralized training run. The pattern continued through Q2 2026: NEAR became the sector breakout (+72.1% seven-day surge in late May on an Arthur Hayes endorsement, NVIDIA Inception acceptance, and >\$3B cross-chain Intents volume), with FET (+23.3%) and VVV (+19.4%) leading late-May rallies; the sector cap spiked above \$26B in late May 2026 before retracing to ~\$22B by mid-July 2026 (~\$3.4B daily volume at the May reading), with individual leaders still volatile (recent weekly pops of ~25-35% in FET/RENDER). Revenue-generating DePIN compute networks decoupled from pure-speculation tokens in the selloff – direct support for the fee-revenue thesis.

7.5 Liquidity Network Effects**2025-2026 Paradigm:**

- Market primarily driven by global liquidity dynamics and macroeconomic catalysts
- Traditional four-year halving cycle theory becoming obsolete

Reflexivity Mechanics:

- Leverage, automated liquidations, and ETF flows amplify both directions
- “Reflexivity works to the downside just as much as upside”

Developer Growth Network Effects (v2.0; baseline is the Feb 2026 snapshot):

Network context (v2.3): Gonka compute passed 10,000 H100-equivalents (with ~100M tokens/day across the top-5 inference models), then 12,000, and was reported at ~14,000 H100-equivalents by early February 2026 (~52% monthly growth from ~6,000 at Bitfury’s December 2025 investment), with nodes in ~20 countries and H100/H200/A100-class GPUs making up >80% of compute. **The 14,000 figure did not hold – and neither did the April ~4,648.** As of July 2026 joingonka.ai’s live counter shows **~1,178 GPUs active**, consistent with tracker.gonka.vip’s ~1,214 – far below the April 2026 joingonka.ai reading of ~4,648 GPUs (~113 independent participants, ~582 MLNodes) and CoinMarketCap’s static “~5,000 H100 GPUs” project description, which is stale marketing text. Treat the live counters (~1,200 GPUs) as the active-mining denominator; ~14,000 (Feb 2026) was the announced peak. Real-time explorers (gonka.gg with a free public API, gonkascan.com, gonkahub.com, tracker.gonka.vip) publish participant/GPU counts – use them, not announcement figures, in any per-GPU economics. The 2,200-developer figure below is the February 2026 count (not re-verified for July 2026).

Developer Milestone	Timeline	Driver
2,200 (Feb 2026 baseline)	Baseline	Organic + early adopters
6,000 active	6 months	Free tier + hackathons

Developer Milestone	Timeline	Driver
15,000 active	18 months	Framework partnerships + grants
25,000+ active	36 months	Network effects + organic growth

8. Collateral & Slashing Economics

8.1 Optimal Collateral Ratios

DeFi Standards:

- Typical threshold: 150% (require \$150 collateral to borrow \$100)
- Protects against volatile asset price swings

Liquid Staking Revolution (H1 2025):

- Total LSD TVL: \$58.9 billion across all blockchains
- Ethereum liquid staking: ~\$44.8 billion (94% of market)
- LSD protocols: \$65B+ total market cap

Capital Efficiency via LSDs:

- Multi-layer returns: Staking rewards + DeFi yields
- 3-4% SRR on Ethereum (2025) + lending/yield farming
- Risks: Smart contract vulnerabilities, de-pegging

8.2 Slashing Penalty Calibration

Ethereum's Correlated Penalty Design:

- Severity scales with validators slashed simultaneously
- Minor penalties for isolated incidents
- Escalating penalties for mass events

Penalty Structure:

- Extended downtime: Minor deductions
- Safety violations (double-signing): 5%+ of bonded stake
- Day 18 "correlation penalty" for simultaneous slashing

Inactivity Leak Mechanism:

- Activates if consensus fails to finalize for 4+ epochs
- Validator deposits slowly drain
- Creates conditions for chain recovery

8.3 Unbonding Periods & Liquidity Trade-offs

Network	Unbonding Period
Ethereum	~11 days
Solana	1 epoch (~2-3 days)
Cosmos	21 days
DYDX	30 days

8.4 Insurance & Risk Pooling

Slashing Insurance Expansion:

- Munich Re: Ethereum PoS staking risk insurance
- Unslashed Finance: \$100M+ staked assets with successful claims

8.5 Collateral-Backed Governance

veGNK and Collateral Separation (v2.0):

A key design decision from the governance research: **locked veGNK does NOT count toward the 0.0625 GNK/nonce host collateral requirement.**

Rationale:

- Clean separation: Collateral is for security, veGNK is for governance
- Simpler slashing mechanics
- Prevents governance centralization by large hosts

- Encourages broader GNK holder participation (not just hosts)

Gonka Application:

- 0.0625 GNK collateral per nonce creates voting weight
- 20% base weight + 80% collateral-weighted system
- Aligns economic stake with network security contribution

9. Dynamic Pricing Mechanisms

9.1 EIP-1559 Economic Analysis

Mechanism:

- Automated base fee adjusting based on congestion
- Base fee burned (prevents miner extraction)
- User-set priority fees on top

Empirical Results:

- Reduced fee uncertainty by ~40%
- Improved user experience and fee prediction
- Limited direct fee reduction, but better estimation
- Exhibits chaotic dynamics under optimal conditions

9.2 EIP-4844 Blob Pricing (Proto-Danksharding)

Technical Implementation (March 2024):

- “Blob” transactions: 128 KB data chunks
- Temporary consensus layer storage (~18 days)
- Separate blob gas market with dynamic pricing

Economic Impact:

- 10-100x reduction in L2 data posting costs
- Base: 224% transaction volume increase post-Dencun

9.3 Pectra Upgrade (May 2025)

Blob Capacity Expansion:

- Target: 3 to 6 blobs per block
- Maximum: 6 to 9 blobs per block
- Daily capacity: 5.5GB to 8.15GB

Cost Impact:

- 51% reduction in daily rollup costs
- \$20,660/day to \$11,015/day average

9.4 EIP-1559 Parameter Sensitivity (v2.0)

Gonka's Current Implementation:

- Target utilization zone: 40-60% (stability zone)
- Base fee adjustment: +-2% per block
- Within stability zone: No base fee change
- Above 60%: Base fee increases up to 2% per block
- Below 40%: Base fee decreases up to 2% per block
- Floor: 1 nicoin per AI token

Parameter Sensitivity Analysis (from stress test research):

Adjustment Rate	Convergence Speed	Stability	Recommendation
+/-1%	Very slow (200+ blocks)	Very stable	Too conservative
+/-2% (current)	Moderate (80 blocks for 80% adjustment)	Stable	Acceptable baseline
+/-4% (recommended for testing)	Fast (40 blocks)	Stable	Better market tracking
+/-6%	Very fast (27 blocks)	Approaches instability	Not recommended
+/-8%+	Extremely fast	Unstable oscillations	Dangerous

Key Finding: Academic research shows optimal adjustment rates for compute markets are higher than Ethereum's gas market, because compute demand has higher variance. $\pm 4\%$ provides faster convergence with acceptable volatility.

Dual Volatility Problem: When GNK appreciates 5x, base fee needs $\sim 80\%$ reduction to maintain USD parity. At $\pm 2\%$, this requires ~ 80 consecutive blocks of downward adjustment. At $\pm 4\%$, only ~ 40 blocks – significantly faster market response.

9.5 Congestion Pricing Theory

Price vs Quantity Control:

- Price controls (Ethereum) outperform when:
 - Significant demand volatility
 - Low correlation between marginal costs and demand
 - High validator bargaining power

Queuing Theory Applications:

- Priority queue models derive stability conditions
- Base fee adjustment creates dynamic resource allocation
- Revenue maximization vs welfare maximization trade-offs

10. Decentralized Governance Economics

10.1 Quadratic Voting & Funding

Mechanism:

- Cost of n votes = n^2 credits
- 100-token holder: Only 10 votes (vs 100 in linear)
- Reduces whale domination

Sybil Vulnerability (v2.0):

Quadratic voting's Achilles heel is wallet splitting:

10M GNK in 1 wallet: $\sqrt{10M} = 3,162$ votes
 Split across 100 wallets: $100 * \sqrt{100K} = 31,600$ votes (10x more)

Gonka Application: Quadratic voting only viable for host-gated Community Pool decisions where GPU-based identity provides Sybil resistance. NOT recommended for general governance without proof of personhood.

Stanford Research (2025):

“Unpermissioned blockchains with quadratic mechanisms remain vulnerable to Sybil Attacks through wallet creation strategies.”

10.2 Conviction Voting (Polkadot OpenGov)

Mechanism:

- Voting power = tokens * conviction multiplier
- Longer lock periods = higher conviction = stronger votes
- Novel delegation by conviction and commitment

Advantages:

- Reduces rash decisions through time commitment
- Aligns incentives with long-term protocol health
- Prevents flash loan governance attacks

10.3 Plutocracy Problems

Empirical Evidence:

- Compound: Top 10 voters control 57.86% of voting power
- Uniswap: Top 10 voters control 44.72%
- Decentraland: Only 0.79% average participation per proposal

Governance Concentration Risk for Gonka (v2.0):

Scenario	Founder veGNK	Total veGNK	Founder Control
Founders don't lock	0	300M (others)	0%
Founders lock 1 year	100M	400M	25%
Founders lock 2 years (max)	200M	500M	40%

Mitigation: Voluntary lock caps, delegation to distribute voting power, community distribution before veGNK launch.

10.4 Governance Attack Vectors**Flash Loan Attacks:**

- February 2025: AttackDAO extracted \$67M
- Beanstalk: \$182M governance attack
- Atomic: Borrow, pass proposal, extract, repay in single transaction

veGNK Defense (v2.0): veGNK fully mitigates flash loan governance attacks because locked tokens cannot be borrowed. This is the strongest argument for transitioning from token-weighted to time-weighted governance.

Public Acquisition:

- Justin Sun: 30% of Steem's supply for witness control

10.5 Defense Mechanisms

veGNK provides Tier 2 + Tier 3 defense simultaneously:

Tier 2 - Increase Acquisition Costs:

- Time-locked tokens cannot be borrowed for attacks
- Reduces token liquidity through staking

Tier 3 - Increase Execution Costs:

- Must lock for 1 month minimum to gain any voting power
- 2-year lock for maximum influence
- No early exit – attacker capital trapped

11. Gonka-Specific Applications**11.1 Emission Design Validation****Gonka's Formula:**

```
current_epoch_reward = 323,000 * exp(-0.000475 * epochs)
```

Research Alignment:

- Exponential decay: Optimal balance between early growth and long-term sustainability
- Avoids "halving shock" disruptions of step-function models
- ~4-year halving equivalent provides Bitcoin-comparable scarcity
- Inflation rate converges to <3% within 10 years

11.2 Sprint Consensus Economics**Research Insights Applied:**

- PoW security with productive computation
- Addresses environmental concerns (99.95% efficiency vs hash PoW)
- Verification efficiency (1-10% vs 100% redundancy)

Economic Advantages:

- Compute resources generate network utility (AI inference)
- Miners earn from both block rewards and compute provision
- Sustainable security model beyond pure token inflation

11.3 Collateral System Optimization**Current Design:**

- 20% base weight (unconditional)
- 80% requires 0.0625 GNK collateral per nonce

- 180-epoch grace period

Research Validation:

- Collateral-weighted systems align economic incentives with network security
- Grace periods enable bootstrapping (similar to zero pricing phases)
- Slashing mechanics (20% malicious, 10% poor performance) align with best practices

11.4 Governance Parameter Design

Current Parameters:

- 33.4% quorum threshold
- 50% majority threshold
- 33.4% veto threshold

Recommended Enhancement (v2.0): Transition to veGNK-based governance in 3 phases:

1. **Phase 1 (Q2 2026):** Basic lock + voting power (1 month - 2 year range)
2. **Phase 2 (Q4 2026):** Boost mechanics for AI Training Fund yield + delegation
3. **Phase 3 (2027):** Advanced features based on adoption data

11.5 Dynamic Pricing Validation

Current Design:

- EIP-1559 inspired
- 40-60% utilization stability zone
- +-2% max change per block

Recommendation (v2.0): Test +-4% adjustment rate for faster market convergence. Implement oracle-based USD pricing with GNK settlement to eliminate dual volatility problem (GPU deflation + GNK volatility).

11.6 Market Position Analysis

Updated Competitive Landscape (v2.1 – July 2026):

Feature	Gonka	Akash	Render	Bittensor	io.net
OpenAI-compatible API	Yes	No	No	No	No
Productive compute	~98%	Variable	~90%	~40%	High
Real yield to token holders	Planned (5% + surplus)	Take Pool to stakers	None	None	None
Buyback mechanism	Planned (5% burn)	Live (BME, Mar 2026)	Burn-mint	None	IDE live Jun 2026 (burn currently emission-funded)
Dynamic pricing	EIP-1559	Reverse auction	Fixed	Subnet-dependent	Variable
Floor price defense	Planned (TWAP)	None	None	None	None
H100 \$/hr (est.)	\$1.50-2.50 target band (set pre-2026 price stall; revisit)	\$1.80-2.80	AI compute onboarding (RNP-021)	N/A	\$1.50-2.50

Note: Chutes (Bittensor Subnet 64 – the largest decentralized inference provider, ~120B tokens/day steady-state, reachable through OpenRouter with zero config; see §4.5), Aethir (\$127.8M FY2025 revenue, enterprise contracts), Prime Intellect (>100M annualized revenue, notoken), Gensyn (AI token, mainnet Apr 2026), and BitTorrent's BTTInferGrid (Jun 2026) now also belong in any full landscape view – see §4.6. In particular, the “OpenAI-compatible API: Yes only for Gonka” row above is no longer a clean differentiator: Chutes serves 9 models via OpenRouter's standard API (as of Jul 2026).

Decentralized AI Compute Market:

- \$9 billion (2024) to \$100 billion (2032) projection
- Gonka positioned for enterprise/professional workloads
- Cost advantages: 60-80% vs hyperscalers, at parity with decentralized peers but with superior API quality

11.7 Enhanced Revenue Allocation (v2.0)

Current Model:

Inference Revenue: 100%
 70% → Hosts
 20% → AI Training Fund
 10% → Unallocated

Proposed Model:

Inference Revenue: 100%
 70% → Hosts (UNCHANGED)
 20% → AI Training Fund (UNCHANGED, surplus to veGNK holders)
 5% → GNK Buyback and Burn (NEW – continuous TWAP)
 5% → veGNK Staker Yield Pool (NEW – real yield)

Revenue Impact at Different Scales:

Annual Revenue	Buyback (5%)	Yield Pool (5%)	Monthly Buyback
\$5M	\$250K	\$250K	~\$21K
\$25M	\$1.25M	\$1.25M	~\$104K
\$100M	\$5M	\$5M	~\$417K
\$500M	\$25M	\$25M	~\$2.1M

11.8 Implementation Roadmap (v2.0)

Priority-Ordered Recommendations:

Priority	Recommendation	Timeline	Complexity
1	Oracle-based USD pricing with GNK settlement	Q2 2026	Medium
2	veGNK Phase 1 (basic lock + voting)	Q2 2026	Medium
3	Enhanced revenue split (20/70/5/5)	Q2-Q3 2026	Low
4	POL deployment (Phase 1: 22M GNK)	Q3 2026	Medium
5	Continuous TWAP buyback-and-burn	Q3 2026	Medium
6	Floor price defense mechanism	Q3-Q4 2026	Medium
7	veGNK Phase 2 (boost + delegation)	Q4 2026	High
8	Developer growth program (free tier + grants)	Ongoing	Low
9	EIP-1559 parameter testing (+-4%)	Q2 2026	Low
10	veGNK Phase 3 (advanced features)	2027	High

(v2.1 note: the Q2 2026 target dates above have passed; implementation status was not re-verified in the July 2026 update. Treat timelines as the original v2.0 plan, not current status.)

12. Protocol-Owned Liquidity Strategy

(New section – v2.0, sourced from research/01-pol-and-liquidity.md)

12.1 POL vs Liquidity Mining: Quantified Comparison

Traditional liquidity mining retains only 10-25% of incentivized liquidity after emissions cease. Protocol-Owned Liquidity provides permanent, fee-generating positions.

Cost Analysis:

- **Liquidity Mining:** 10M GNK emissions to attract \$5M TVL. After emissions end, \$1M retained. Cost: **\$10 per \$1 retained.**

- **POL Deployment:** 20M GNK deployed with paired assets for \$40M TVL. 100% retained. Cost: **\$0.50 per \$1 TVL.**

POL is 20x more capital efficient.

12.2 Leading Protocol Benchmarks

Protocol	Treasury Size	POL Allocation	% of Treasury	Liquidity Depth
Olympus DAO	\$30M	\$28M	93%	\$15M
Frax Finance	\$250M	\$45M	18%	\$60M
GMX	\$180M	\$32M	18%	\$55M
Gains Network	\$85M	\$22M	26%	\$35M
Average	-	-	15-35%	-

12.3 Gonka POL Deployment Plan

Recalibration required (v2.1): This plan was modeled at \$1.00 GNK. As of July 18, 2026, GNK trades ~\$0.13 (CoinMarketCap: \$0.1307, mcap ~\$13.85M, circulating ~105.92M of 1B max; ATL \$0.1258 on July 17, 2026), after an all-time high of \$2.61 on January 16, 2026 – a ~95% drawdown. GNK is still not listed on major CEXs (the anticipated MEXC/Gate listings never materialized); it trades OTC on HEX Exchange, on SafeTrade (GNK/USDT), and as wrapped GNK. (Caution: aggregator prices diverge wildly on thin liquidity – Crypto.com \$0.133, Bitget \$0.276, CryptoRank \$0.2756 (which also shows a different ATL: \$0.1462 on Jul 16, 2026), Coinpaprika \$0.44 – use CMC as canonical; and “Gonka AI” tokens on Solana/PumpSwap at ~\$0.0001-0.0005 are unofficial/fake and must not be cited.) At \$0.13, 22M GNK is ~\$2.9M of protocol-side value, not \$22M; USD depth targets, paired-asset needs, and all ranges below must be rebased to the live price before deployment.

Recommended Allocation (original \$1.00-GNK model, for structure only):

- **Total:** 22M GNK (18.3% of 120M Community Pool)
- **GNK/USDC pair (60%):** 13.2M GNK + matching USDC
- **GNK/ETH pair (40%):** 8.8M GNK + matching ETH

Uniswap v3 Parameters:

- Fee tier: 0.3% (standard for medium-volatility governance tokens)
- Range width: +-25-35% around spot (the \$0.75-\$1.35 band in v2.0 assumed \$1.00 spot; recenter on live price)
- Capital efficiency: 3.8-4.2x vs full range

Expected Outcomes (scale with deployed USD value):

- <1% slippage for trades sized ~0.1% of pool depth
- LP fee revenue: 3-6% APR on deployed value
- Rebalancing cost: negligible vs revenue

12.4 Phased Deployment

Phase	Timeline	Action	Success Metrics
Phase 1	Month 1-2	Deploy 22M GNK with paired assets	TVL/volume targets rebased to live GNK price
Phase 2	Month 3-6	Optimize ranges, evaluate fee tiers	<4 rebalances, >3% APR
Phase 3	Month 7-12	Consider additional 8M GNK or new pairs	Enhanced diversification

12.5 LP Fee Revenue Strategy

- **Years 1-2:** Reinvest LP fees into deepening POL positions
- **Years 3+:** Governance vote on fee usage: reinvest, distribute to veGNK holders, or burn

12.6 Paired Asset Challenge

Gonka may lack the USDC/ETH needed for pairing (~\$20-25M at the \$1.00-GNK model; ~\$2.9M at the July 2026 price of ~\$0.13). Mitigations:

1. Phased deployment with available assets
2. OTC sales of GNK for stablecoins
3. Bitfury/investor negotiation for paired asset contribution
4. Single-sided deployment via Tokemak Autopilot (Phase 2 backup)

13. Real Yield & Revenue Distribution

(New section – v2.0, sourced from research/02-real-yield-and-buybacks.md)

13.1 Enhanced Revenue Allocation Model

From 20/80 to 20/70/5/5:

```
Inference Revenue: 100%
|
+-- 70% -> Hosts (compute providers)      [UNCHANGED]
|
+-- 20% -> AI Training Fund                [UNCHANGED]
|
|      +-- [Surplus above 6-month runway] -> veGNK holders
|
+-- 5% -> GNK Buyback and Burn             [NEW]
|
+-- 5% -> veGNK Staker Yield Pool          [NEW]

Base Fee: 100% BURNED (EIP-1559 mechanism) [UNCHANGED]
```

13.2 AI Training Fund Surplus Mechanism

Runway-Based Threshold (MakerDAO Surplus Buffer pattern):

```
Monthly_Fund_Expenses = average monthly AI development spending
Runway_Target = 6 months
Threshold = Monthly_Fund_Expenses * 6
Surplus = Fund_Balance - Threshold
```

```
If surplus > 0: distribute to veGNK holders
If below threshold: all incoming revenue stays in fund
```

Governance-adjustable: 3 months (aggressive) to 12 months (conservative).

13.3 Buyback-and-Burn Mechanism

Design:

- Continuous TWAP orders, executing small buys every ~15 minutes (96 buys/day)
- Maximum 0.5% slippage per order
- All purchased GNK sent to burn address (permanent supply reduction)
- Minimum \$10 per order to avoid dust

Opportunistic Dip Buying:

- 3x accelerated buyback when GNK drops >20% below 30-day TWAP
- Draws from reserve buffer for opportunistic accumulation

Governance Emergency Pause:

- Requires 33.4% quorum + >50% majority
- Paused funds redirect to AI Training Fund
- Auto-resumes after 30 days unless re-paused

13.4 Total Deflationary Pressure

```
Net Supply Change = New_Emissions - Base_Fee_Burns - Buyback_Burns
```

Three deflationary mechanisms:

1. Base Fee Burns (EIP-1559) -> Usage-proportional (every transaction)
2. Buyback Burns (5% revenue) -> Revenue-proportional (TWAP)
3. Emission Decay -> Time-proportional (exponential)

Expected crossover to net deflationary: Year 3-5

13.5 Tax and Regulatory Considerations

Distribution Method	Tax Treatment	Securities Risk (pre-2026 framing)
Buyback-and-burn	No tax event until sale	Lowest
Auto-compound (xGNK)	Capital gains at unstake	Low
Direct claim to veGNK	Ordinary income at receipt	Higher

2026 US Regulatory Shift (v2.1): The risk differential above predates a major 2026 change in US posture:

- SEC “Project Crypto” under Chair Atkins (A-C-T strategy); Jan 28, 2026 SEC staff token-taxonomy statement; Mar 11, 2026 SEC-CFTC MOU; Mar 17, 2026 Atkins token-safe-harbor speech. Governance/utility tokens are now largely classified as “digital tools” – not securities when sold for use rather than speculation – materially lowering the risk premium previously assigned to direct veGNK yield distribution.
- **Pending SEC token-sale rulemaking (v2.3):** The SEC’s “Regulation Crypto” proposal is confirmed for the July 2026 rulemaking slot – actually three linked proposals: token offer/sale rules (a time-limited “Innovation Exemption” allowing raises up to \$75M in any 12-month period, a 12-36 month safe harbor, and an exit-from-securities-status mechanism), broker-dealer compliance amendments, and crypto market-structure/ATS rules. Two large caveats before relying on it: (1) **Timeline** – the proposal is still under review at the White House OIRA; after publication comes a comment period, final rule, and compliance dates, so an operative rule is quarters away, not weeks. (2) **Eligibility** – the exemption targets early-stage projects (reported criteria include valuation under \$5M within the first four years), which Gonka (mcap ~\$14M, launched 2025) may not satisfy. The rulemaking *may* eventually open a compliant sale route, but should not be treated as an available option in treasury planning.
- The CLARITY Act passed the House but remains at Senate Calendar No. 423 with no cloture motion as of July 18, 2026 – and **no floor vote is scheduled (v2.4)**. Bipartisan talks on the ethics/law-enforcement provisions collapsed; prediction-market odds of passage fell from the low 70s to roughly 43% (near coin-flip); the July 17 event was only a House Financial Services *field hearing* in New York – a pressure/messaging event that cannot pass anything. Senators Murphy, Van Hollen, and Merkley remain formally opposed, the bill still needs ~7 Democratic votes for cloture, and August 7 is the final session day before recess. Three disputes still block passage: ethics provisions, Section 604 (opposed by the National District Attorneys’ Association), and stablecoin yield (Coinbase earns ~\$1.35B/yr in USDC rewards). The statutory picture is unresolved even as agency posture has softened.
- **GENIUS Act implementation (v2.5 – deadline missed):** July 18, 2026 was the statutory deadline for six federal agencies (OCC, FDIC, Treasury, FinCEN, et al.) to finalize stablecoin implementing rules under the GENIUS Act. **It was missed:** eight proposed rules have been published but zero finalized, and on June 22 regulators effectively conceded the miss by publishing three more proposals whose comment periods extend past July 18. The Act’s effective date now defaults toward the statutory backstop of **January 18, 2027** (18 months from enactment), with issuers getting ~120 days to comply after rules eventually finalize. Key proposed terms so far: OCC \$5M minimum capital floor with a three-tier liquidity framework (10% same-day redemption), FDIC confirming no deposit insurance for token holders, and par redemption within two business days. For Gonka’s treasury/POL planning (USDC pairing legality and yield treatment, §12): treat the stablecoin regime as **not operative until late 2026/early 2027 at the earliest**, not as an input landing this summer. Also still relevant to the stablecoin-yield dispute blocking the CLARITY Act.

Recommendation (revised): Buyback-and-burn remains tax-efficient for holders, but the 2026 SEC posture weakens the case for avoiding direct veGNK yield on securities-risk grounds alone. Continue positioning veGNK yield as governance participation rewards, and monitor the SEC proposal’s comment period (and Gonka’s likely ineligibility for the Innovation Exemption as drafted), the stalled CLARITY Act (no vote scheduled; ~43% odds), and the delayed GENIUS Act stablecoin rules (no operative regime before late 2026/early 2027) before finalizing distribution mechanics.

14. Fee-to-Emission Transition Stress Test

(New section – v2.0, sourced from research/04-fee-transition-stress-test.md)

14.1 The Critical Transition

As epoch rewards decay exponentially, inference fee revenue must scale proportionally or hosts will exit. This is Gonka’s most important economic challenge.

14.2 Comparable Network Transitions

Network	Emission Model	Fee Revenue % (2026)	Crossover Timeline
Bitcoin	Step halving (4yr)	10-15%	Never (declining)
Ethereum	PoS issuance	30-40%	N/A (PoS transition)
Filecoin	Exponential decay	~45%	Year 8-9
Akash	Linear decay	~40%	Year 4-5
Gonka	Exponential decay	0% (Year 1) target 50%+	Year 4-10

14.3 Three-Scenario Revenue Model

Scenario	Dev Growth	Year 4 Revenue	Year 8 Revenue	Crossover Year
A: Conservative	10%/yr	\$96.9M	\$142.4M	Year 4 (\$1 GNK) to Year 8-9 (\$5 GNK)
B: Moderate	25%/yr	\$1.24B	\$12.1B	Year 1-3
C: Aggressive	50%/yr	\$26.3B	Very high	Year 1-2

Target trajectory: Scenario B (moderate). Plan for Scenario A (conservative) as baseline.

14.4 Host Profitability Thresholds

Breakeven GNK Price by Year (Conservative Scenario):

Year	Epoch Reward/Day (GNK per host)	Inference Fee/Day (per GPU)	Breakeven GNK Price
1	361	\$8	\$0.15
4	170	\$15	\$0.35
8	80	\$25	\$3.30
12	44	\$35	\$1.50*
20	15	\$65	\$3.50*

At Year 8 conservative scenario, GNK must exceed \$3.30 for hosts to beat traditional GPU rental (\$1,937/month). The v2.0 caveat that GPU deflation could cut this to \$0.62 (H100 at \$1.50/hr by 2028) is now doubtful: rental deflation stalled in 2026 (see §5.2), which keeps the breakeven threshold high. Note also that this table was built on the ~6,000 H100-equivalent network of early 2026. The denominator matters enormously: any break-even math computed on the February 2026 announcement figure of ~14,000 H100-equivalents overstates dilution – live trackers show only **~1,178-1,214 active GPUs as of July 2026* (§7.5), roughly 12x the per-GPU emission share versus the 14,000 assumption and ~4x versus a 5,000-GPU assumption (on the order of ~250-270 GNK/day per GPU rather than ~21), pulling host break-even GNK price down roughly 4x further than the ~\$0.80 estimate computed on a 5,000-GPU network (toward ~\$0.20). At the July 2026 GNK price (~\$0.13) hosts remain marginally below traditional rental rates even on the live ~1,200-GPU denominator – but the gap is far narrower than the 5,000-GPU math implied.

14.5 Crossover Sensitivity to GNK Price

Scenario	GNK = \$1	GNK = \$2	GNK = \$5
Conservative	Year 4 crossover	Year 6	Year 8-9
Moderate	Year 1	Year 2	Year 3
Aggressive	Year 1	Year 1	Year 2

Key Insight: Higher GNK price paradoxically pushes crossover later (epoch rewards worth more in USD), but also attracts more hosts. The system self-balances through the supply-demand equilibrium.

14.6 Contingency Plans

If Year 4 fee revenue falls below \$50M (vs \$96.9M conservative baseline):

1. **Governance-activated tail emissions:** Small perpetual emission to maintain host incentives
2. **Enhanced developer subsidies:** Redirect Community Pool funds to developer growth programs
3. **Host efficiency programs:** Help hosts reduce operating costs (power, cooling optimization)
4. **EIP-1559 parameter adjustment:** Faster fee convergence to attract developers

14.7 Early Warning Indicators

Indicator	Green	Yellow	Red
Developer growth rate	>20%/yr	10-20%/yr	<10%/yr

Indicator	Green	Yellow	Red
Fee % of host income	>30% by Year 4	15-30%	<15%
Host churn rate	<5%/yr	5-15%/yr	>15%/yr
Network utilization	40-60%	20-40%	<20%

15. Competitive Positioning & Developer Growth

(New section – v2.0, sourced from research/05-gpu-economics-and-developer-growth.md)

15.1 Gonka's Target Market Position

```

Hyperscalers ($7-12/hr)      <- Enterprise, compliance-heavy
|
Specialized Cloud ($1.99-3.50) <- AI-focused startups
|
[GONKA TARGET] ($1.50-2.50)  <- Cost-optimized, API-compatible
|
Decentralized Low ($0.80-1.50) <- Spot/best-effort, unreliable

```

(v2.1: bands updated to July 2026 market – H100 median \$2.29-3.12/hr, hyperscaler on-demand \$7-8. Centralized clouds have been raising, not cutting, prices in 2026, which widens Gonka's cost advantage but also raises the floor under the target band.)

Differentiated Value Proposition:

Feature	Hyperscaler	Specialized Cloud	Gonka	Others (Decentralized)
OpenAI-compatible API	No	Partial	Yes	Rarely
Censorship resistance	No	No	Yes	Yes
Dynamic pricing	No (fixed tiers)	Partial	Yes (EIP-1559)	Partial
Cost per H100/hr	\$7-12	\$1.99-3.50	\$1.50-2.50	\$0.80-3.00
On-chain verifiable	No	No	Yes	Partial

15.2 Oracle-Based USD Pricing

The Dual Volatility Problem: When GNK appreciates while GPU costs deflate, Gonka becomes progressively more expensive unless pricing adjusts. Fixed GNK-denominated pricing fails.

Solution: Oracle-based USD pricing with GNK settlement:

1. Set inference pricing in USD based on competitive market rates
2. Accept GNK payment at real-time oracle exchange rate
3. Adjust dynamically via EIP-1559 on the USD-denominated base price

Recommended Oracle Stack:

Oracle	Purpose	Speed	Security
Pyth	Real-time GNK price	400ms	Medium
Chainlink	Heartbeat GNK price	1 hour	High
UMA	GPU benchmark pricing	Weekly	Flexible

15.3 Developer Onboarding Strategy

Migration Path (OpenAI to Gonka):

```

# Only 2 lines change:
from openai import OpenAI
client = OpenAI(
    api_key="gnk-...",
    base_url="https://api.gonka.network/v1"
)
# All existing code works unchanged

```

Growth Program:

Phase	Timeline	Budget	Target Developers
Foundation	Months 1-6	1.85M GNK	3,800 active
Growth	Months 7-18	1.7M GNK	9,070 active
Scale	Months 19-36	5M GNK	10,000+ active

Blended acquisition cost: \$150-500 per active developer (2026 benchmark).

Marketing Message: “Same API. 70% Less Cost. Censorship-Resistant.”

15.4 Floor Price Defense Mechanism

Architecture: Programmatic TWAP buybacks triggered by price oracle, funded by Community Pool and inference revenue.

Trigger Conditions:

Tier	Trigger	Daily Buyback	Duration
1	GNK < 75% of 30-day TWAP	0.5% of treasury	Up to 30 days
2	GNK < \$0.45 absolute	1.0% of treasury	Up to 60 days
3	GNK < \$0.30 (crisis)	2.0% of treasury	Up to 90 days

Treasury Allocation: Up to 5% of Community Pool annually (6M GNK) + 5-10% of inference revenue.

Bitfury Anchor (v2.1 – corrected and updated): On December 1, 2025 Bitfury announced a **\$50M total investment commitment** to Gonka – the community-pool-approved \$12M GNK purchase at \$0.60/GNK was only the first tranche – publicly framed as the first strike in a ~\$1B plan to decentralize AI compute. The \$0.60 level was treated as an institutional Schelling point, with the Tier 2 trigger set at \$0.45 (25% below).

The Schelling point did not hold. GNK hit an all-time low of \$0.1258 on July 17, 2026 and trades ~\$0.13 – ~78% below the Bitfury entry and far below both the \$0.45 and \$0.30 triggers. The floor-defense mechanism was not deployed in time to be tested; had it been, the “sustained –60%+ bear market lasting 6+ months” treasury-depletion scenario – which is now the live base case, not a tail risk (see \$7.4) – would have been the operative condition.

Critical Limitation: Floor defense is a speed bump, not a wall. It slows declines and provides time for fundamentals to recover. The 2026 drawdown demonstrates that trigger levels anchored to a single strategic purchase price are fragile; any future implementation should anchor triggers to rolling TWAP levels and treasury runway rather than absolute dollar floors.

16. Research Sources & Citations

Academic Papers

1. Bitcoin & Emissions:

- [Bitcoin Return Prediction: Stock-to-Flow, Metcalfe's Law, Technical Analysis](#) - MDPI (2024)
- [Dissecting the stock to flow model for Bitcoin](#) - ResearchGate
- [The Imminent Security Risk of Bitcoin Halving](#) - SSRN (Sedlmeir et al.)

2. Consensus Mechanisms:

- [A Detailed Comparative Analysis of Blockchain Consensus Mechanisms](#) - arXiv
- [A Game Theoretic Analysis of Validator Strategies in Ethereum 2.0](#) - arXiv
- [Blockchain Consensus Mechanisms Primer](#) - IMF (2025)

3. Token Economics:

- [Tokenomics: Optimal Monetary and Fee Policies](#) - SSRN (Jermann, Xiang)
- [Single-Token vs Two-Token Blockchain Tokenomics](#) - DAGSTUHL
- [Tokenomics: When Tokens Beat Equity](#) - Management Science
- [Evaluating Token Economics for Web3 Infrastructure: Emission Schedules](#) - 1kx Network (2025)

4. DeFi Mechanisms:

- [Mechanism Design for Automated Market Makers](#) - arXiv
- [Price Discovery and Efficiency in Uniswap Liquidity Pools](#) - Journal of Futures Markets
- [DAO voting mechanism resistant to whale and collusion problems](#) - Frontiers
- [Token Buybacks in Web3: Trends, Strategies, and Impact](#) - DWF Labs (2025)
- [Buyback, Burning, and Supply: How Deflationary Tokenomics Shape the Crypto Market](#) - OKX Research (2025)

5. Fee Mechanisms:

- [Analysis of Dynamic Transaction Fee Blockchain Using Queueing Theory](#) - MDPI
- [Transaction Fee Mechanism Design in a Post-MEV World](#) - ePrint
- [EIP-4844 Economics and Rollup Strategies](#) - arXiv
- [The Future of Bitcoin Mining Incentives](#) - Sedlmeir et al. (2024)

6. Governance:

- [Going Parabolic: Analyzing Sybil Resistance in Quadratic Voting](#) - Stanford
- [Decentralizing governance: exploring dynamics of DAOs](#) - Frontiers
- [Balancing Security and Liquidity: Time-Weighted Snapshot Framework](#) - arXiv

Industry Reports

7. Market Analysis:

- [State of Akash Q3 2025](#) - Messari
- [Understanding io.net](#) - Messari
- [State of DeFi 2025](#) - DL News

8. Compute Markets:

- [AI Infrastructure Market Size](#) - Markets and Markets
- [GPU Pricing Trends 2025](#) - ACCIO
- [AI Inference Market Size](#) - Markets and Markets
- [NVIDIA H100 Deep Dive](#) - Fluence
- [GPU Cloud Market Size](#) - Grand View Research (2024)

9. Staking & Collateral:

- [Ethereum Staking: Second Half of 2025 Outlook](#) - Figment
- [EigenLayer Rewards v2 and Slashing](#) - Kiln

10. Network Effects:

- [Crypto's Market Penetration Tipping Point](#) - CoinDesk
- [Tokenized Marketplaces: Bootstrapping and Scaling](#) - Variant Fund
- [2026 Crypto Market Outlook](#) - Coinbase

POL & Liquidity Sources (v2.0)

11. Protocol-Owned Liquidity:

- [Olympus DAO Documentation](#)
- [Berachain Proof-of-Liquidity Documentation](#)
- [Tokemak v2 Documentation](#)
- [Balancer 80/20 Weighted Pools](#)
- [Uniswap v3 Whitepaper](#)
- [POL vs Liquidity Mining Analysis](#) - Gauntlet Research

Real Yield & Buyback Sources (v2.0)

12. Revenue Distribution:

- [GMX Documentation](#) - GMX Stats
- [Aave Governance Forum - Aavenomics Update](#)
- [Hyperliquid HYPE Analysis](#) - MEXC
- [MakerDAO Smart Burn Engine](#) - MakerBurn
- [Curve Finance Fee Distribution](#)
- [BNB Auto-Burn FAQ](#) - Binance

Governance & ve-Tokenomics Sources (v2.0)

13. Vote-Escrowed Models:

- [Curve Finance Analytics](#) - veCRV data
- [Convex Finance](#) - DeFi Llama
- [Velodrome/Aerodrome ve\(3,3\)](#) - The Block
- [Balancer veBAL Documentation](#)
- [PancakeSwap veCAKE](#)
- [Frax veFXS Documentation](#)

GPU & Developer Sources (v2.0)

14. GPU Economics:

- [NVIDIA H100 Deep Dive](#) - Fluence
- [Bittensor Halving Analysis](#) - Crypto.com

Conclusion

This comprehensive v2.0 research synthesis demonstrates that Gonka's tokenomics design aligns with cutting-edge macro-economic research across multiple dimensions:

1. **Emission Design:** Exponential decay provides optimal balance between growth incentives and long-term sustainability, validated by fee transition stress testing across three scenarios
2. **Consensus Mechanism:** Sprint Consensus addresses PoW sustainability concerns while maintaining security properties
3. **Collateral System:** Research-validated approach to aligning economic incentives with network security
4. **Dynamic Pricing:** EIP-1559-inspired mechanisms with recommended +-4% testing for faster market convergence
5. **Market Position:** Positioned within a rapidly growing decentralized AI compute market (\$9B to \$100B by 2032)

v2.0 recommendations for enhancement (with v2.1 status):

- **Protocol-Owned Liquidity:** Deploy 22M GNK via concentrated liquidity for permanent depth (rebase USD targets to the ~\$0.13 July 2026 price)
- **Real Yield:** Enhanced 20/70/5/5 revenue split with continuous buyback-and-burn + veGNK staker yield
- **veGNK Governance:** 3-phase rollout, fully mitigates flash loan attacks (original Q2 2026 start date has passed; status unverified)
- **Floor Price Defense:** Programmatic TWAP buybacks with transparent on-chain triggers (rebase triggers – the \$0.45/\$0.30 levels were breached in the 2026 drawdown)
- **Oracle Integration:** Pyth + Chainlink + UMA stack for USD-pegged pricing with GNK settlement
- **Developer Growth:** 8.55M GNK across three phases targeting 25K+ active developers by Month 36
- **Contingency Plans:** Governance-activated tail emissions and enhanced subsidies if fee growth lags

Gonka's combination of AI-productive PoW, exponential emission decay, and the proposed tokenomics enhancements positions it as one of the most comprehensive decentralized AI compute networks. With Akash's BME live and io.net's IDE nominally live as of H1 2026 (though io.net's burn is so far emission-funded rather than revenue-funded), the differentiator is no longer buyback/burn – it is direct fee-revenue distribution to veGNK lockers, funded by real revenue, which remains rare in the category, backed by growing 2026 evidence that revenue-generating compute networks outperform pure-speculation tokens.

July 2026 Verification Sources (v2.1-v2.5)

v2.5 additions (verified July 18, 2026):

- [OpenRouter Chutes provider page](#) – 9 models listed (Jul 18, 2026)
- [tao.media: The Investor's Guide to Chutes \(Feb 24, 2026\)](#) / [SimplyTAO: Subnet 64 guide](#) – corrected Chutes usage/revenue figures
- [CoinGape: Grayscale S-1 for Bittensor TAO ETF](#) – Bitwise TAO Strategy ETF N-1A noted in the same August 2026 window
- [Cryptonews: Bittensor Subnet 108 Frontier Compute](#)
- [Aethir 12-month strategic roadmap \(chain migration, v2, ATH Vault\)](#)
- [DailyCoin: GENIUS Act anniversary, rules missing](#) / [Crypto Impact Hub: GENIUS Act rulemaking sprint](#) / [Paul Hastings: GENIUS Act guide](#)
- [CryptoRank: Gonka](#) – aggregator-divergence datapoint (\$0.2756; ATL \$0.1462 Jul 16, 2026)

v2.4 additions (verified July 18, 2026):

- [joiningonka.ai live GPU counter](#) – ~1,178 GPUs active (Jul 2026)
- [Messari: Understanding the Render Network](#) / ["Understanding Dispersed"](#) (Jul 7, 2026) / [Salad x Render integration](#) / [RenderCon 2026: MCP + Salad subnet](#)
- [Grayscale GTAO S-1 amendment](#) / [8-K \(May 15, 2026\)](#) / [Phemex: Grayscale Bittensor spot ETF](#)
- [Kraken lists Bittensor subnet alpha tokens](#) (Jun 29, 2026)
- [io.net July 2026 unlock \(15.96M IO\)](#)
- [Nosana grant: Voight](#) (Jul 8, 2026) / [Nosana blog](#)
- [CLARITY Act: Senate showdown](#) / [vote in doubt](#) / [TechTimes: ethics impasse](#)
- [TrendForce raises Q3/Q4 DRAM forecasts](#) (Jul 9, 2026) / [ADATA Q3 memory pricing](#)
- [Benzinga/KeyBanc: Vera Rubin ramp slightly delayed](#) / [TechTimes: Rubin Ultra cancelled/scaled back](#)
- [AIMultiple GPU index](#) (Jul 2026) / [IntuitionLabs H100 rental comparison](#)
- [AI-crypto sector cap ~\\$22B \(mid-Jul 2026\)](#)
- [OpenRouter \\$113M Series B, 25T tokens/week](#) (Businesswire, May 26, 2026)

v2.3 additions (verified July 18, 2026):

- [joingonka.ai: What is Gonka / tracker.gonka.vip live dashboard](#) – live network GPU/participant counts (also [gonka.gg](#), [gonkascan.com](#), [gonkahub.com](#))
- [The Investor's Guide to Chutes – Bittensor's Inference Layer / tao.app Subnet 64](#)
- [Render RNP-021 text / Render latest updates \(CMC AI\)](#)
- [Wormhole bridges canonical TAO to Solana \(May 5, 2026\)](#)
- [Nosana: 3M jobs milestone / Nosana-OpenGPU \(Feb 5, 2026\)](#)
- [Gensyn on CoinGecko](#) – \$AI market status
- [TrendForce \(Jul 3, 2026\): Q3 memory price deceleration](#)
- [CLARITY Act stalls in Senate / CLARITY Act nears Senate floor](#)
- [SEC "Regulation Crypto" July 2026 slot / Innovation Exemption founder's guide](#)
- [OpenRouter State of AI \(100T-token study\) / a16z State of AI](#)
- [NVIDIA Rubin first shipments July 2026](#)
- [Thunder Compute B200 pricing / getdeploying B200](#)
- [Tianrong \(TIPS\) DEPIN Studios launch](#)
- [Phemex: AI tokens Q1 2026 sector data / SpotedCrypto: AI token rankings mid-2026](#)

v2.1-v2.2 sources:

- [Gonka on CoinMarketCap](#) – GNK price, supply, market cap
- [Bitfury \\$50M Investment in Gonka](#) – Business Wire (Dec 2025)
- [Akash Network Q1 2026 Report / Messari State of Akash Q1 2026](#)
- [io.net IDE Launch](#) (Jun 2026)
- [Aethir Q3 2025 Results](#)
- [Render Network Foundation Monthly Report, March 2026](#)
- [Bittensor Halving / Bittensor halving docs](#)
- [Thunder Compute H100 Pricing \(2026\) / B200 Pricing](#)
- [TrendForce: Blackwell shipment mix & memory supercycle](#)
- [Deloitte TMT Predictions 2026: AI compute](#)
- [SEC Chair Atkins, remarks on crypto asset regulation \(Mar 17, 2026\)](#)
- [Prime Intellect \\$130M Series A – TechCrunch \(Jul 2026\) / Prime Intellect Series A blog](#) (revenue, customer count)
- [CoinGecko: Bittensor / CoinGecko: Render](#) – July 2026 price/mcap corrections
- [io.net IDE independent burn analysis / CoinDesk io.net IDE press release](#)
- [Bittensor Robin τ subnet expansion / TAO ETFs and institutional bets](#)
- [Gensyn mainnet and \\$AI token / BitTorrent BTTInferGrid launch](#)
- [NVIDIA Vera Rubin enters production / CNBC on Kyber rack delay](#)
- [Thunder Compute GPU rental market trends \(Jul 2026\)](#)
- [Futurum: AI capex 2026 – the \\$690B infrastructure sprint](#)
- [SEC July 2026 token-sale rulemaking plans / Ledger Insights on Atkins exemptions](#)
- [Goldman Sachs: AI agents forecast to boost token usage](#)
- [2026 memory supercycle trajectory](#)
- [Nosana-OpenGPU integration and relaunch / Aethir 2025 wrap-up](#)

Research compiled from 10 initial parallel investigation agents and 5 deep-dive research agents covering academic papers, industry reports, and empirical data from 2024-2026. v2.0 updated February 2026; v2.1 verified against live sources July 2026; v2.2-v2.5 re-verified July 18, 2026.