

Gonka Product-Led Growth Model

Version: 1.3 **Date:** 2026-07-18 **Classification:** Internal – actionable growth model for Gonka developer acquisition
Dependencies: Phase 15 (competitive analysis), Phase 16 (developer personas), Phase 17 (messaging), Phase 18 (channel strategy), Phase 19 (partnership playbook) **Requirement:** GTM-03

Executive Summary

This document synthesizes all Phase 15-19 research into a single actionable growth model for Gonka’s developer adoption. It defines the PLG funnel (how developers move from discovery to payment), the free tier design spec (what developers get without paying), and the time-to-first-inference plan (how developers go from zero to their first API response in under 5 minutes).

Every conversion rate, usage limit, and design decision is grounded in evidence from the research corpus. Where data is indirect (industry benchmarks rather than Gonka-specific measurements), confidence levels are noted. The model is designed so that leadership can read this document and know exactly what Gonka must build, what the targets are, and how to measure progress.

All developer-facing copy in this document follows the Phase 17 vocabulary guidelines (gonka_message_house.md Section 6). No terms from the 16-item never-say list appear in any developer-facing recommendation.

July 2026 revision note: This revision re-bases the model on the current Moonshot lineup – Kimi K2.5 has been superseded by K2.6 (April 2026) and K2.7-Code (June 2026); K3 launched July 16, 2026, and Moonshot has closed K2.5 to new users, is redirecting K2.5 API traffic to K2.6, and will sunset K2.5 platform-wide on August 31, 2026 – any K2.5-anchored comparison has ~6 weeks of shelf life. It also corrects the Kimi context window (256K, not 131K), refreshes competitor free-tier and pricing facts, flags that prompt caching is now universal across compared providers (which shrinks the session-persistence savings deltas computed in April), and adds ClawHub security and x402 agent-payment context. **Pricing re-base:** the April 2026 Scenario B rates (\$0.35/\$1.75 per 1M) are now fully obsolete, not merely thin. Gonka’s live network rate is ~\$0.0003/1M tokens (per-block dynamic pricing; pricepertoken.com/endpoints/gonka). Broker retail via Gonka24 is now model-specific: Kimi K2.6 \$0.055/\$0.32 per 1M, GLM-5.2 \$0.095/\$0.30, MiniMax M2.7 from \$0.018/\$0.072 – the widely quoted “\$0.018/\$0.072” is the M2.7 floor rate only, not a flat Gonka rate. All dollar figures inherited from gonka_agent_pricing_analysis.md that still assume Scenario B are marked obsolete below; revenue projections must be rebuilt on the live per-block pricing mechanism plus a normalized-rate band that accounts for current utilization subsidy.

PLG Funnel Model

Funnel Overview

The PLG funnel maps Gonka’s developer acquisition to the AAARRRP framework from Phase 16 (gonka_developer_personas.md). Six stages capture the complete journey from first encounter to paying customer, with each stage producing a measurable outcome and a clear handoff to the next.

Stage	AAARRRP	Definition	Target Rate	Metric
Discover	Awareness	Developer encounters Gonka through OpenClaw ecosystem channels, Reddit, GitHub, or technical content	– (top of funnel)	Unique visitors to docs.gonka.ai
Explore	Acquisition	Developer visits documentation, reads quickstart guide, reviews pricing page	40-50% of Discover	Page views on docs.gonka.ai; time on site > 2 minutes

Stage	AAARRRP	Definition	Target Rate	Metric
Sign Up	Acquisition	Email-only API key creation – no wallet, no OAuth, no credit card	15-25% of Explore	API keys created per week
First Inference	Activation	First successful API call through Gonka network	70-80% of Sign Up	First API call within 24 hours of signup
Habitual Use	Retention	100+ API calls in 30 days – developer has integrated Gonka into a real workflow	20-30% of First Inference	API-active developers per month (gonka_channel_strategy.md primary KPI)
Paid Conversion	Revenue	Developer exceeds free tier limits and upgrades to paid plan	5-10% of Habitual Use	Paying users; monthly revenue per user

Why these rates: The target conversion rates are calibrated against the January 2026 ChartMogul/ProductLed/Growth Unhinged free-to-paid conversion study (200 B2B products, conversion measured within 6 months): overall median free-to-paid is 8%; opt-in no-credit-card trials convert at ~8.9% median; credit-card-required trials at ~30-31.4% (more than 5x no-card trials); freemium products rarely convert above 15%, and developer-tool freemium benchmarks sit at 2-4%. Top-performing PLG companies target 40-60% activation (best-in-class 70%+) – an April 2026 industry estimate we have not re-verified. Gonka's 70-80% signup-to-first-inference target is aggressive but defensible because the activation step is a single config paste, not an SDK integration. The 5-10% paid conversion rate reflects the typical PLG pattern where most value is delivered free and only production-scale workloads trigger payment (gonka_developer_personas.md: Weekend Builder stays on free tier; Startup CTO converts to paid).

Per-Persona Conversion Drivers

Each funnel stage has a different primary driver depending on the developer persona. Understanding these drivers determines what content, features, and messaging to prioritize at each stage.

Stage 1: Discover

Persona	Primary Conversion Driver	#1 Drop-off Risk	Mitigation
Weekend Builder	Reddit r/LocalLLaMA cost comparison post showing Gonka as the cheapest tracker-listed provider for K2.6/K2.7 (pricepertoken.com-verifiable; note the current network rate reflects low utilization and is effectively subsidized – do not quote the April 2026 “\$40/mo to \$13/mo” Scenario B figures)	“Never heard of Gonka – is this a crypto scam?” Landing page looks like a DeFi protocol	Landing page must resemble Vercel/Supabase, not a token project. Lead with “Cut your agent costs by 73%” not “Decentralized AI inference” (gonka_message_house.md VP1; PITFALLS.md Pitfall 1)
Startup CTO	Hacker News technical deep-dive on agent heartbeat overhead problem with Gonka cited as the session persistence solution	“Decentralized infrastructure means unreliable. Akash and Render both failed at scale”	Technical white paper with uptime data and p95 latency benchmarks. Must read like a Cloudflare blog post, not a crypto whitepaper (gonka_developer_personas.md Startup CTO Awareness stage)
Privacy-First Builder	GitHub repository README or technical article on privacy-preserving inference architectures, found while researching alternatives to self-hosted vLLM	“Every cloud provider claims no logging. How is Gonka different architecturally?”	Honest capabilities page: current guarantees (no central log aggregation, open-weight model, no content filtering) vs roadmap (TEE-based encrypted inference) (gonka_developer_personas.md Privacy-First Awareness stage)

Source: gonka_developer_personas.md AAARRRP Journey Maps, Awareness stage for each persona.

Stage 2: Explore

Persona	Primary Conversion Driver	#1 Drop-off Risk	Mitigation
Weekend Builder	Quickstart page showing the 3-field openclaw.json snippet with a copy button – “Add Gonka in 90 seconds”	Page mentions tokens, staking, or wallet before showing the config snippet	Docs homepage leads with the config snippet, not network architecture. Follow never-say list (gonka_message_house.md Section 6.1)
Startup CTO	Architecture overview explaining reliability through redundant nodes and health-checked routing, plus a cost comparison calculator	Docs lack uptime data, SLA information, or team account management	Enterprise evaluation page with p95 latency data, multi-key team accounts, and architecture explanation using cloud infrastructure language (gonka_developer_personas.md Startup CTO Acquisition stage)
Privacy-First Builder	Data handling policy page and open-source infrastructure code link (github.com/mitgor/gonka-ai-infrastructure) – reviewable before signup	Privacy documentation is vague or makes claims without architectural backing	Published data handling policy: what is logged, what is not, retention periods, node operator agreements. Link to open-source code for independent audit (gonka_developer_personas.md Privacy-First Acquisition stage)

Source: gonka_developer_personas.md AAARRRP Journey Maps, Acquisition stage for each persona.

Stage 3: Sign Up

Persona	Primary Conversion Driver	#1 Drop-off Risk	Mitigation
Weekend Builder	Single-CTA “Get API Key” button leading to email-only form. API key displayed immediately after email verification. Zero friction	Form asks for company name, use case, or any information beyond email address	Minimal form: email only. Verification link sent instantly. API key displayed on verification page with one-click copy. No approval queue, no waitlist (CONTEXT.md locked decision: email-only signup)
Startup CTO	Team account creation with multiple API keys for dev/staging/prod environments. Architecture docs reviewable pre-signup	Signup requires individual accounts for each team member; no team management	Team signup flow: one email creates team account, generates scoped API keys per environment. Billing attached to team, not individual (gonka_developer_personas.md Startup CTO Acquisition stage)
Privacy-First Builder	Data handling policy reviewed and satisfactory before email submission. Open-source code independently auditable	Privacy policy is generic boilerplate that does not address inference-specific concerns (prompt logging, content filtering, data residency)	Inference-specific privacy policy linked prominently on signup page. “We do not log your prompts” with architectural explanation, not just a policy claim (gonka_developer_personas.md Privacy-First Acquisition stage)

Source: gonka_developer_personas.md AAARRRP Journey Maps; CONTEXT.md decisions.

Stage 4: First Inference

Persona	Primary Conversion Driver	#1 Drop-off Risk	Mitigation
Weekend Builder	Pre-built openclaw.json snippet on docs page with copy button. “Add Gonka to OpenClaw in 90 seconds” tutorial	The two-step provider gotcha: developer defines provider but forgets to add model to agent allowlist, causing silent failure	Quickstart guide explicitly shows both steps (provider definition AND model allowlisting) with a troubleshooting section for the

Persona	Primary Conversion Driver	#1 Drop-off Risk	Mitigation
			gotcha (STACK.md Section 1: provider selection is a TWO-STEP process)
Startup CTO	Integration test results from v1.2 showing OpenClaw, CrewAI, and LangGraph compatibility. Staging deployment guide	First test request times out or returns unexpected error, destroying trust at first contact	Sub-200ms p95 latency on first call. Clear error messages with actionable troubleshooting. Integration test suite available for the team to run against their staging environment (gonka_developer_personas.md Startup CTO Activation stage; PITFALLS.md Pitfall 3)
Privacy-First Builder	Test prompt that would be blocked by OpenAI/Anthropic content filters completes successfully on Gonka. Aha moment: "It processed my prompt without filtering"	First prompt triggers an unexpected error or content filter, contradicting "no content policy" claim	The served Kimi models (K2.6, K2.7-Code – Modified MIT, open-weight) have no built-in content filtering. Verify no content filtering middleware in the gateway. Test with adversarial prompts during QA (gonka_competitive_feature_matrix.md: Content filtering – Gonka "None (open)")

Source: gonka_developer_personas.md AAARRRP Journey Maps, Activation stage; STACK.md.

Stage 5: Habitual Use

Persona	Primary Conversion Driver	#1 Drop-off Risk	Mitigation
Weekend Builder	Usage dashboard showing actual spend vs the cheapest listed competitor route, cached-adjusted (retire the April 2026 "\$40 vs \$13" Scenario B framing; at live network rates the dashboard delta is dramatic but must carry the utilization/subsidy caveat). Cost comparison is visceral and shareable	After initial excitement, the developer forgets about Gonka and defaults back to OpenRouter because it is built-in	Weekly email digest showing cost savings. X-Gonka-Usage-Remaining header in every response keeps savings visible. Cost comparison calculator in dashboard (gonka_developer_personas.md Weekend Builder Retention stage)
Startup CTO	Weekly cost report showing session persistence savings across all agents. Gradual migration from 2 agents to 4 to all 6 as confidence builds	One timeout or error during production usage triggers revert to previous provider. First impressions are permanent for infrastructure trust	Real-time status page (status.gonka.ai) with per-model availability and historical uptime. Incident report publication process. Gradual migration path with dual-provider configuration (gonka_developer_personas.md Startup CTO Retention stage; PITFALLS.md Pitfall 3)
Privacy-First Builder	Ongoing use for sensitive workloads with periodic privacy posture auditing via open-source verification tools	Gonka updates terms of service or infrastructure in a way that weakens privacy guarantees without communicating the change	Privacy-specific changelog (separate from feature releases). Automated privacy audit tool. Community-maintained watchdog (gonka_developer_personas.md Privacy-First Retention stage)

Source: gonka_developer_personas.md AAARRRP Journey Maps, Retention stage.

Stage 6: Paid Conversion

Persona	Primary Conversion Driver	#1 Drop-off Risk	Mitigation
Weekend Builder	Exceeds free tier credit (\$5-10 trial). Clear pricing page showing per-token rates with	Fear of runaway costs – "OpenClaw agents can spike unpredictably, \$300+ in 2	Configurable daily spending cap. Usage alerts at 50%, 80%, 100% of limit. "No

Persona	Primary Conversion Driver	#1 Drop-off Risk	Mitigation
	session savings calculator. Credit card entry with spending cap	days.” This fear is validated at industry scale: a 2026 review of 127 enterprise agentic-AI implementations found 73% went over budget (some by 2.4x+), and blended enterprise cost per 1M tokens fell 67% YoY (\$18.40 to \$6.07, Q1 2025 to Q1 2026); the FinOps Foundation’s State of FinOps 2026 survey separately reports 98% of FinOps teams now manage AI spend (up from 31% in 2024)	surprise bills” guarantee. Pricing comparison table: Gonka vs OpenRouter vs Together AI at Casual tier (gonka_developer_personas.md Weekend Builder Revenue stage)
Startup CTO	Commits to paid plan after 2-4 weeks of staging + gradual production rollout. Volume pricing for 161M+ tokens/month	Needs formal SLA with downtime credits to justify to co-founder and investors	Enterprise pricing tier with committed-use discounts. SLA framework: 99.5% uptime target with credit mechanism. Invoice and receipt system for accounting (gonka_developer_personas.md Startup CTO Revenue stage; gonka_competitive_feature_matrix.md Section 7)
Privacy-First Builder	Converts to paid plan valuing session persistence for reducing how often sensitive context traverses the network. Willing to pay premium for TEE when available	No privacy-specific pricing tier or TEE timeline commitment	Privacy tier with optional TEE-backed inference at premium when available. Roadmap commitment on encrypted inference with quarterly updates. Early access program for privacy beta features (gonka_developer_personas.md Privacy-First Revenue stage)

Source: gonka_developer_personas.md AAARRRP Journey Maps, Revenue stage; gonka_agent_pricing_analysis.md.

Funnel Math

If 1,000 developers discover Gonka (visit docs.gonka.ai or encounter Gonka in the OpenClaw ecosystem), here is how many become paying users at each scenario:

Stage	Pessimistic	Target	Optimistic
Discover	1,000	1,000	1,000
Explore (40-50% of Discover)	400	450	500
Sign Up (15-25% of Explore)	60	90	125
First Inference (70-80% of Sign Up)	42	67	100
Habitual Use (20-30% of First Inference)	8	17	30
Paid Conversion (5-10% of Habitual Use)	0.4 (~0)	1.3 (~1)	3

Interpretation:

- **Pessimistic:** ~0.04% end-to-end conversion. From 1,000 developers who discover Gonka, zero to one becomes a paying customer. This scenario reflects a cold start with no ecosystem presence, minimal content, and the crypto perception barrier fully active. Recovery requires aggressive investment in P0 channels (OpenClaw Provider Directory, GitHub, Discord – gonka_channel_strategy.md; the official OpenClaw Discord stands at ~175K members as of July 2026, and OpenClaw hit 247K GitHub stars and 47.7K forks by March 2, 2026 – the fastest-growing repo in GitHub history – validating these as P0 channels. A fork ecosystem is emerging, including sandboxed-execution security forks, Chinese adaptations wired to DeepSeek and WeChat,

and OpenClaw-based commercial services from Tencent and Z.ai – notable because Z.ai’s GLM-5.2 is on Gonka’s network).

- **Target:** ~0.13% end-to-end conversion. From 1,000 discovers, approximately 1 paying user and 17 API-active developers generating usage data, cost savings testimonials, and referral potential. At the Month 6 target of 200 API-active developers (gonka_channel_strategy.md), this implies needing approximately 12,000 discover-stage developers.
- **Optimistic:** ~0.3% end-to-end conversion. From 1,000 discovers, 3 paying users and 30 API-active developers. Achievable if Gonka secures built-in OpenClaw provider status (Tier 3 in gonka_partnership_playbook.md), which eliminates the manual configuration friction at the Sign Up and First Inference stages.

Key insight: The funnel is widest at the top (Discover to Explore) and steepest at the Habitual Use stage. The 20-30% conversion from First Inference to Habitual Use is the critical gate – this is where developers decide if Gonka’s session persistence savings are real enough to justify switching from their default provider. The X-Gonka-Usage-Remaining header and cost comparison dashboard are the primary retention mechanisms at this stage.

Revenue projection at target rates (obsolete – rebuild required):

The April 2026 projection (~\$733/month at Month 6 from a 1,000-discover cohort) was computed at Scenario B pricing (\$0.35/\$1.75 per 1M). Scenario B is now obsolete: Gonka’s live network rate is ~\$0.0003/1M tokens under the per-block dynamic pricing mechanism, and broker retail on Gonka24 is model-specific (Kimi K2.6 \$0.055/\$0.32 per 1M; \$0.018/\$0.072 is the MiniMax M2.7 floor rate only). At those rates:

- A Casual-tier developer (~31.5M tokens/month) costs under one cent per month at the network rate – Scenario B overstated near-term revenue-per-developer by roughly 2-3 orders of magnitude.
- Near-term inference revenue from this cohort is therefore negligible; the current rate reflects low network utilization and is effectively a supply-side subsidy, not a sustainable price.

The projection must be rebuilt on the live per-block pricing mechanism plus a normalized-rate band (what the rate converges to at target utilization). Until that re-run exists, treat funnel volume metrics (API-active developers), not revenue, as the Month 6 success measure. Revenue scales with discover-stage volume, which is driven by P0 channel investment (gonka_channel_strategy.md).

Free Tier Design Spec

Design Principles

The free tier must accomplish three goals simultaneously:

1. **Eliminate all friction for the Weekend Builder** – generous enough that a solo developer running one agent on one channel can use Gonka indefinitely without paying, experiencing session persistence savings firsthand and becoming a referral source (gonka_developer_personas.md: Weekend Builder monthly spend \$20-50 on OpenRouter). Note (July 2026): at the live network rate (~\$0.0003/1M) a Casual month costs under a cent, so “must the free tier cover a Casual agent” is currently moot on cost grounds – the limits below matter as abuse control and as the upgrade-trigger structure for when rates normalize.
2. **Create a natural upgrade trigger for the Startup CTO** – limited enough that a team running multiple agents across channels hits the ceiling within their first week of serious evaluation, triggering the paid conversion conversation (gonka_developer_personas.md: Startup CTO monthly spend \$200-500; free tier should cover approximately 1 agent, not 6)
3. **Remove the crypto barrier entirely** – Email-only signup with no wallet, no token purchase, no OAuth complexity. The word “wallet” must never appear before “API key” on any page (gonka_message_house.md Section 6.1: Never-Say List – “wallet” is the single most alienating word for Web2 developers)

Scope note (July 2026): the “API keys only, never crypto” rule applies to the human Web2 personas above. For the agent-as-customer segment (gonka_agent_native_pitch.md), agent-native payment rails went mainstream in 2026: the x402 Foundation reached operational launch under the Linux Foundation on July 14, 2026 with 40 member organizations – 17 premier members: Adyen, AWS, American Express, Circle, Cloudflare, Coinbase, Fiserv, Google, Mastercard, Monad Foundation, MoonPay, Ripple, Shopify, Solana Foundation, Stellar Development Foundation, Stripe, and Visa (general members include Polygon Labs, NEAR Foundation, Fireblocks, and KakaoPay). In the 30 days ending ~July 15, 2026 the protocol processed ~75M transactions moving ~\$24M (~\$800K/day) between ~94K buyers and ~22K sellers, predominantly settled in USDC, with 95% of payment value now in transactions above \$1 (Chainalysis) – the earlier “mostly testing/gamed micropayments” caveat no longer holds. BlockRunAI’s ClawRouter already authenticates OpenClaw agents with wallet signatures and pays for inference via x402 USDC

micropayments. Gonka – already a crypto network – has a natural x402 story here and a competitor occupying the position; this deserves a separate track that never leaks into the human onboarding flow.

Signup Flow

- Step 1: Developer lands on docs.gonka.ai
Single CTA: "Get Your API Key" button
No distracting navigation, no "Learn about our network" sidebar
- Step 2: Email-only form
One field: email address
No company name, no use case, no phone number
Submit button: "Send me my API key"
- Step 3: Verification email arrives (target: under 60 seconds)
Transactional email service (SendGrid, Postmark, or SES)
Subject: "Your Gonka API Key"
Body: One-click verification link, nothing else
- Step 4: Developer clicks link, lands on key display page
API key displayed prominently with one-click copy button
Below the key: the exact openclaw.json snippet pre-filled with their key
Below the snippet: "Paste this into ~/.openclaw/openclaw.json and you are done"
- Step 5: Developer is using Gonka
No approval queue. No waitlist. No manual review.
Free tier limits active immediately.

Per CONTEXT.md locked decision: Email-only signup. No wallet or crypto knowledge required.

Usage Limits

Free tier limits are calibrated against two data points:

- Competitor free tiers:** OpenRouter offers free-tier access with rate limiting and quality degradation under load. Google cut Gemini API free-tier quotas 50-80% starting December 2025 and removed Pro-series models from the free tier in April 2026 (Flash/Flash-Lite only, ~250-1,000 requests/day); Google also now imposes a mandatory monthly spending cap that auto-pauses the API, and the free tier disappears once billing is enabled. No Kimi-family free tier exists among compared providers (gonka_competitive_feature_matrix.md; gonka_provider_landscape_map.md Segment 3: Together AI and DeepInfra are pay-per-use only).
- Persona workload profiles:** The Weekend Builder consumes ~31.5M tokens/month at Casual tier with 30-minute heartbeats (gonka_agent_pricing_analysis.md Section 3, April 2026 estimate). With Gonka's session persistence reducing effective token consumption by ~80% (estimate; competitors' universal prompt caching now captures part of this delta), the Weekend Builder's actual Gonka token usage would be approximately 12.6M tokens/month. The free tier should cover most of this.

Limit	Free Tier Value	Rationale
Daily requests	1,000 requests/day	Covers 1 agent with 50 user messages + 150 tool calls + 48 heartbeats = ~248 requests/day. Generous buffer for burst testing. Startup CTO with 6 channel-agents needs ~1,488 requests/day – exceeds limit within first day of serious evaluation (gonka_agent_pricing_analysis.md Section 3: Active tier daily breakdown)
Monthly tokens	15,000,000 tokens/month	Covers Weekend Builder's session-optimized usage (~12.6M tokens/month) with headroom. Startup CTO at Active tier (~56.3M session-optimized tokens/month) exceeds this in week 1 (gonka_agent_pricing_analysis.md Section 5: Gonka hidden cost analysis)
Concurrent sessions	2 active sessions	Covers 1 agent on 1-2 channels (Weekend Builder). Startup CTO with 6 channel-agents needs 6 concurrent

Limit	Free Tier Value	Rationale
		sessions – hits upgrade trigger immediately when deploying team workload
Models available	K2.6 lite and mid tiers only	Full-precision K2.6 (and a K3 frontier tier – K3 launched July 16, 2026; open-weight availability on Gonka to be confirmed) reserved for paid tier. Lite and mid tiers are sufficient for the Weekend Builder's personal assistant workload. Startup CTO evaluating quality for production needs full tier access (gonka_competitive_feature_matrix.md 3-tier quantization concept, re-based from K2.5 to K2.6)
Rate limit	60 requests/minute	Prevents abuse while allowing burst testing (a developer testing their agent configuration may send 10-20 requests in quick succession). Production agents with 30-minute heartbeats generate ~1 request/minute baseline. 60 RPM provides comfortable headroom
Context window	32,768 tokens	Reduced from the full 262,144 (Kimi K2.6 and K2.7-Code both ship 256K contexts). Sufficient for most agent workloads (system prompt + context is ~9,600 tokens). Startup CTO running complex multi-step workflows with deep context needs the full window

Upgrade Triggers

Three specific mechanisms notify developers when they approach or exceed free tier limits, driving paid conversion without hard-blocking their workflow:

Trigger 1: Usage Limit Approaching

When a developer reaches 80% of their daily request limit or monthly token limit, every API response includes a warning header:

```
X-Gonka-Usage-Remaining: 200
X-Gonka-Usage-Limit: 1000
X-Gonka-Usage-Reset: 2026-07-19T00:00:00Z
X-Gonka-Upgrade-URL: https://docs.gonka.ai/pricing
```

At 100%, requests return HTTP 429 with a clear message: "Free tier daily limit reached. Resets at midnight UTC. Upgrade at docs.gonka.ai/pricing for unlimited requests."

No cryptic error codes. No silent failures. The developer knows exactly what happened, when it resets, and how to fix it.

Trigger 2: Session Limit Reached

When a developer attempts to create a third concurrent session (free tier allows 2), the API returns HTTP 403 with:

```
{
  "error": "session_limit_reached",
  "message": "Free tier supports 2 concurrent sessions. You have 2 active sessions. Upgrade to paid for unlimited sessions.",
  "current_sessions": 2,
  "max_sessions": 2,
  "upgrade_url": "https://docs.gonka.ai/pricing"
}
```

This trigger is designed specifically for the Startup CTO persona: they have proven Gonka works with 1-2 agents on staging and now want to add their production agents. The session limit creates a natural conversation point for team account and volume pricing discussion.

Trigger 3: Model Tier Access

When a developer requests the full-precision model (available only on paid tier), the API routes to the mid tier and includes a header:

```
X-Gonka-Tier-Served: mid
X-Gonka-Tier-Requested: full
X-Gonka-Tier-Upgrade: "Full-precision models available on paid tier. See docs.gonka.ai/pricing"
```

The request succeeds (no error) but with a lower-tier model, and the header transparently communicates the substitution. This avoids breaking the developer's workflow while making the quality difference discoverable.

What Is NOT in the Free Tier

Features reserved for paid plans to create clear value differentiation:

Feature	Free Tier	Paid Tier	Why Reserved
Full-precision K2.6 / K3 frontier tier	No (lite and mid only)	Yes	Quality differentiation for production workloads
Unlimited sessions	2 concurrent max	Unlimited	Multi-agent production requires paid commitment
SLA guarantee	Best effort	99.5% uptime target with credits	SLA has cost implications; reserved for committed customers
Priority routing	Standard queue	Priority queue with dedicated capacity	Ensures paid customers get consistent low latency
Webhook notifications	No	Yes	Advanced agent feature for production use (gonka_competitive_feature_matrix.md: Tier 2 differentiator)
Memory API	No	Yes	Persistent key-value memory for long-running agents
Usage API / Dashboard	Basic (current usage only)	Full (historical trends, cost comparison, export)	Detailed analytics drive retention for paid users
Team accounts	No (individual only)	Yes (multi-key, scoped permissions)	Team management is an enterprise feature
Spending caps	N/A (free)	Configurable daily/monthly limits	Cost control for production workloads (addresses "\$300+ in 2 days" fear – gonka_developer_personas.md Weekend Builder Revenue stage)

Competitive Positioning of Free Tier

Provider	Free Tier	Signup Friction	Limits
OpenRouter	Free-tier access to select models	GitHub OAuth	Rate-limited; quality degraded under load; free requests queued behind paid (gonka_competitive_feature_matrix.md Section 5)
Google Gemini	Flash/Flash-Lite only (Pro removed April 2026)	Google account	Quotas cut 50-80% since Dec 2025; ~250-1,000 requests/day; free tier disappears once billing is enabled
Together AI	No free tier (signup credits retired July 2025)	\$5 minimum credit purchase	Pay-per-use only (gonka_provider_landscape_map.md Segment 3); Startup Accelerator credits exist for qualifying startups
DeepInfra	No free tier	Credit card required	Pay-per-use only
Gonka (proposed)	15M tokens/month, 1,000 requests/day, 2 sessions	Email only	Most generous free tier among Kimi-family providers; lowest signup friction; session persistence active on free tier

Competitive advantage: Gonka's free tier is the only one that (a) requires no credit card or OAuth, (b) includes session persistence (the primary cost-saving feature), and (c) provides enough capacity for a solo developer's always-on agent. The Weekend Builder can experience the full session persistence savings without paying, creating the cost comparison data point ("what I paid on OpenRouter vs free on Gonka" – use the developer's own tracked spend, not the retired April 2026 estimates) that drives organic referrals (gonka_developer_personas.md Weekend Builder Referral stage).

Copy Guidelines

All developer-facing copy on signup pages, error messages, and upgrade prompts must follow the Phase 17 vocabulary guidelines (gonka_message_house.md Section 6). Specifically:

- **Never use:** wallet, staking, mining, DePIN, Web3, tokenomics, consensus mechanism, on-chain, smart contract, gas, slashing, governance proposals, epochs, validators, token (as cryptocurrency), decentralized (as headline)
- **Always use:** "API key" (not "wallet"), "account" (not "wallet"), "compute credits" (not "GNK tokens"), "GPU providers" (not "miners" or "validators")
- **Error messages:** Must be actionable and specific. "Free tier daily limit reached. Resets at midnight UTC." not "Insufficient GNK balance."

Time-to-First-Inference Plan

Objective

Per CONTEXT.md locked decision: target under 5 minutes from email submission to first successful API response through the Gonka network. This is the single most important PLG metric – every minute of friction between "I want to try Gonka" and "my agent is running on Gonka" loses developers who will never return (PITFALLS.md Pitfall 3: first impressions are permanent for infrastructure trust).

Step-by-Step Breakdown

Step	Action	Time Budget	Blocker Risk	What Must Be Built
1	Land on docs.gonka.ai	0s (assumed)	Page must exist. Currently no docs site – must-close gap from Phase 15 (gonka_provider_landscape.md 5 critical gaps)	docs.gonka.ai with SSL, CDN, sub-2s page load. Single-page-app or static site (Hugo, Astro, or Next.js). Landing page with one CTA
2	Click "Get API Key"	10s	Single CTA must be above the fold. No competing navigation, no "Learn about our network" distractions. Developer must not need to scroll to find the button	CTA button component. Landing page design that passes the "is this crypto?" test (gonka_message_house.md VP3; PITFALLS.md Pitfall 1)
3	Enter email, submit	30s	Email-only form. No wallet creation, no OAuth redirect, no credit card. One field, one button. Any additional friction at this step causes the steepest drop-off in the funnel (gonka_developer_personas.md: 95% funnel drop at wallet step per PITFALLS.md Pitfall 2)	Email submission endpoint. Rate limiting on submissions (prevent abuse). CAPTCHA only if abuse is detected, not by default
4	Receive verification email	60s	Must arrive within 60 seconds. Delayed verification emails (>2 minutes) cause	Email sending infrastructure. Transactional email

Step	Action	Time Budget	Blocker Risk	What Must Be Built
			developers to lose context and abandon. Use a transactional email service (SendGrid, Postmark, or AWS SES) with dedicated IP and domain authentication (SPF, DKIM, DMARC) to avoid spam folders	template. Email deliverability monitoring
5	Click verification link, see API key	15s	Key must be displayed immediately on the page with a one-click copy button. No “check your dashboard” redirect. No additional login step. The key is shown once, prominently, with a “Copy” button that provides visual confirmation	Key generation service. Key display page with clipboard API integration. Key storage (hashed) in database
6	Paste config into openclaw.json	90s	The exact config snippet must be pre-filled with the developer’s API key and ready to copy-paste. Show both steps of the two-step process (provider definition AND model allowlisting) to avoid the silent failure gotcha (STACK.md Section 1)	Pre-filled config snippet generator. Documentation page showing the exact JSON to paste
7	Send first message to agent	60s	Agent responds via Gonka inference. The first response must arrive within 2-3 seconds (sub-200ms TTFT + streaming). Any timeout or error at this step permanently damages trust	Live API endpoint at api.gonka.ai. Health-checked GPU hosts. vLLM inference with K2.6 model loaded and warm
Total		~4 min 15s		

Step 6 Detail: The Config Snippet

After the developer copies their API key, the verification page displays the exact `openclaw.json` configuration they need to paste. This is not a generic template – it is pre-filled with their specific API key:

```
// Step 1: Add this to your ~/.openclaw/openclaw.json under models.providers:
{
  "gonka": {
    "baseUrl": "https://api.gonka.ai/v1",
    "apiKey": "gnk_sk_abc123...", // Your API key (already filled in)
    "api": "openai-completions",
    "models": [
      {
        "id": "kimi-k2.6",
        "name": "Kimi K2.6",
        "contextWindow": 262144,
        "maxTokens": 8192,
        "cost": { "input": 0.055, "output": 0.32 } // K2.6 Gonka24 retail; set from live rate at
        publish time
      }
    ]
  }
}
```

```

    }
  ]
}
}

// Step 2: Add this to your agents.defaults.models to enable the model:
{
  "gonka/kimi-k2.6": { "alias": "k2.6" }
}

```

Why both steps are shown: OpenClaw requires both provider definition (Step 1) AND model allowlisting (Step 2). Missing Step 2 causes silent failure – the agent ignores the provider entirely with no error message. This is a known gotcha in the OpenClaw community (STACK.md Section 1: “Missing either step causes silent failure”). Showing both steps on the same page with clear “Step 1” and “Step 2” labels prevents the most common configuration error.

Config snippet format: Based on the OpenClaw `openclaw.json` configuration format documented in STACK.md. The `api: "openai-completions"` value tells OpenClaw to use the standard `/v1/chat/completions` endpoint format, which Gonka’s gateway already implements (shipped in v1.2). Cost values shown use the July 2026 Gonka24 retail rate for Kimi K2.6 (\$0.055/\$0.32 per 1M; the raw network rate is ~\$0.0003/1M, and \$0.018/\$0.072 applies only to MiniMax M2.7) and must be regenerated from the live rate at publish time – the April 2026 Scenario B values (\$0.35/\$1.75) are obsolete. Competitor reference points: Together AI K2.6 \$1.20/\$4.50, K2.7-Code \$0.95/\$4.00. (OpenRouter still lists K2.5 at \$0.375/\$2.025 and DeepInfra at \$0.45/\$2.25, but Moonshot is sunsetting K2.5 on August 31, 2026 – do not anchor comparisons on K2.5 rates.)

Current Gaps vs Target Flow

Step	Current Status	Gap	Engineering Work Required	Cross-Reference
1	No docs site exists	Must-close gap	Static site deployment (Hugo/Astro/Next.js) with domain, SSL, CDN	gonka_provider_landscape_map.md: critical gap #2
2	No landing page exists	Must-close gap	Landing page design and implementation	gonka_provider_landscape_map.md: critical gap #2
3	No signup flow exists	Must-close gap	Email submission API, rate limiting, form validation	gonka_provider_landscape_map.md: critical gap #3
4	No email infrastructure	Must-close gap	Transactional email service integration, domain auth	New requirement (identified in this plan)
5	No key generation service	Must-close gap	API key generation, storage, display page	gonka_provider_landscape_map.md: critical gap #3
6	Config snippet documented in STACK.md (research only)	Needs productization	Pre-filled snippet generator on docs site	gonka_partnership_playbook.md: requirement #13
7	API endpoint exists (v1.2) but not publicly accessible	Must-close gap	Public deployment of api.gonka.ai with domain, SSL, health checks	gonka_provider_landscape_map.md: critical gap #1

Key observation: Steps 1-6 are entirely new engineering work. Only Step 7 (the actual inference) has existing infrastructure from v1.2. The entire developer onboarding flow – from landing page to API key delivery – must be built from scratch. This is why the `gonka_provider_landscape_map.md` identifies these as “not feature gaps but table-stakes infrastructure gaps that prevent developers from evaluating Gonka at all.”

The Atomic Growth Unit

The `openclaw.json` config snippet is Gonka’s atomic growth unit – the smallest shareable artifact that converts a non-user into a user.

Why the config snippet is viral:

1. **It is copy-pasteable.** A developer who has configured Gonka can share their provider config block (minus the API key) in a Reddit comment, Discord message, or GitHub issue. The recipient copies it, adds their own API

key, and is running on Gonka in 90 seconds. No documentation reading required. No signup flow navigation. Just paste and go.

2. **It lives in version control.** `openclaw.json` is committed to Git repositories. When a developer adds Gonka as a provider, their teammates see the config change in the next pull request. The config snippet propagates through team workflows organically. When the Startup CTO adds Gonka to the team's shared config, all 3-8 engineers inherit it automatically (`gonka_developer_personas.md`: Startup CTO team size 3-8).
3. **It is portable across OpenClaw installations.** The same config snippet works on any OpenClaw deployment – personal laptop, VPS, cloud server, Docker container. A developer who configures Gonka on their local machine can copy the same config to production. This portability means the snippet carries across environments without modification (`STACK.md`: OpenClaw's declarative configuration model).
4. **It embeds Gonka's cost advantage.** The `cost` field in the config snippet shows Gonka's per-token rates directly in the developer's configuration file. Every time a developer opens `openclaw.json`, they see Gonka's rates next to their other providers. At July 2026 rates the gap is unmissable: Gonka24 K2.6 retail \$0.055/\$0.32 per 1M (network rate ~\$0.0003/1M) vs Together AI's \$1.20/\$4.50 (K2.6) and OpenRouter's cheapest Kimi route at \$0.375/\$2.025 (K2.5 – being sunset by Moonshot August 31, 2026) – Gonka is the cheapest tracker-listed provider for K2.6 by roughly an order of magnitude. Frame external claims as “cheapest listed provider for K2.6/K2.7” with the caveat that the current rate reflects low utilization and is effectively subsidized.
5. **ClawHub skills amplify the snippet – with a security caveat.** The Gonka Provider Skill submitted to ClawHub (`gonka_partnership_playbook.md`: ClawHub Submission Plan) teaches agents how to use Gonka-specific features. When a developer installs the skill, it includes the config snippet as part of the setup instructions. The skill and the snippet work together: the snippet connects the agent to Gonka, and the skill teaches the agent to use sessions, tiering, and memory (`gonka_partnership_playbook.md`: SKILL.md draft outline). **Caveat (July 2026):** the early-2026 ClawHavoc supply-chain campaign seeded ClawHub with malicious skills delivering crypto-wallet-stealing malware (Atomic macOS Stealer), with ~\$2.3M in cryptocurrency reported stolen across ~247K confirmed installs. Koi Security's first audit (Feb 1, 2026) found 341 malicious skills among 2,632 audited (registry then ~2,857); Koi's count rose to 824 by mid-February as the marketplace passed 10,700 skills, and by Feb 19 Antiy CERT had uncovered at least 1,184 malicious skill packages tied to 12 compromised publisher accounts – published estimates of the flagged-malicious-or-suspicious share of the registry range from ~7.6% to ~20% (“approximately one in five packages at the time of discovery”), depending on registry-size denominator and date; payloads included Windows VMProtect-packed infostealers alongside AMOS. Unit 42 (Palo Alto) has since published its own AI supply-chain threat analysis of the marketplace. ClawHub now screens every submission via VirusTotal plus code-level ClawScan analysis (3,016+ samples analyzed). The campaign was compounded by a one-click RCE (CVE-2026-25253, patched). Developers now distrust third-party skills, registry vetting is heightened, and – because the malware targeted crypto wallets – this is a credibility landmine for a crypto-adjacent provider. The Gonka skill submission must lead with security: signed releases, auditable source, minimal permissions.

The viral loop:

```
Developer A configures Gonka (pastes config snippet)
-> Saves on heartbeat costs (session persistence)
-> Shares cost comparison screenshot on Reddit/Twitter
-> Developer B sees post, copies config snippet from comment
-> Developer B configures Gonka in 90 seconds
-> Developer B saves on heartbeat costs
-> Developer B shares their own comparison...
```

This loop is driven by the session persistence savings – the cost reduction was modeled at 73% at Active tier in April 2026 (`gonka_agent_pricing_analysis.md` Section 6), computed against mostly-uncached competitor baselines. Prompt caching is now universal (OpenAI/Anthropic 90% off cached input, Together default-on, Moonshot cache-hit discounts on current models, DeepSeek \$0.0028; Moonshot's first-party K2.5 rates are no longer quotable – the model is closed to new users and sunsets August 31, 2026), so the honest delta is smaller and must be recomputed before any shared comparison claims 73%. Even at a reduced delta, the savings plus the config snippet keep sharing frictionless. No referral program is needed to trigger the initial loop; the savings create natural word-of-mouth. The referral program (`gonka_developer_personas.md` Weekend Builder Referral stage: “\$5 credit for each referred developer who makes 100+ API calls”) amplifies an already-functioning viral mechanism.

What Competitors Require

Gonka's time-to-first-inference target (under 5 minutes, email-only signup) creates a measurable friction advantage over the competitors that Weekend Builders and Startup CTOs are most likely evaluating.

OpenRouter (Weekend Builder's Current Provider)

Step	OpenRouter Requirement	Time	Friction
1	Visit openrouter.ai	0s	None
2	Click "Sign In"	10s	None
3	GitHub OAuth login	30-60s	Requires GitHub account. OAuth redirect. Permission grant screen. Some developers hesitate to grant OAuth permissions to unknown services
4	Credit card for paid tier	120-180s	Credit card entry form. Address verification. Payment processing. This step does not exist for Gonka's free tier
5	Browse 400+ model catalog (60+ providers), select model	60-120s	Model selection paralysis. "Which model do I pick?" No guidance for agent workloads. Developer must research models independently
6	Copy API key, configure application	60-90s	API key is available immediately after signup (no email verification)
Total		~5-8 min (free tier) or ~8-12 min (paid)	GitHub OAuth + credit card + model selection

Gonka advantage: No OAuth, no credit card, no model selection paralysis. Gonka offers one model family (Kimi K2.6, with a K3 frontier tier pending open-weight availability) in quantization tiers with automatic routing – the developer does not need to choose. Email-only signup eliminates the OAuth permission screen that causes hesitation. The free tier eliminates the credit card step entirely.

Source: gonka_competitive_feature_matrix.md (OpenRouter catalog re-verified July 2026: 400+ active models from 60+ providers, per openrouter.ai/models); gonka_developer_personas.md (Weekend Builder: OpenRouter is the default, zero-config provider).

Together AI (Startup CTO's Primary Alternative)

Step	Together AI Requirement	Time	Friction
1	Visit api.together.ai	0s	None
2	Click "Sign Up"	10s	None
3	OAuth login (Google or GitHub)	30-60s	OAuth redirect. Permission grant. Account creation
4	Credit card required for all usage	120-180s	No free tier. Credit card is mandatory. Some developers abandon at this step for evaluation purposes
5	Select model from 200+ catalog	60-120s	Less paralysis than OpenRouter (200 vs 400+ models) but still requires research. The Kimi family is one of many options
6	Copy API key, configure application	60-90s	API key generation is straightforward
Total		~5-8 min	OAuth + mandatory credit card + model selection

Gonka advantage: No credit card requirement (free tier exists). No model selection step (K2.6 with auto-tiering). Email-only signup (no OAuth). The Startup CTO can evaluate Gonka on staging without entering payment information or requesting procurement approval from their finance team.

Source: gonka_agent_pricing_analysis.md Section 4 (Together AI pricing: \$0.50/\$2.50); gonka_provider_landscape_map.md Segment 3 (Together AI: pay-per-use only, no free tier).

Self-Hosted vLLM (Privacy-First Builder's Current Approach)

Step	vLLM Self-Hosting Requirement	Time	Friction
1	Acquire GPU hardware (\$2,000-10,000+)	Days to weeks	Hardware procurement. Budget approval. Delivery time
2	Install CUDA, Python, vLLM	30-60 min	Driver compatibility issues. Python environment management. Dependency conflicts
3	Download K2.6 model weights	20-60 min	Model weights are 10-50 GB depending on quantization. Requires HuggingFace account. Network bandwidth dependent
4	Configure vLLM server	15-30 min	Configuration options for GPU memory, quantization, batch size. Trial and error for optimal settings
5	Start server, verify inference works	5-10 min	Health checks, first inference warm-up, verify output quality
6	Configure OpenClaw to point to local vLLM	5-10 min	Same openclaw.json config as Gonka, but with localhost URL
Total		Hours to days	Hardware + software + model download + configuration

Gonka advantage: Under 5 minutes vs hours to days. No hardware procurement. No driver installation. No model download. The Privacy-First Builder gets the same open-weight K2.6 model with no content filtering, served on distributed infrastructure with no central logging – and they can start using it in 4 minutes instead of 4 hours. The tradeoff: Gonka’s privacy guarantees are architectural (no central log aggregation) rather than absolute (no TEE yet). The honest capabilities page (gonka_developer_personas.md Privacy-First Activation stage) communicates this clearly.

Source: gonka_developer_personas.md Privacy-First Pain Point #2 (self-hosting costs and maintenance); gonka_competitive_feature_matrix.md (K2.6: open-weight, no content filtering).

Failure Modes and Recovery

Each step in the time-to-first-inference plan has a specific failure mode that would break the under-5-minute target. These must be engineered against:

Step	Failure Mode	Impact	Prevention
3	Email submission rate-limited or blocked by spam filter	Developer cannot sign up	Use transactional email service with dedicated IP. Monitor deliverability rates. Provide “did not receive email?” link with resend option
4	Verification email delayed >2 minutes	Developer loses context, opens new tab, forgets about Gonka	Alert on email delivery latency >30s. Use multiple email service providers with failover. Display “email sent, check your inbox” with a countdown timer
5	API key page fails to load or copy button does not work	Developer cannot retrieve their key	Client-side key display (no server round-trip after page load). Fallback: display key as selectable text if clipboard API fails. Email the key as backup
6	Developer pastes config but misses Step 2 (model allowlisting)	Silent failure – agent ignores Gonka provider entirely	Show both steps on the same page with numbered instructions. Verification endpoint: GET /v1/verify that returns 200 if the key is valid, helping developers test before configuring their agent

Step	Failure Mode	Impact	Prevention
7	First API call times out or returns 500	Permanent trust damage – developer will not try again (PITFALLS.md Pitfall 3)	Health-checked routing ensures requests go only to responsive nodes. Cold-start mitigation: keep K2.6 warm on minimum N nodes. Error response includes troubleshooting link

Cross-References

Every claim in this document traces to a specific Phase 15-19 output document. This section lists each source and what was drawn from it.

Document	What Was Used
gonka_developer_personas.md (Phase 16)	Three personas (Weekend Builder, Startup CTO, Privacy-First Builder) with decision drivers, pain points, adoption triggers, objections, and complete AAARRRP journey maps. Per-persona conversion drivers at each funnel stage. Drop-off risk analysis. Referral mechanisms.
gonka_agent_pricing_analysis.md (Phase 15)	Workload tier definitions (Casual: 31.5M tokens/month, Active: 161.5M, Heavy: 2.94B). Heartbeat overhead percentages (44%, 51%, 85%). Scenario B pricing (\$0.35/\$1.75). Session persistence savings (80% token reduction). Monthly cost projections per provider per tier. Competitor pricing data (Together AI, DeepInfra, OpenRouter, OpenAI, Anthropic).
gonka_competitive_feature_matrix.md (Phase 15)	8-dimension competitive comparison. WIN verdicts on Agent Sessions and Model Tiering. LOSE verdicts on Model Breadth and Uptime/Reliability. Competitor free tier analysis. Content filtering comparison (Gonka: "None (open)"). Quantization-tier concept (originally modeled on K2.5, re-based to K2.6).
gonka_provider_landscape_map.md (Phase 15)	5 critical must-close gaps (not built-in provider, no docs site, no self-serve signup, no pricing page, in-memory sessions). 4-segment landscape framework. Gonka's unique position as only decentralized provider with agent-native extensions.
gonka_message_house.md (Phase 17)	Core positioning statement (73% cost reduction via session persistence). 16-item never-say list. Approved vocabulary mapping. VP1-VP3 ranked value propositions. Competitive differentiation statements per persona.
gonka_channel_strategy.md (Phase 18)	Primary KPI: API-active developers (>100 calls/month). Target trajectory (10 Month 1, 50 Month 3, 200 Month 6, 1000 Month 12). 70/30 channel split. P0 channels (OpenClaw Provider Directory, GitHub, Discord). Anti-metrics (follower counts, Discord members, GitHub stars).
gonka_partnership_playbook.md (Phase 19)	Four-tier OpenClaw integration roadmap (Listed -> Plugin -> Built-In -> Preferred). ClawHub submission plan. Technical requirements checklist (33 items). Community-first PR strategy timeline (3-6 months).
gonka_objection_playbook.md (Phase 17)	ACE objection framework. 12 documented objections with persona mapping. P0 severity objections: awareness, crypto perception, SLA gap, prompt privacy.
gonka_agent_native_pitch.md (Phase 17)	Agent-as-customer thesis. 7 task profiles with provider match analysis. Gonka wins 3/7 (long-running agents, multi-agent systems, privacy-sensitive).
.planning/research/STACK.md	OpenClaw <code>openclaw.json</code> configuration format. Two-step provider selection process (provider definition + model allowlisting). Built-in vs custom provider distinction. Key rotation and failover behavior.

Document: *gonka_plg_growth_model.md* | Version 1.3 | 2026-07-18 Capstone deliverable for Phases 15-19.

Companion document: *gonka_v14_engineering_backlog.md* (Plan 20-02)