

Gonka Product-Led Growth Model

Version: 1.0 **Date:** 2026-04-01 **Classification:** Internal – actionable growth model for Gonka developer acquisition
Dependencies: Phase 15 (competitive analysis), Phase 16 (developer personas), Phase 17 (messaging), Phase 18 (channel strategy), Phase 19 (partnership playbook) **Requirement:** GTM-03

Executive Summary

This document synthesizes all Phase 15-19 research into a single actionable growth model for Gonka’s developer adoption. It defines the PLG funnel (how developers move from discovery to payment), the free tier design spec (what developers get without paying), and the time-to-first-inference plan (how developers go from zero to their first API response in under 5 minutes).

Every conversion rate, usage limit, and design decision is grounded in evidence from the research corpus. Where data is indirect (industry benchmarks rather than Gonka-specific measurements), confidence levels are noted. The model is designed so that leadership can read this document and know exactly what Gonka must build, what the targets are, and how to measure progress.

All developer-facing copy in this document follows the Phase 17 vocabulary guidelines (gonka_message_house.md Section 6). No terms from the 16-item never-say list appear in any developer-facing recommendation.

PLG Funnel Model

Funnel Overview

The PLG funnel maps Gonka’s developer acquisition to the AAARRRP framework from Phase 16 (gonka_developer_personas.md). Six stages capture the complete journey from first encounter to paying customer, with each stage producing a measurable outcome and a clear handoff to the next.

Stage	AAARRRP	Definition	Target Rate	Metric
Discover	Awareness	Developer encounters Gonka through OpenClaw ecosystem channels, Reddit, GitHub, or technical content	– (top of funnel)	Unique visitors to docs.gonka.ai
Explore	Acquisition	Developer visits documentation, reads quickstart guide, reviews pricing page	40-50% of Discover	Page views on docs.gonka.ai; time on site > 2 minutes
Sign Up	Acquisition	Email-only API key creation – no wallet, no OAuth, no credit card	15-25% of Explore	API keys created per week
First Inference	Activation	First successful API call through Gonka network	70-80% of Sign Up	First API call within 24 hours of signup
Habitual Use	Retention	100+ API calls in 30 days – developer has integrated Gonka into a real workflow	20-30% of First Inference	API-active developers per month (gonka_channel_strategy.md primary KPI)
Paid Conversion	Revenue	Developer exceeds free tier limits and upgrades to paid plan	5-10% of Habitual Use	Paying users; monthly revenue per user

Why these rates: The target conversion rates are calibrated against developer tool industry benchmarks (Stripe, Twilio, and Vercel report 10-25% signup-to-activation rates; Gonka targets higher because the activation step is a single config paste, not an SDK integration). The 5-10% paid conversion rate reflects the typical PLG pattern where most value is delivered free and only production-scale workloads trigger payment (gonka_developer_personas.md: Weekend Builder stays on free tier; Startup CTO converts to paid).

Per-Persona Conversion Drivers

Each funnel stage has a different primary driver depending on the developer persona. Understanding these drivers determines what content, features, and messaging to prioritize at each stage.

Stage 1: Discover

Persona	Primary Conversion Driver	#1 Drop-off Risk	Mitigation
Weekend Builder	Reddit r/LocalLLaMA cost comparison post showing Gonka's heartbeat savings vs OpenRouter (\$40/mo to \$13/mo)	"Never heard of Gonka – is this a crypto scam?" Landing page looks like a DeFi protocol	Landing page must resemble Vercel/Supabase, not a token project. Lead with "Cut your agent costs by 73%" not "Decentralized AI inference" (gonka_message_house.md VP1; PITFALLS.md Pitfall 1)
Startup CTO	Hacker News technical deep-dive on agent heartbeat overhead problem with Gonka cited as the session persistence solution	"Decentralized infrastructure means unreliable. Akash and Render both failed at scale"	Technical white paper with uptime data and p95 latency benchmarks. Must read like a Cloudflare blog post, not a crypto whitepaper (gonka_developer_personas.md Startup CTO Awareness stage)
Privacy-First Builder	GitHub repository README or technical article on privacy-preserving inference architectures, found while researching alternatives to self-hosted vLLM	"Every cloud provider claims no logging. How is Gonka different architecturally?"	Honest capabilities page: current guarantees (no central log aggregation, open-weight model, no content filtering) vs roadmap (TEE-based encrypted inference) (gonka_developer_personas.md Privacy-First Awareness stage)

Source: gonka_developer_personas.md AAARRRP Journey Maps, Awareness stage for each persona.

Stage 2: Explore

Persona	Primary Conversion Driver	#1 Drop-off Risk	Mitigation
Weekend Builder	Quickstart page showing the 3-field openclaw.json snippet with a copy button – "Add Gonka in 90 seconds"	Page mentions tokens, staking, or wallet before showing the config snippet	Docs homepage leads with the config snippet, not network architecture. Follow never-say list (gonka_message_house.md Section 6.1)
Startup CTO	Architecture overview explaining reliability through redundant nodes and health-checked routing, plus a cost comparison calculator	Docs lack uptime data, SLA information, or team account management	Enterprise evaluation page with p95 latency data, multi-key team accounts, and architecture explanation using cloud infrastructure language (gonka_developer_personas.md Startup CTO Acquisition stage)
Privacy-First Builder	Data handling policy page and open-source infrastructure code link (github.com/mitgor/gonka-ai-infrastructure) – reviewable before signup	Privacy documentation is vague or makes claims without architectural backing	Published data handling policy: what is logged, what is not, retention periods, node operator agreements. Link to open-source code for independent audit (gonka_developer_personas.md Privacy-First Acquisition stage)

Source: gonka_developer_personas.md AAARRRP Journey Maps, Acquisition stage for each persona.

Stage 3: Sign Up

Persona	Primary Conversion Driver	#1 Drop-off Risk	Mitigation
Weekend Builder	Single-CTA "Get API Key" button leading to email-only form. API key displayed immediately after email verification. Zero friction	Form asks for company name, use case, or any information beyond email address	Minimal form: email only. Verification link sent instantly. API key displayed on verification page with one-click copy. No approval queue, no waitlist (CONTEXT.md locked decision: email-only signup)

Persona	Primary Conversion Driver	#1 Drop-off Risk	Mitigation
Startup CTO	Team account creation with multiple API keys for dev/staging/prod environments. Architecture docs reviewable pre-signup	Signup requires individual accounts for each team member; no team management	Team signup flow: one email creates team account, generates scoped API keys per environment. Billing attached to team, not individual (gonka_developer_personas.md Startup CTO Acquisition stage)
Privacy-First Builder	Data handling policy reviewed and satisfactory before email submission. Open-source code independently auditable	Privacy policy is generic boilerplate that does not address inference-specific concerns (prompt logging, content filtering, data residency)	Inference-specific privacy policy linked prominently on signup page. “We do not log your prompts” with architectural explanation, not just a policy claim (gonka_developer_personas.md Privacy-First Acquisition stage)

Source: gonka_developer_personas.md AAARRRP Journey Maps; CONTEXT.md decisions.

Stage 4: First Inference

Persona	Primary Conversion Driver	#1 Drop-off Risk	Mitigation
Weekend Builder	Pre-built openclaw.json snippet on docs page with copy button. “Add Gonka to OpenClaw in 90 seconds” tutorial	The two-step provider gotcha: developer defines provider but forgets to add model to agent allowlist, causing silent failure	Quickstart guide explicitly shows both steps (provider definition AND model allowlisting) with a troubleshooting section for the gotcha (STACK.md Section 1: provider selection is a TWO-STEP process)
Startup CTO	Integration test results from v1.2 showing OpenClaw, CrewAI, and LangGraph compatibility. Staging deployment guide	First test request times out or returns unexpected error, destroying trust at first contact	Sub-200ms p95 latency on first call. Clear error messages with actionable troubleshooting. Integration test suite available for the team to run against their staging environment (gonka_developer_personas.md Startup CTO Activation stage; PITFALLS.md Pitfall 3)
Privacy-First Builder	Test prompt that would be blocked by OpenAI/Anthropic content filters completes successfully on Gonka. Aha moment: “It processed my prompt without filtering”	First prompt triggers an unexpected error or content filter, contradicting “no content policy” claim	K2.5 is open-weight with no built-in content filtering. Verify no content filtering middleware in the gateway. Test with adversarial prompts during QA (gonka_competitive_feature_matrix Content filtering – Gonka “None (open)”)

Source: gonka_developer_personas.md AAARRRP Journey Maps, Activation stage; STACK.md.

Stage 5: Habitual Use

Persona	Primary Conversion Driver	#1 Drop-off Risk	Mitigation
Weekend Builder	Usage dashboard showing session persistence savings: “\$40/month on OpenRouter vs \$13/month on Gonka for the same agent.” Cost comparison is visceral and shareable	After initial excitement, the developer forgets about Gonka and defaults back to OpenRouter because it is built-in	Weekly email digest showing cost savings. X-Gonka-Usage-Remaining header in every response keeps savings visible. Cost comparison calculator in dashboard (gonka_developer_personas.md Weekend Builder Retention stage)

Persona	Primary Conversion Driver	#1 Drop-off Risk	Mitigation
Startup CTO	Weekly cost report showing session persistence savings across all agents. Gradual migration from 2 agents to 4 to all 6 as confidence builds	One timeout or error during production usage triggers revert to previous provider. First impressions are permanent for infrastructure trust	Real-time status page (status.gonka.ai) with per-model availability and historical uptime. Incident report publication process. Gradual migration path with dual-provider configuration (gonka_developer_personas.md Startup CTO Retention stage; PITFALLS.md Pitfall 3)
Privacy-First Builder	Ongoing use for sensitive workloads with periodic privacy posture auditing via open-source verification tools	Gonka updates terms of service or infrastructure in a way that weakens privacy guarantees without communicating the change	Privacy-specific changelog (separate from feature releases). Automated privacy audit tool. Community-maintained watchdog (gonka_developer_personas.md Privacy-First Retention stage)

Source: gonka_developer_personas.md AAARRRP Journey Maps, Retention stage.

Stage 6: Paid Conversion

Persona	Primary Conversion Driver	#1 Drop-off Risk	Mitigation
Weekend Builder	Exceeds free tier credit (\$5-10 trial). Clear pricing page showing per-token rates with session savings calculator. Credit card entry with spending cap	Fear of runaway costs – “OpenClaw agents can spike unpredictably, \$300+ in 2 days”	Configurable daily spending cap. Usage alerts at 50%, 80%, 100% of limit. “No surprise bills” guarantee. Pricing comparison table: Gonka vs OpenRouter vs Together AI at Casual tier (gonka_developer_personas.md Weekend Builder Revenue stage)
Startup CTO	Commits to paid plan after 2-4 weeks of staging + gradual production rollout. Volume pricing for 161M+ tokens/month	Needs formal SLA with downtime credits to justify to co-founder and investors	Enterprise pricing tier with committed-use discounts. SLA framework: 99.5% uptime target with credit mechanism. Invoice and receipt system for accounting (gonka_developer_personas.md Startup CTO Revenue stage; gonka_competitive_feature_matrix. Section 7)
Privacy-First Builder	Converts to paid plan valuing session persistence for reducing how often sensitive context traverses the network. Willing to pay premium for TEE when available	No privacy-specific pricing tier or TEE timeline commitment	Privacy tier with optional TEE-backed inference at premium when available. Roadmap commitment on encrypted inference with quarterly updates. Early access program for privacy beta features (gonka_developer_personas.md Privacy-First Revenue stage)

Source: gonka_developer_personas.md AAARRRP Journey Maps, Revenue stage; gonka_agent_pricing_analysis.md.

Funnel Math

If 1,000 developers discover Gonka (visit docs.gonka.ai or encounter Gonka in the OpenClaw ecosystem), here is how many become paying users at each scenario:

Stage	Pessimistic	Target	Optimistic
Discover	1,000	1,000	1,000
Explore (40-50% of Discover)	400	450	500
Sign Up (15-25% of Explore)	60	90	125
First Inference (70-80% of Sign Up)	42	67	100

Stage	Pessimistic	Target	Optimistic
Habitual Use (20-30% of First Inference)	8	17	30
Paid Conversion (5-10% of Habitual Use)	0.4 (~0)	1.3 (~1)	3

Interpretation:

- **Pessimistic:** ~0.04% end-to-end conversion. From 1,000 developers who discover Gonka, zero to one becomes a paying customer. This scenario reflects a cold start with no ecosystem presence, minimal content, and the crypto perception barrier fully active. Recovery requires aggressive investment in P0 channels (OpenClaw Provider Directory, GitHub, Discord – gonka_channel_strategy.md).
- **Target:** ~0.13% end-to-end conversion. From 1,000 discovers, approximately 1 paying user and 17 API-active developers generating usage data, cost savings testimonials, and referral potential. At the Month 6 target of 200 API-active developers (gonka_channel_strategy.md), this implies needing approximately 12,000 discover-stage developers.
- **Optimistic:** ~0.3% end-to-end conversion. From 1,000 discovers, 3 paying users and 30 API-active developers. Achievable if Gonka secures built-in OpenClaw provider status (Tier 3 in gonka_partnership_playbook.md), which eliminates the manual configuration friction at the Sign Up and First Inference stages.

Key insight: The funnel is widest at the top (Discover to Explore) and steepest at the Habitual Use stage. The 20-30% conversion from First Inference to Habitual Use is the critical gate – this is where developers decide if Gonka’s session persistence savings are real enough to justify switching from their default provider. The X-Gonka-Usage-Remaining header and cost comparison dashboard are the primary retention mechanisms at this stage.

Revenue projection at target rates:

- 17 API-active developers at Casual tier (~31.5M tokens/month each) at Scenario B pricing (\$1.05/1M blended): ~\$563/month total inference revenue
- 1 paying Startup CTO at Active tier (~161.5M tokens/month) at Scenario B pricing: ~\$170/month
- Combined Month 6 revenue from 1,000 discover-stage cohort: ~\$733/month

These projections use Scenario B pricing (\$0.35/\$1.75 per 1M tokens input/output) from gonka_agent_pricing_analysis.md. Revenue scales with discover-stage volume, which is driven by P0 channel investment (gonka_channel_strategy.md).

Free Tier Design Spec

Design Principles

The free tier must accomplish three goals simultaneously:

1. **Eliminate all friction for the Weekend Builder** – generous enough that a solo developer running one agent on one channel can use Gonka indefinitely without paying, experiencing session persistence savings firsthand and becoming a referral source (gonka_developer_personas.md: Weekend Builder monthly spend \$20-50 on OpenRouter; free tier must cover their reduced-cost Gonka usage)
2. **Create a natural upgrade trigger for the Startup CTO** – limited enough that a team running multiple agents across channels hits the ceiling within their first week of serious evaluation, triggering the paid conversion conversation (gonka_developer_personas.md: Startup CTO monthly spend \$200-500; free tier should cover approximately 1 agent, not 6)
3. **Remove the crypto barrier entirely** – Email-only signup with no wallet, no token purchase, no OAuth complexity. The word “wallet” must never appear before “API key” on any page (gonka_message_house.md Section 6.1: Never-Say List – “wallet” is the single most alienating word for Web2 developers)

Signup Flow

- Step 1: Developer lands on docs.gonka.ai
Single CTA: "Get Your API Key" button
No distracting navigation, no "Learn about our network" sidebar
- Step 2: Email-only form
One field: email address
No company name, no use case, no phone number
Submit button: "Send me my API key"
- Step 3: Verification email arrives (target: under 60 seconds)
Transactional email service (SendGrid, Postmark, or SES)
Subject: "Your Gonka API Key"
Body: One-click verification link, nothing else
- Step 4: Developer clicks link, lands on key display page
API key displayed prominently with one-click copy button
Below the key: the exact openclaw.json snippet pre-filled with their key
Below the snippet: "Paste this into ~/.openclaw/openclaw.json and you are done"
- Step 5: Developer is using Gonka
No approval queue. No waitlist. No manual review.
Free tier limits active immediately.

Per CONTEXT.md locked decision: Email-only signup. No wallet or crypto knowledge required.

Usage Limits

Free tier limits are calibrated against two data points:

- Competitor free tiers:** OpenRouter offers free-tier access with rate limiting and quality degradation under load. Google Gemini offers a generous free tier. No K2.5-specific free tier exists among compared providers (gonka_competitive_feature_matrix.md; gonka_provider_landscape_map.md Segment 3: Together AI and Deep-Infra are pay-per-use only).
- Persona workload profiles:** The Weekend Builder consumes ~31.5M tokens/month at Casual tier with 30-minute heartbeats (gonka_agent_pricing_analysis.md Section 3). With Gonka's session persistence reducing effective token consumption by ~80%, the Weekend Builder's actual Gonka token usage would be approximately 12.6M tokens/month. The free tier should cover most of this.

Limit	Free Tier Value	Rationale
Daily requests	1,000 requests/day	Covers 1 agent with 50 user messages + 150 tool calls + 48 heartbeats = ~248 requests/day. Generous buffer for burst testing. Startup CTO with 6 channel-agents needs ~1,488 requests/day – exceeds limit within first day of serious evaluation (gonka_agent_pricing_analysis.md Section 3: Active tier daily breakdown)
Monthly tokens	15,000,000 tokens/month	Covers Weekend Builder's session-optimized usage (~12.6M tokens/month) with headroom. Startup CTO at Active tier (~56.3M session-optimized tokens/month) exceeds this in week 1 (gonka_agent_pricing_analysis.md Section 5: Gonka hidden cost analysis)
Concurrent sessions	2 active sessions	Covers 1 agent on 1-2 channels (Weekend Builder). Startup CTO with 6 channel-agents needs 6 concurrent sessions – hits upgrade trigger immediately when deploying team workload
Models available	K2.5 lite and mid tiers only	Full-precision K2.5 reserved for paid tier. Lite and mid tiers are sufficient for the Weekend Builder's personal assistant workload. Startup CTO evaluating quality for production needs full tier access (gonka_competitive_feature_matrix.md: 3-tier K2.5 quantization)

Limit	Free Tier Value	Rationale
Rate limit	60 requests/minute	Prevents abuse while allowing burst testing (a developer testing their agent configuration may send 10-20 requests in quick succession). Production agents with 30-minute heartbeats generate ~1 request/minute baseline. 60 RPM provides comfortable headroom
Context window	32,768 tokens	Reduced from K2.5's full 131,072. Sufficient for most agent workloads (system prompt + context is ~9,600 tokens). Startup CTO running complex multi-step workflows with deep context needs the full window

Upgrade Triggers

Three specific mechanisms notify developers when they approach or exceed free tier limits, driving paid conversion without hard-blocking their workflow:

Trigger 1: Usage Limit Approaching

When a developer reaches 80% of their daily request limit or monthly token limit, every API response includes a warning header:

```
X-Gonka-Usage-Remaining: 200
X-Gonka-Usage-Limit: 1000
X-Gonka-Usage-Reset: 2026-04-02T00:00:00Z
X-Gonka-Upgrade-URL: https://docs.gonka.ai/pricing
```

At 100%, requests return HTTP 429 with a clear message: "Free tier daily limit reached. Resets at midnight UTC. Upgrade at docs.gonka.ai/pricing for unlimited requests."

No cryptic error codes. No silent failures. The developer knows exactly what happened, when it resets, and how to fix it.

Trigger 2: Session Limit Reached

When a developer attempts to create a third concurrent session (free tier allows 2), the API returns HTTP 403 with:

```
{
  "error": "session_limit_reached",
  "message": "Free tier supports 2 concurrent sessions. You have 2 active sessions. Upgrade to paid for unl",
  "current_sessions": 2,
  "max_sessions": 2,
  "upgrade_url": "https://docs.gonka.ai/pricing"
}
```

This trigger is designed specifically for the Startup CTO persona: they have proven Gonka works with 1-2 agents on staging and now want to add their production agents. The session limit creates a natural conversation point for team account and volume pricing discussion.

Trigger 3: Model Tier Access

When a developer requests the full-precision K2.5 model (available only on paid tier), the API routes to mid-tier K2.5 and includes a header:

```
X-Gonka-Tier-Served: mid
X-Gonka-Tier-Requested: full
X-Gonka-Tier-Upgrade: "Full-precision K2.5 available on paid tier. See docs.gonka.ai/pricing"
```

The request succeeds (no error) but with a lower-tier model, and the header transparently communicates the substitution. This avoids breaking the developer's workflow while making the quality difference discoverable.

What Is NOT in the Free Tier

Features reserved for paid plans to create clear value differentiation:

Feature	Free Tier	Paid Tier	Why Reserved
Full-precision K2.5	No (lite and mid only)	Yes	Quality differentiation for production workloads
Unlimited sessions	2 concurrent max	Unlimited	Multi-agent production requires paid commitment
SLA guarantee	Best effort	99.5% uptime target with credits	SLA has cost implications; reserved for committed customers
Priority routing	Standard queue	Priority queue with dedicated capacity	Ensures paid customers get consistent low latency
Webhook notifications	No	Yes	Advanced agent feature for production use (gonka_competitive_feature_matrix.md: Tier 2 differentiator)
Memory API	No	Yes	Persistent key-value memory for long-running agents
Usage API / Dashboard	Basic (current usage only)	Full (historical trends, cost comparison, export)	Detailed analytics drive retention for paid users
Team accounts	No (individual only)	Yes (multi-key, scoped permissions)	Team management is an enterprise feature
Spending caps	N/A (free)	Configurable daily/monthly limits	Cost control for production workloads (addresses “\$300+ in 2 days” fear – gonka_developer_personas.md Weekend Builder Revenue stage)

Competitive Positioning of Free Tier

Provider	Free Tier	Signup Friction	Limits
OpenRouter	Free-tier access to select models	GitHub OAuth	Rate-limited; quality degraded under load; free requests queued behind paid (gonka_competitive_feature_matrix.md Section 5)
Google Gemini	Generous free tier	Google account	Rate limits vary by model; generous for experimentation
Together AI	No free tier	Credit card required	Pay-per-use only (gonka_provider_landscape_map.md Segment 3)
DeepInfra	No free tier	Credit card required	Pay-per-use only
Gonka (proposed)	15M tokens/month, 1,000 requests/day, 2 sessions	Email only	Most generous free tier among K2.5 providers; lowest signup friction; session persistence active on free tier

Competitive advantage: Gonka’s free tier is the only one that (a) requires no credit card or OAuth, (b) includes session persistence (the primary cost-saving feature), and (c) provides enough capacity for a solo developer’s always-on agent. The Weekend Builder can experience the full session persistence savings without paying, creating the cost comparison data point (“\$40/month on OpenRouter vs free on Gonka”) that drives organic referrals (gonka_developer_personas.md Weekend Builder Referral stage).

Copy Guidelines

All developer-facing copy on signup pages, error messages, and upgrade prompts must follow the Phase 17 vocabulary guidelines (gonka_message_house.md Section 6). Specifically:

- **Never use:** wallet, staking, mining, DePIN, Web3, tokenomics, consensus mechanism, on-chain, smart contract, gas, slashing, governance proposals, epochs, validators, token (as cryptocurrency), decentralized (as headline)
- **Always use:** “API key” (not “wallet”), “account” (not “wallet”), “compute credits” (not “GNK tokens”), “GPU providers” (not “miners” or “validators”)
- **Error messages:** Must be actionable and specific. “Free tier daily limit reached. Resets at midnight UTC.” not “Insufficient GNK balance.”

Time-to-First-Inference Plan

Objective

Per CONTEXT.md locked decision: target under 5 minutes from email submission to first successful API response through the Gonka network. This is the single most important PLG metric – every minute of friction between “I want to try Gonka” and “my agent is running on Gonka” loses developers who will never return (PITFALLS.md Pitfall 3: first impressions are permanent for infrastructure trust).

Step-by-Step Breakdown

Step	Action	Time Budget	Blocker Risk	What Must Be Built
1	Land on docs.gonka.ai	0s (assumed)	Page must exist. Currently no docs site – must-close gap from Phase 15 (gonka_provider_landscape_map.6EA 5 critical gaps)	docs.gonka.ai with SSL, CDN, sub-2s page load. Single-page-app or static site (Hugo, Astro, or Next.js). Landing page with one CTA
2	Click “Get API Key”	10s	Single CTA must be above the fold. No competing navigation, no “Learn about our network” distractions. Developer must not need to scroll to find the button	CTA button component. Landing page design that passes the “is this crypto?” test (gonka_message_house.md VP3; PITFALLS.md Pitfall 1)
3	Enter email, submit	30s	Email-only form. No wallet creation, no OAuth redirect, no credit card. One field, one button. Any additional friction at this step causes the steepest drop-off in the funnel (gonka_developer_personas.md: 95% funnel drop at wallet step per PITFALLS.md Pitfall 2)	Email submission endpoint. Rate limiting on submissions (prevent abuse). CAPTCHA only if abuse is detected, not by default
4	Receive verification email	60s	Must arrive within 60 seconds. Delayed verification emails (>2 minutes) cause developers to lose context and abandon. Use a transactional email service (SendGrid, Postmark, or AWS SES) with dedicated IP and domain authentication (SPF, DKIM, DMARC) to avoid spam folders	Email sending infrastructure. Transactional email template. Email deliverability monitoring
5	Click verification link, see API key	15s	Key must be displayed immediately on the page with a one-click copy button. No “check your dashboard” redirect. No additional login step. The key is shown once, prominently, with a “Copy” button that provides visual confirmation	Key generation service. Key display page with clipboard API integration. Key storage (hashed) in database
6	Paste config into openclaw.json	90s	The exact config snippet must be pre-filled with the developer’s API key and ready to copy-paste. Show both steps of the two-step process (provider definition AND model allowlisting) to avoid the silent failure gotcha (STACK.md Section 1)	Pre-filled config snippet generator. Documentation page showing the exact JSON to paste
7	Send first message to agent	60s	Agent responds via Gonka inference. The first response must arrive within 2-3 seconds (sub-200ms TTFT + streaming). Any timeout or error at this step permanently damages trust	Live API endpoint at api.gonka.ai. Health-checked GPU hosts. vLLM inference with K2.5 model loaded and warm
Total		~4 min 15s		

Step 6 Detail: The Config Snippet

After the developer copies their API key, the verification page displays the exact `openclaw.json` configuration they need to paste. This is not a generic template – it is pre-filled with their specific API key:

```
// Step 1: Add this to your ~/.openclaw/openclaw.json under models.providers:
{
  "gonka": {
    "baseUrl": "https://api.gonka.ai/v1",
    "apiKey": "gnk_sk_abc123...", // Your API key (already filled in)
```

```

"api": "openai-completions",
"models": [
  {
    "id": "kimi-k2.5",
    "name": "Kimi K2.5",
    "contextWindow": 131072,
    "maxTokens": 8192,
    "cost": { "input": 0.35, "output": 1.75 }
  }
]
}
}

// Step 2: Add this to your agents.defaults.models to enable the model:
{
  "gonka/kimi-k2.5": { "alias": "k2.5" }
}

```

Why both steps are shown: OpenClaw requires both provider definition (Step 1) AND model allowlisting (Step 2). Missing Step 2 causes silent failure – the agent ignores the provider entirely with no error message. This is a known gotcha in the OpenClaw community (STACK.md Section 1: “Missing either step causes silent failure”). Showing both steps on the same page with clear “Step 1” and “Step 2” labels prevents the most common configuration error.

Config snippet format: Based on the OpenClaw `openclaw.json` configuration format documented in STACK.md. The `api: "openai-completions"` value tells OpenClaw to use the standard `/v1/chat/completions` endpoint format, which Gonka’s gateway already implements (shipped in v1.2). Cost values reflect Scenario B pricing (\$0.35/\$1.75 per 1M tokens input/output) from `gonka_agent_pricing_analysis.md`.

Current Gaps vs Target Flow

Step	Current Status	Gap	Engineering Work Required	Cross-Reference
1	No docs site exists	Must-close gap	Static site deployment (Hugo/Astro/Next.js) with domain, SSL, CDN	gonka_provider_landscape_map.md: critical gap #2
2	No landing page exists	Must-close gap	Landing page design and implementation	gonka_provider_landscape_map.md: critical gap #2
3	No signup flow exists	Must-close gap	Email submission API, rate limiting, form validation	gonka_provider_landscape_map.md: critical gap #3
4	No email infrastructure	Must-close gap	Transactional email service integration, domain auth	New requirement (identified in this plan)
5	No key generation service	Must-close gap	API key generation, storage, display page	gonka_provider_landscape_map.md: critical gap #3
6	Config snippet documented in STACK.md (research only)	Needs productization	Pre-filled snippet generator on docs site	gonka_partnership_playbook.md: requirement #13
7	API endpoint exists (v1.2) but not publicly accessible	Must-close gap	Public deployment of api.gonka.ai with domain, SSL, health checks	gonka_provider_landscape_map.md: critical gap #1

Key observation: Steps 1-6 are entirely new engineering work. Only Step 7 (the actual inference) has existing infrastructure from v1.2. The entire developer onboarding flow – from landing page to API key delivery – must be built from scratch. This is why the `gonka_provider_landscape_map.md` identifies these as “not feature gaps but table-stakes infrastructure gaps that prevent developers from evaluating Gonka at all.”

The Atomic Growth Unit

The `openclaw.json` config snippet is Gonka’s atomic growth unit – the smallest shareable artifact that converts a non-user into a user.

Why the config snippet is viral:

1. **It is copy-pasteable.** A developer who has configured Gonka can share their provider config block (minus the API key) in a Reddit comment, Discord message, or GitHub issue. The recipient copies it, adds their own API key, and is running on Gonka in 90 seconds. No documentation reading required. No signup flow navigation. Just paste and go.

2. **It lives in version control.** `openclaw.json` is committed to Git repositories. When a developer adds Gonka as a provider, their teammates see the config change in the next pull request. The config snippet propagates through team workflows organically. When the Startup CTO adds Gonka to the team’s shared config, all 3-8 engineers inherit it automatically (`gonka_developer_personas.md`: Startup CTO team size 3-8).
3. **It is portable across OpenClaw installations.** The same config snippet works on any OpenClaw deployment – personal laptop, VPS, cloud server, Docker container. A developer who configures Gonka on their local machine can copy the same config to production. This portability means the snippet carries across environments without modification (`STACK.md`: OpenClaw’s declarative configuration model).
4. **It embeds Gonka’s cost advantage.** The `cost` field in the config snippet shows Gonka’s per-token rates directly in the developer’s configuration file. Every time a developer opens `openclaw.json`, they see Gonka’s rates next to their other providers. If Gonka’s rates are lower (Scenario B: \$0.35/\$1.75 vs Together AI: \$0.50/\$2.50 – `gonka_agent_pricing_analysis.md` Section 4), the cost advantage is visible without visiting a pricing page.
5. **ClawHub skills amplify the snippet.** The Gonka Provider Skill submitted to ClawHub (`gonka_partnership_playbook.m` ClawHub Submission Plan) teaches agents how to use Gonka-specific features. When a developer installs the skill, it includes the config snippet as part of the setup instructions. The skill and the snippet work together: the snippet connects the agent to Gonka, and the skill teaches the agent to use sessions, tiering, and memory (`gonka_partnership_playbook.md`: SKILL.md draft outline).

The viral loop:

Developer A configures Gonka (pastes config snippet)

- > Saves 73% on heartbeat costs (session persistence)
- > Shares cost comparison screenshot on Reddit/Twitter
- > Developer B sees post, copies config snippet from comment
- > Developer B configures Gonka in 90 seconds
- > Developer B saves 73% on heartbeat costs
- > Developer B shares their own comparison...

This loop is driven by the session persistence savings – the cost reduction is dramatic enough (73% at Active tier – `gonka_agent_pricing_analysis.md` Section 6) to motivate organic sharing, and the config snippet makes sharing frictionless. No referral program is needed to trigger the initial loop; the savings create natural word-of-mouth. The referral program (`gonka_developer_personas.md` Weekend Builder Referral stage: “\$5 credit for each referred developer who makes 100+ API calls”) amplifies an already-functioning viral mechanism.

What Competitors Require

Gonka’s time-to-first-inference target (under 5 minutes, email-only signup) creates a measurable friction advantage over the competitors that Weekend Builders and Startup CTOs are most likely evaluating.

OpenRouter (Weekend Builder’s Current Provider)

Step	OpenRouter Requirement	Time	Friction
1	Visit <code>openrouter.ai</code>	0s	None
2	Click “Sign In”	10s	None
3	GitHub OAuth login	30-60s	Requires GitHub account. OAuth redirect. Permission grant screen. Some developers hesitate to grant OAuth permissions to unknown services
4	Credit card for paid tier	120-180s	Credit card entry form. Address verification. Payment processing. This step does not exist for Gonka’s free tier
5	Browse 500+ model catalog, select model	60-120s	Model selection paralysis. “Which model do I pick?” No guidance for agent workloads. Developer must research models independently
6	Copy API key, configure application	60-90s	API key is available immediately after signup (no email verification)
Total		~5-8 min (free tier) or ~8-12 min (paid)	GitHub OAuth + credit card + model selection

Gonka advantage: No OAuth, no credit card, no model selection paralysis. Gonka offers one model (K2.5) in three tiers with automatic routing – the developer does not need to choose. Email-only signup eliminates the OAuth permission screen that causes hesitation. The free tier eliminates the credit card step entirely.

Source: gonka_competitive_feature_matrix.md (OpenRouter: 500+ models); gonka_developer_personas.md (Week-end Builder: OpenRouter is the default, zero-config provider).

Together AI (Startup CTO’s Primary Alternative)

Step	Together AI Requirement	Time	Friction
1	Visit api.together.ai	0s	None
2	Click “Sign Up”	10s	None
3	OAuth login (Google or GitHub)	30-60s	OAuth redirect. Permission grant. Account creation
4	Credit card required for all usage	120-180s	No free tier. Credit card is mandatory. Some developers abandon at this step for evaluation purposes
5	Select model from 200+ catalog	60-120s	Less paralysis than OpenRouter (200 vs 500+ models) but still requires research. K2.5 is one of many options
6	Copy API key, configure application	60-90s	API key generation is straightforward
Total		~5-8 min	OAuth + mandatory credit card + model selection

Gonka advantage: No credit card requirement (free tier exists). No model selection step (K2.5 with auto-tiering). Email-only signup (no OAuth). The Startup CTO can evaluate Gonka on staging without entering payment information or requesting procurement approval from their finance team.

Source: gonka_agent_pricing_analysis.md Section 4 (Together AI pricing: \$0.50/\$2.50); gonka_provider_landscape_map.md Segment 3 (Together AI: pay-per-use only, no free tier).

Self-Hosted vLLM (Privacy-First Builder’s Current Approach)

Step	vLLM Self-Hosting Requirement	Time	Friction
1	Acquire GPU hardware (\$2,000-10,000+)	Days to weeks	Hardware procurement. Budget approval. Delivery time
2	Install CUDA, Python, vLLM	30-60 min	Driver compatibility issues. Python environment management. Dependency conflicts
3	Download K2.5 model weights	20-60 min	Model weights are 10-50 GB depending on quantization. Requires HuggingFace account. Network bandwidth dependent
4	Configure vLLM server	15-30 min	Configuration options for GPU memory, quantization, batch size. Trial and error for optimal settings
5	Start server, verify inference works	5-10 min	Health checks, first inference warm-up, verify output quality
6	Configure OpenClaw to point to local vLLM	5-10 min	Same openclaw.json config as Gonka, but with localhost URL
Total		Hours to days	Hardware + software + model download + configuration

Gonka advantage: Under 5 minutes vs hours to days. No hardware procurement. No driver installation. No model download. The Privacy-First Builder gets the same open-weight K2.5 model with no content filtering, served on distributed infrastructure with no central logging – and they can start using it in 4 minutes instead of 4 hours. The tradeoff: Gonka’s privacy guarantees are architectural (no central log aggregation) rather than absolute (no TEE yet). The honest capabilities page (gonka_developer_personas.md Privacy-First Activation stage) communicates this clearly.

Source: gonka_developer_personas.md Privacy-First Pain Point #2 (self-hosting costs and maintenance); gonka_competitive_feature_matrix.md (K2.5: open-weight, no content filtering).

Failure Modes and Recovery

Each step in the time-to-first-inference plan has a specific failure mode that would break the under-5-minute target. These must be engineered against:

Step	Failure Mode	Impact	Prevention
3	Email submission rate-limited or blocked by spam filter	Developer cannot sign up	Use transactional email service with dedicated IP. Monitor deliverability rates. Provide “did not receive email?” link with resend option
4	Verification email delayed >2 minutes	Developer loses context, opens new tab, forgets about Gonka	Alert on email delivery latency >30s. Use multiple email service providers with failover. Display “email sent, check your inbox” with a countdown timer
5	API key page fails to load or copy button does not work	Developer cannot retrieve their key	Client-side key display (no server round-trip after page load). Fallback: display key as selectable text if clipboard API fails. Email the key as backup
6	Developer pastes config but misses Step 2 (model allowlisting)	Silent failure – agent ignores Gonka provider entirely	Show both steps on the same page with numbered instructions. Verification endpoint: <code>GET /v1/verify</code> that returns 200 if the key is valid, helping developers test before configuring their agent
7	First API call times out or returns 500	Permanent trust damage – developer will not try again (PITFALLS.md Pitfall 3)	Health-checked routing ensures requests go only to responsive nodes. Cold-start mitigation: keep K2.5 warm on minimum N nodes. Error response includes troubleshooting link

Cross-References

Every claim in this document traces to a specific Phase 15-19 output document. This section lists each source and what was drawn from it.

Document	What Was Used
gonka_developer_personas.md (Phase 16)	Three personas (Weekend Builder, Startup CTO, Privacy-First Builder) with decision drivers, pain points, adoption triggers, objections, and complete AAARRRP journey maps. Per-persona conversion drivers at each funnel stage. Drop-off risk analysis. Referral mechanisms.
gonka_agent_pricing_analysis.md (Phase 15)	Workload tier definitions (Casual: 31.5M tokens/month, Active: 161.5M, Heavy: 2.94B). Heartbeat overhead percentages (44%, 51%, 85%). Scenario B pricing (\$0.35/\$1.75). Session persistence savings (80% token reduction). Monthly cost projections per provider per tier. Competitor pricing data (Together AI, DeepInfra, OpenRouter, OpenAI, Anthropic).
gonka_competitive_feature_matrix.md (Phase 15)	8-dimension competitive comparison. WIN verdicts on Agent Sessions and Model Tiering. LOSE verdicts on Model Breadth and Uptime/Reliability. Competitor free tier analysis. Content filtering comparison (Gonka: “None (open)”). K2.5 quantization tiers.
gonka_provider_landscape_map.md (Phase 15)	5 critical must-close gaps (not built-in provider, no docs site, no self-serve signup, no pricing page, in-memory sessions). 4-segment landscape framework. Gonka’s unique position as only decentralized provider with agent-native extensions.
gonka_message_house.md (Phase 17)	Core positioning statement (73% cost reduction via session persistence). 16-item never-say list. Approved vocabulary mapping. VP1-VP3 ranked value propositions. Competitive differentiation statements per persona.
gonka_channel_strategy.md (Phase 18)	Primary KPI: API-active developers (>100 calls/month). Target trajectory (10 Month 1, 50 Month 3, 200 Month 6, 1000 Month 12). 70/30 channel split. P0 channels (OpenClaw Provider Directory, GitHub, Discord). Anti-metrics (follower counts, Discord members, GitHub stars).
gonka_partnership_playbook.md (Phase 19)	Four-tier OpenClaw integration roadmap (Listed -> Plugin -> Built-In -> Preferred). ClawHub submission plan. Technical requirements checklist (33 items). Community-first PR strategy timeline (3-6 months).
gonka_objection_playbook.md (Phase 17)	ACE objection framework. 12 documented objections with persona mapping. P0 severity objections: awareness, crypto perception, SLA gap, prompt privacy.
gonka_agent_native_pitch.md (Phase 17)	Agent-as-customer thesis. 7 task profiles with provider match analysis. Gonka wins 3/7 (long-running agents, multi-agent systems, privacy-sensitive).
.planning/research/STACK.md	OpenClaw <code>openclaw.json</code> configuration format. Two-step provider selection process (provider definition + model allowlisting). Built-in vs custom provider distinction. Key rotation and failover behavior.

Document: *gonka_plg_growth_model.md* / Version 1.0 / 2026-04-01 Capstone deliverable for Phases 15-19. Companion document: *gonka_v14_engineering_backlog.md* (Plan 20-02)